

World Models for Coding: A Critical Survey of Execution-Grounded Code Models, Agentic Evaluation, and Overloaded Terminology

Keon Kim

May 2026

Contents

World Models for Coding: A Critical Survey of Execution-Grounded Code Models, Agentic Evaluation, and Overloaded Terminology	1
Abstract	2
Theses Summary (Read First)	2
1. Introduction	3
2. Methodology	4
3. Defining a World Model for Coding	6
4. Twelve Years of Code World Models	7
5. Taxonomy: Three Axes of Code World Models	10
6. Foundations: Neural Execution as Implicit World Modeling	13
7. The Trace-Pretraining and CWM Lineage	13
8. World Models for Code Agents	18
9. RL with Execution as the World Signal	20
10. Planning and Search with Code World Models	21
11. JEPA, Dreamer, and the Latent-Action Gap	22
12. Specialized Domains	24
13. Reasoning, Process Rewards, Memory	24
14. Verification, Probing, Safety	25
15. Benchmarks and the Evaluation Gap	26
16. Empirical Landscape	27
17. Critical Perspectives	32
18. Open Problems	35
19. Conclusion	36
Appendix · Glossary	37
References	37

World Models for Coding: A Critical Survey of Execution-Grounded Code Models, Agentic Evaluation, and Overloaded Terminology

Last updated 2026-05

Genre note. This document is a hybrid: §§6–16 are a conventional survey of the field (foundations, lineages, benchmarks, empirical landscape), while §17 is an explicitly argumentative critical synthesis. Readers who want only the catalog can skip §17; readers who want only the criticism can read the Theses Summary below plus §17 and skip §§6–16. The two halves are meant to be read together — neither alone captures the field’s current shape.

Abstract

A *world model* is an internal predictor an agent maintains over the dynamics of its environment, used to imagine the consequences of actions. In coding, the environment is the program itself — its runtime state, its execution trace, the filesystem it manipulates, the tests it must satisfy, the developer task it is solving. Twelve years after Zaremba & Sutskever asked whether a network could execute code (Zaremba and Sutskever 2014), and seven months after Meta FAIR released CWM (team et al. 2025) as the first openly-released LLM explicitly branded a Code World Model, the question has changed. It is no longer controversial that internal models of execution are *learnable*; whether they are *necessary*, *architecturally separable*, or *causally responsible* for coding-agent gains remains unsettled. The open questions are whether the *named artifact* “code world model” is a structural commitment or a marketing label, whether trace pretraining buys what it claims, and whether the field’s Dreamer-shaped vocabulary will survive when the empirical evidence comes in.

This survey synthesizes 184 papers around the world-model lens. We define the object of study with two distinct definitions — a permissive *descriptive* one (what the field currently calls a code WM) and a strict *normative* one (architectural commitment to forward state prediction); trace its twelve-year arc; build a three-axis taxonomy (functionality × temporal × representation); produce technical system cards for thirteen representative systems; assemble protocol-stratified benchmark tables for SWE-bench, CRUX-Eval, web agents, and formal verification; develop seven critical theses where the field overclaims; and identify the open problems where new work would matter most. The single most defensible empirical claim is that *execution-grounded supervision improves code agents on runtime-reasoning-heavy benchmarks*; the single most defensible critical claim is that broader inferences from this to “world models in coding” remain underdetermined by the evidence at this date.

Theses Summary (Read First)

The survey is a hybrid: catalog plus critique. Readers pressed for time can read this summary and §17 (Critical Perspectives) and skip the lineage chapters (§§6–14). Readers entering the field should read §§1–5 first, then §§6–16, then §17. The catalog and the critique are meant to be read together.

The seven critical theses developed in §17, each grounded in a specific disconfirmation from the corpus:

1. The “world model” label has become overloaded. Under the field’s permissive descriptive definition (D, §3.1), CWM, TRACED, SemCoder, LLM-JEPA, and DyMo are all “world models” despite being architecturally standard LLMs with enriched training. A strict normative definition (N1/N2/N3, §3.1) admits only systems with explicit forward-prediction machinery: neural dynamics models (N1, no code exemplar), latent-action rollout models such as CoLA (N2), and synthesized executable simulators such as GIF-MCTS and ARC executable WMs (N3).

2. Trace pretraining has a causal-isolation problem. “Do Code Semantics Help?” (J. Wang et al. 2025) finds no trace representation consistently improves code synthesis across multiple backbones; CWM’s 65.8% SWE-bench Verified is confounded between trace mid-training, ForagerAgent trajectories, and RL.
3. The Dreamer-for-code gap may be a non-problem. Vision needed latent rollouts because pixels are expensive; Python state is small and CPython is free, so the architectural pressure does not transfer. CWM in token space reaches state-of-the-art at 32B.
4. PRMs are critics, not world models. A learned evaluator of partial trajectories cannot roll out future states; the conceptual conflation dilutes the world-model vocabulary.
5. “General Agents Contain World Models” (Richens et al. 2025) is weaker than its title. The theorem applies under restrictive assumptions (full observability, stationarity, deep-composite-goal regret bounds) that SWE agents demonstrably violate.
6. The verifier-grounded lineage is the actual leading edge, but scoped. ATLAS, Re:Form, CLEVER, VeriStruct, AutoRocq produce machine-checkable correctness. Scoped: they lead on synthesis-from-spec (e.g., ATLAS DafnySynthesis 65.8% pass@5); they do *not* lead on end-to-end verified codegen from NL (CLEVER 71/161 Lean).
7. The evaluation gap is the structural reason the field looks confused. No benchmark holds policy fixed and varies WM quality. Models with 85–98% output-prediction accuracy still produce traces with systematic errors throughout — outcome accuracy decouples from process fidelity.

These theses are the survey’s main contribution. §§6–16 catalog the work; §§16–17 provide the empirical and critical synthesis that grounds the theses; §18 lists open problems.

1. Introduction

Autoregressive code LLMs generate tokens conditioned on syntactic context. Correct programs live in two worlds: a *syntactic* world of tokens and a *semantic* world of values, control flow, side effects, and developer intent. The world-model framing — imported from model-based reinforcement learning, where it names an internal predictor of environment dynamics used to imagine action outcomes — is the field’s bet that the gap between these two worlds closes when the network has been trained on what code does rather than only on what code looks like.

Two adjacent surveys cover non-overlapping ground. A Survey on LLMs for Code Generation (Jiang et al. 2024) maps the code-LLM space without the world-model lens. Understanding World or Predicting Future (J. Ding et al. 2024) maps world models in general without the code lens. A Comprehensive Survey on World Models for Embodied AI (X. Li et al. 2025) maps embodied world models but excludes the coding domain. The intersection — the subject of this document — has cohered only recently into a recognizable program.

We aim for three things at once: comprehensive coverage of the corpus, technical depth on the canonical systems, and an honest accounting of where the prevailing rhetoric outruns the evidence. The first two are owed to readers entering the field; the third is owed to those already in it.

2. Methodology

Corpus construction. The 184-PDF corpus enumerated in `papers.json` was assembled in four iterative passes between March and May 2026. The seed paper was CWM (team et al. 2025). Each pass expanded the corpus along a different axis: (i) snowball search on cited and citing papers via Semantic Scholar’s reference API and arxiv listings; (ii) targeted topic searches on subdomains where the previous pass was thin (latent-action WMs for LLMs; safety/malicious-code; symbolic-execution and formal verification; agent memory; REPL-grounded models); (iii) a 2026-specific sweep to capture papers the earlier passes missed because of arxiv-ID ambiguity; (iv) a focused fourth pass on reasoning models, process reward models, test-generation, decompilation, diffusion code models, mech-interp, ARC synthesis, self-improvement, Verilog/RTL, and major 2026 capstones. Each pass produced a written rationale for accepted and rejected candidates.

Inclusion criteria. A paper was included if it credibly intersected *both* (a) world-modeling, state-tracking, or environment-prediction architectures, and (b) code generation, debugging, repair, or agentic coding. Pure code-LLM papers without a world-model angle and pure world-model papers without a code angle were excluded. Where the boundary was ambiguous, we erred on the side of inclusion when the paper introduced a representation, training objective, or evaluation that could plausibly be adopted by code-WM work.

Exclusion criteria. Excluded: pure transformer scaling papers; pure vision world models (DreamerV1/V2/V3, V-JEPA, Genie) except where cited as architectural precedent; generic RL theory; classical PL foundations; benchmark surveys without methodological contribution. We retained DreamerV1–V3 as named citations in §11 because they ground the latent-action discussion, but did not include their PDFs in the corpus.

Date cutoff. Papers with arxiv submission dates through 2026-05-15. Arxiv-ID format YYMM.NNNNN means IDs beginning 2601, 2602, 2603, 2604, 2605 are valid 2026 papers (Jan–May 2026), not future-dated. An earlier survey pass mistakenly excluded these; the fourth pass corrected the omission.

Selection bias and reproducibility. Two biases are explicit. First, anglocentric arxiv-first selection — we did not systematically cover Chinese-language preprint servers, workshop proceedings without arxiv copies, or industry technical reports. Second, recency bias — 60% of the corpus is from 2025 onward, reflecting the field’s actual growth curve but also our search recency. The full corpus is enumerated in `papers.json` with arxiv IDs, titles, and filenames; readers can regenerate the bibliography from this index.

Sources reconciled with adjacent surveys. We compared our taxonomy and corpus against three awesome-list repositories (knightnemo/Awesome-World-Models, JiahuaDong/Awesome-World-Models, tsinghua-fib-lab/World-Model) and two prior surveys (Ding et al. (J. Ding et al. 2024) on world models in general; Li et al. (X. Li et al. 2025) on world models for embodied AI). The three-axis taxonomy in §5 is adapted from Li et al. with code-specific representation classes added. Where our judgment diverges from these adjacent works — for example, in our classification of CWM as architecturally a token-policy rather than a world-model proper (§17.1) — we mark the divergence explicitly.

Taxonomy coding. Each paper was assigned at least one lineage label (§§6–13) and at least one representation label (§5.3 table). The assignments were made by one of us and are recorded in the per-section enumerations and `README.md`. We have not yet performed inter-rater reliability checks; readers should treat the assignments as one curator’s best judgment, not consensus.

Search procedure (for reproducibility). Seed: (team et al. 2025). Search engines: arxiv listing API (cs.SE, cs.PL, cs.LG, cs.CL submission categories), Semantic Scholar Graph API (/paper/arXiv:{id}/references and /paper/arXiv:{id}/citations, fields=title,externalIds,abstract,year), and WebSearch with query variants enumerated below. Date cutoff: arxiv submissions through 2026-05-15.

Representative search queries used across the four passes (non-exhaustive): - “world model code generation,” “code world model” - “neural code execution,” “learning to execute” - “execution trace pre-training LLM,” “trace-aware code model” - “Dreamer code,” “JEPA LLM,” “latent action LLM” - “SWE agent reinforcement learning,” “agentic coding world model” - “self-debugging execution simulation,” “trace-driven program repair” - “symbolic execution LLM,” “verified code generation Dafny,” “Lean code agent” - “diffusion code model,” “decompilation LLM semantic equivalence” - “Verilog RL world model,” “ARC executable world model” - “process reward model code,” “long-CoT code reasoning”

Inclusion / exclusion counts. Pass 1 considered ~70 candidates, accepted 31. Pass 2 considered ~40, accepted 20 (focused on 2026 papers earlier passes missed). Pass 3 considered ~80, accepted 52 via citation BFS from 11 central papers. Pass 4 considered ~60, accepted 28 across 5 thin niches. Final corpus: 184 PDFs (one was added late: (X. Li et al. 2025)). Total candidates considered: ~250. Rejection reasons recorded per-paper in the pass reports.

Evidence grading. Each citation in this survey can be classified by source type. We do not annotate inline (it would clutter the prose), but the following hierarchy is the implicit standard:

Tier	Definition	Survey use
A	Peer-reviewed venue (NeurIPS, ICML, ICLR, ACL, OOPSLA, FSE, etc.)	Highest weight; load-bearing claims
B	Major arxiv preprint with code/checkpoints + multi-author institutional backing	Treated as A for technical accuracy, less for community vetting
C	Solo-author or single-institution arxiv preprint	Cited for ideas; numerical claims hedged
D	Technical report / blog / benchmark leaderboard claim	Cited as evidence of capability, not as scientific finding

Roughly 60% of the corpus is currently Tier B; ~25% Tier A (mostly pre-2024); ~12% Tier C; ~3% Tier D. The recency skew toward 2025–2026 means many headline claims are Tier B and have not yet been independently replicated. Numbers in §16 should be read accordingly — they are the best available evidence at this date, not settled science.

Limitations of this methodology. Four. (1) No formal inter-rater reliability metric — one curator did the taxonomy coding. (2) No systematic verification of the headline numerical claims in every cited paper; §16 numbers were verified by reading the source PDFs but §§7–14 inline claims rely on abstracts and our prior reading. (3) The corpus is unstable — between drafts several 2026 papers we initially excluded as off-topic turned out to be relevant. A fifth pass would change the count by ± 15 papers without materially changing the survey’s conclusions. (4) Selection bias toward arxiv-first, anglophone, code-LLM-adjacent sources; we did not systematically cover Chinese-language preprint servers, workshop proceedings without arxiv copies, or industry technical reports.

3. Defining a World Model for Coding

The literature uses “world model” in two distinct senses that this survey will carefully separate.

3.1 Two definitions

Definition D (descriptive, permissive). A *world model for coding* is any code LLM whose training objective concretely encodes program semantics — execution traces, runtime state, environment feedback, or simulated outcomes — in addition to or instead of source-token prediction. This is the field’s current usage. Under D, CWM, TRACED, SemCoder, LLM-JEPA, DyMo, RLEF, and most papers in the corpus are world models.

Definition N (normative, strict). A *world model for coding* is a system with an explicit, separable forward-prediction mechanism. We split N into three sub-types:

- N1 — Neural dynamics model. A learned $W : (\text{state}, \text{action}) \rightarrow \text{next_state}$ instantiated as a distinct architectural component with latent recurrent state, separate dynamics head, or inverse model. Dreamer-class. No code exemplar in the corpus.
- N2 — Latent-action model. A learned action abstraction over a base LLM, with the LLM serving as transition model in compressed action space, supporting rollout or tree search over latent actions. CoLA (Jia et al. 2025) is the canonical example.
- N3 — Synthesized executable simulator. The world model is an executable program (Python, DSL) synthesized by an LLM and run against ground-truth transitions during planning. GIF-MCTS (Dainese et al. 2024), WorldCoder, Executable WMs for ARC-AGI-3 (Rodionov 2026).

Each N-subtype has a distinct architectural commitment. N1 commits to learned latent dynamics; N2 commits to a discrete latent-action space; N3 commits to program synthesis as the dynamics. Conflating them — as the loose “Dreamer-for-LLMs gesture” framing did — obscures which architectural bet is being made. Under any N-subtype, CWM, TRACED, SemCoder, and most of the corpus do *not* qualify.

The two definitions agree on the empirical fact that execution-grounded supervision helps. They disagree on whether that grounding is correctly named a *world model* in the model-based-RL sense.

Aspect	Definition D (descriptive)	Definition N (normative)
Granted to	Any LLM trained with execution-related signal	Only systems with explicit forward-prediction module
Includes CWM	Yes	No (architecturally a Llama-class decoder)
Includes TRACED, SemCoder, NExT	Yes	No (auxiliary heads on a Transformer)
Includes CoLA, GIF-MCTS	Yes	Yes
Includes PRMs (ExecVerify, ThinkPRM)	Often grouped	No (critics, not forward predictors)
Includes verifiers (Lean, Dafny, Z3)	Sometimes grouped	No (deterministic oracles, not learned WMs)
Includes Dreamer / V-JEPA (vision precedents)	Yes	Yes

This survey uses D throughout the catalog sections (§§6–16) because that is how the literature talks, and N in §17 because the critical synthesis depends on it. We mark the switch each time it matters. The split has been called for in prior reviews of the field and the two-definition framework is, in our reading, the cleanest resolution.

3.2 What is modeled

Orthogonal to D-vs-N, a fourth question asks *what* is modeled. The corpus splits across:

- variable values and stack frames (CWM, CodeExecutor)
- linear or branching execution traces (NExT, SemCoder, TRACED)
- test outcomes and runtime errors (LEVER, RLEF)
- environment/OS/web state (WebDreamer, Dyna-Think)
- repository state (RepoGraph, Understanding by Reconstruction)
- developer task or specification (ATLAS, Re:Form)
- adversarial behavior (Double Life of CWMs)

A given system typically commits strongly to one or two of these.

3.3 Behavioral vs architectural classification

CWM occupies an instructive position. *Behaviorally*, it emits stack frames and variable bindings, which feels like an “explicit symbolic world model.” *Architecturally*, it is a 32B Llama-class decoder with no separate dynamics head, RSSM, or inverse model. The two readings are both correct: CWM is behaviorally explicit but architecturally implicit. Earlier framings in the literature (and earlier drafts of this survey) collapsed these into a single “explicit WM” label; we recommend separating them. SemCoder, NExT, and TRACED follow the same pattern.

4. Twelve Years of Code World Models

The lineage is best understood as a sequence of inheriting questions. Each era’s answer dissolved the previous era’s bottleneck and exposed the next.

(For the text-only version of this diagram see *SURVEY.md* on the repository.)

Diamonds (◊) mark moments where “world model” enters the *name* of the contribution.

Pre-2020 — Can a network execute code? Zaremba & Sutskever’s Learning to Execute (Zaremba and Sutskever 2014) handed an LSTM character-level Python and asked it to predict output. The model worked on straight-line programs with bounded loops, but only with a curated curriculum, and only because the LSTM’s constant memory was just enough to simulate the interpreter when the interpreter ran in constant memory too. Everything that followed in this lineage tried to escape that curse. Neural Programmer-Interpreters (Reed and Freitas 2015) and the Differentiable Forth Interpreter built differentiable program counters and call stacks — the bet that the right architecture would close the gap. Dynamic Neural Program Embedding (K. Wang, Singh, and Su 2017) made the inverse move: run the real interpreter, embed the resulting state traces. Neural Code Fusion (Shi et al. 2019) and IPA-GNN (Bieber et al. 2020) extended the GNN-over-execution playbook to the point where attention played the role of the program counter. By 2020 the lineage had answered its question — yes, a neural network can play interpreter, but only when the interpreter is encoded into its architecture, and these architectures




Figure 1: Twelve-year arc of code world models (2014–2026). Three parallel lineages — neural execution, world models / RL, and code LLMs — converge in 2025–2026 around the named artifact CWM. [?] marks the first paper to use “world model” / “Dreamer” in a contribution name; [?] marks CWM as the named open-weights artifact.

did not transfer to Python, C, or assembly at corpus scale. Ha & Schmidhuber’s World Models paper (Ha and Schmidhuber 2018) had already named for vision and RL exactly the pattern this lineage was reaching for. The vocabulary existed; the coding community had not yet borrowed it.

2020–2022 — Can the network’s training include execution? The era’s reframing was simple: instead of *can the network execute*, ask *can we train a normally-shaped Transformer on enough execution evidence that semantics seep into its weights*. Codex (M. Chen et al. 2021) and MBPP (Austin et al. 2021) made code generation a real engineering target. Show Your Work / Scratchpads (Nye et al. 2021) made the decisive move: a Transformer that could not predict a program’s output could predict it perfectly if it was allowed to emit the intermediate computation first. Same trick Dynamic Neural Program Embedding had pulled, now at LLM scale, in token space, without architectural surgery. The neural-as-interpreter program of 2014–2020 did not scale beyond toy languages; the next era’s move was to bake execution into training data rather than architecture.

2023 — Trace pretraining as a named recipe. CodeExecutor (C. Liu et al. 2023) made the recipe explicit: mutate competitive-programming submissions, run them in a sandbox, capture per-line state tokens, train a transformer to emit the trace from source. TRACED (Y. Ding et al. 2023) generalized this to a pretraining auxiliary that any code-LLM could absorb. CRUXEval (A. Gu et al. 2024) provided the canonical input/output-prediction benchmark and suddenly there was a number that captured “does this model understand what code does, as opposed to what code looks like.” In parallel, Reflexion (Shinn et al. 2023) and Self-Debug (X. Chen et al. 2023) showed that an LLM’s mistakes could be fed back to itself as natural-language critiques, LEVER (Ni et al. 2023) used execution to verify candidate generations during decoding, and RAP (Hao et al. 2023) framed the LLM itself as a world model and ran MCTS over its imagined rollouts. The year’s unifying insight: execution traces are an auxiliary objective, not a separate model.

2024 — From models that simulate to agents that act. SWE-bench (Jimenez et al. 2023) replaced “write a function that passes a unit test” with “fix a real GitHub issue in a real repository.” CodeAct (X. Wang et al. 2024) claimed the agent’s action space should be Python code itself. SWE-agent (John Yang et al. 2024) shipped the harness. NExT (Ni et al. 2024) inlined traces into the agent loop. RLEF (Gehring et al. 2024) fed execution outcomes back as RL rewards. WebDreamer (Y. Gu et al. 2024) transplanted Dreamer-style imagination to digital agents. Generating Code World Models via MCTS (Dainese et al. 2024) introduced the literal phrase “Code World Models.” What unified the year was a structural shift in *where* the world model lives: in 2023 it lived in the weights, surfaced through trace prediction; in 2024 it lived in the loop, in the agent’s behavior under environmental feedback.

2025 — The CWM moment. DeepSeek-R1 (DeepSeek-AI et al. 2025) opened the year by showing that pure reasoning-RL on verifiable rewards could reach the frontier on math and code. SWE-RL (Y. Wei, Duchenne, et al. 2025) applied the same recipe to full SWE traces. CoLA (Jia et al. 2025) made the first concrete attempt at a Dreamer-for-LLMs: inverse-dynamics over latent actions, then RL over a learned codebook. LLM-JEPA (Huang, LeCun, and Balestriero 2025) ported LeCun’s joint-embedding predictive objective to language. General Agents Contain World Models (Richens et al. 2025) supplied a theorem: any agent satisfying a regret bound on goal-conditioned tasks must have learned a predictive model of its environment. Then in October, Meta FAIR released CWM (team et al. 2025) — a 32B open-weights model mid-trained on 5T tokens of Python execution traces and ForagerAgent trajectories from Dockerized repositories. The thing the 2014 LSTM was trying to be was now a downloadable checkpoint.

2026 — Stress-testing and broadening. With the artifact in hand, the field pivoted to critique and generalization. Debugging Code World Models (Rahmani 2026) catalogs CWM’s failures on long traces and

string-state representation. Demystifying Errors in LLM Reasoning Traces (Abdollahi et al. 2025) audits where trace-trained models hallucinate. Industrial CWM (Jian Yang et al. 2026) and Parallel-Code WMs (Singh et al. 2026) generalize the recipe to Verilog/GPU and parallelism semantics. Executable World Models for ARC-AGI-3 (Rodionov 2026) brings the generative-environment flavor to abstract visual reasoning. Reinforcement World Model Learning for LLM Agents (Yu et al. 2026) flips the standard recipe by training the WM rather than the policy.

The arc. Across twelve years: *can a network execute code?* [?] *can a network’s training include execution?* [?] *can an LLM agent simulate its environment?* [?] *is the world model a named artifact rather than a metaphor?* Each era’s answer dissolved the previous bottleneck and exposed the next. Architecture gave way to data; data gave way to agency; agency gave way to artifacts. What was a half-philosophical question in 2014 — *does this network understand what code does* — became, in 2025, an operational one with an open-weights baseline. The field has not finished. The 2026 critique wave shows the artifact is brittle in ways the 2014 LSTM was never asked to be. But the trajectory is now legible.

5. Taxonomy: Three Axes of Code World Models

Two cuts at the taxonomy are useful, and they are complementary. The first is a *lineage* cut — which research thread produced the system — and is the basis for §§6–13. The second, more durable cut adapts the three-axis framework of Li et al. ((X. Li et al. 2025), *A Comprehensive Survey on World Models for Embodied AI*) to the code domain. The lineage map is below; the three axes are developed in §§5.1–5.3.




Figure 2: Taxonomy of world models for coding. The three primary modeling axes (code, agents, tasks) are bridged by the JEPA / Dreamer / latent-action discussion in §11. Specialized domains and synthesis chapters sit below.

Adjacent WM surveys converge on overlapping splits that the three-axis framework subsumes as projections. Ding et al. (J. Ding et al. 2024) split top-level by *implicit representation* vs *future prediction*. JiahuaDong’s awesome-list organizes by *paradigm*: RL-based / observation-generative / latent-space / object-centric. knightnemo’s list surfaces *pixel* vs *mesh* vs *latent* as cross-cutting tags. The three axes — functionality, temporal modeling, and representation — capture the choices a system makes regardless of which lineage it belongs to.

5.1 Axis 1 — Functionality

- Decision-coupled WMs model only the slice of the world relevant to acting on it. CWM (team et al. 2025), RLEF (Gehring et al. 2024), and WebDreamer (Y. Gu et al. 2024) are decision-coupled — their WMs exist to enable code generation, RL planning, or web navigation respectively. CWM does not predict global filesystem state; it predicts the next Python frame *because* the next action depends on it.
- General-purpose WMs model the environment without reference to a particular task. The “general agents contain world models” theorem (Richens et al. 2025) is the abstract limit. In the corpus, only the largest CWM-class models with broad mid-training approach generality; most “code world models” are decision-coupled to a sub-task (repair, completion, agent control).

5.2 Axis 2 — Temporal Modeling

- Sequential simulation/inference. Step-by-step autoregressive rollout. CWM, NExT, SemCoder, all the trace-pretraining systems, and most LLM-as-WM planners (RAP, WebDreamer in its MPC loop) live here. The state is updated one timestep at a time. The vision analog is RSSM (Hafner et al., DreamerV1–V3).
- Global difference prediction. Predict the entire future state at once, in parallel. The vision analog is video-diffusion or masked-JEPA. In code, this fits diffusion code models (DiffuCoder, (Gong et al. 2025); Dream-Coder 7B, (Xie et al. 2025)) where the next state is sampled jointly rather than autoregressively, and the “specification is the program” framing (Monperrus 2026) where the entire trace is the spec.
- Static, no-trace. Some systems (SemCoder’s static mode, the trace-free baselines in “Do Code Semantics Help?”) explicitly drop temporal modeling at inference, reducing to single-shot prediction.

5.3 Axis 3 — Representation

This is the axis where the WM literature has converged most strongly, and where the code-WM literature is most uneven. Adapting Li et al.’s four-category split (GLV / TFS / SLG / DRR) to code yields five classes, of which only two are well-populated.

Class	Encodes the world as	Vision analog	Code exemplars
Token Sequence (TS)	Discrete or continuous token streams with execution traces, variable bindings, or rationales interleaved with source	Token-as-pixel (IRIS, TWM, Genie, Sora)	CWM (team et al. 2025), CodeExecutor (C. Liu et al. 2023), TRACED (Y. Ding et al. 2023), NExT (Ni et al. 2024), SemCoder (Y. Ding et al. 2024) — the dominant code-WM mode
Global Latent Vector (GLV)	A compact vector updated recurrently, encoding the entire program/agent state	RSSM (Hafner et al., DreamerV1–V3)	No clean exemplar. CoLA (Jia et al. 2025) introduces a learned action codebook but is otherwise a standard LLM, not RSSM-style
Spatial / Structural Grid (SLG)	A geometric or structural grid (BEV/voxel in vision; AST, call-graph, CFG in code)	OccWorld, DriveWorld	No exemplar. RepoGraph (Ouyang et al. 2024) uses a static dependency graph but does not predict over it as a WM
Decomposed Object / Slot (DOR)	Distinct persistent latent slots for objects in the world	SlotFormer and object-centric WMs	No exemplar. No code-WM models variables, scopes, or classes as discrete persistent slots
Synthesized executable (N3)	The world model <i>is</i> an executable program, synthesized rather than learned	(orthogonal to vision)	GIF-MCTS (Dainese et al. 2024), WorldCoder, Executable WMs for ARC-AGI-3 (Rodionov 2026)

Verifiers (Lean, Dafny, Z3) and PRMs are intentionally absent from this table. A verifier is not a *representation* of the world; it is a *grounding oracle* over candidate artifacts. A PRM is a *backward-looking critic*, not a forward predictor. Both belong in §5.4 below.

Three white spaces. The Dreamer-style GLV, the object-centric DOR, and the spatial-grid SLG representations have not been instantiated for code, despite being mature for vision. §18 lists these as open problems, with the caveat that not all three are equally promising — the Dreamer-style gap is contestable (§17.3), while the object-centric and structural-grid gaps look more genuine.

5.4 Grounding mode (orthogonal to representation)

A second classification asks how each system’s predictions are validated. This is the analog of Li et al.’s “reality” column and addresses the decoupling between WM-head accuracy and downstream task success that DyMo (Guo et al. 2025) and §17.7 develop.

Grounding mode	Definition	Code exemplars
None / self-report	Predictions never validated against ground truth	RAP (Hao et al. 2023), basic LLM-as-WM planners
Execution-grounded	Predictions checked against real interpreter / runtime	CWM (team et al. 2025), TRACED (Y. Ding et al. 2023), NExT (Ni et al. 2024), SemCoder (Y. Ding et al. 2024), RLEF (Gehring et al. 2024)
Verifier-grounded	Outputs checked by Lean / Dafny / Verus / Rocq / Z3	ATLAS (Baksys et al. 2025), Re:Form (C. Yan et al. 2025), CLEVER (Thakur et al. 2025), VeriStruct (C. Sun et al. 2025), AutoRocq (Tu et al. 2025)
Synthesized-simulator-checked	Synthesized world model checked against held-out transitions	GIF-MCTS (Dainese et al. 2024), Executable WMs for ARC-AGI-3 (Rodionov 2026)
Critic-grounded (not a WM)	Backward-looking value/quality score, no rollout	ExecVerify (Tang et al. 2026), SWE-PRM (Gandhi et al. 2025), ThinkPRM (Khalifa et al. 2025) — listed for contrast; these are critics, not WMs (§17.4)

Grounding mode is orthogonal to the representation axis: a token-sequence representation can be execution-grounded (CWM) or self-report (RAP); an N3 synthesized-simulator can be checked against transitions (GIF-MCTS) or never validated. The “WM fidelity” question of §17.7 is, structurally, a question about whether a system’s grounding mode is strong enough to falsify a bad WM at training time.

6. Foundations: Neural Execution as Implicit World Modeling

Zaremba & Sutskever’s Learning to Execute (Zaremba and Sutskever 2014) established both feasibility and brittleness. Show Your Work — Scratchpads (Nye et al. 2021) is the hinge moment: by training a Transformer to emit intermediate computation states, the authors recovered much of the LSTM-era execution-prediction performance at scale, presaging the trace-pretraining lineage of §6. CRUXEval (A. Gu et al. 2024) and REval (J. Chen et al. 2024) provide the canonical execution-reasoning benchmarks. The lesson the field absorbed: *replacing* the interpreter with a neural network is harder than *augmenting* a transformer with interpreter-style supervision. Modern systems all take the latter path.

7. The Trace-Pretraining and CWM Lineage

7.1 Trace-pretraining as a recipe

CodeExecutor (C. Liu et al. 2023) trains a Transformer to simulate Python execution token-by-token. TRACED (Y. Ding et al. 2023) adds dynamic-state supervision to a code-LLM pretraining mix. NExT (Ni et al. 2024) formats traces as natural-language rationales, letting a chat-style LLM reason about runtime

behavior via chain-of-thought. SemCoder (Y. Ding et al. 2024) generalizes to “monologue reasoning” linking source-text to execution state.

The 2025 wave consolidated and stress-tested the approach. “What I cannot execute, I do not understand” (Armengol-Estapé et al. 2025) trains and evaluates LLMs explicitly on traces with dynamic scratchpads, pushing Llama-3.1-8B from 37.8% to ~80% on CRUXEval-O. Code Execution as Grounded Supervision (Jung, Zhou, and Chen 2025) repurposes line-by-line traces as verifiable CoT. Self-Execution Simulation (Maimon et al. 2026) lets the model train on its own execution predictions. Demystifying Errors in LLM Reasoning Traces (Abdollahi et al. 2025) audits where trace-trained LLMs fail. “Do Code Semantics Help?” (J. Wang et al. 2025) is the most damaging paper in the lineage: a comprehensive ablation across DeepSeek-Coder, LLaMA-3, and Gemma-2 with five trace representations finds that *no single representation consistently improves code generation*, and in 7 of 9 synthesis settings the no-trace baseline wins or ties.

7.2 Deep dive — TRACED (Y. Ding et al. 2023)

What it models	Per-line program state (concrete variable values quantized into 30 bins crossing data-type × value-range) and per-line execution coverage of a C program, statically predicted from source.
Data	CodeNet C subset: 1,805/1,900 problems, 121,319 training traces collected via gdb stepping through -g -O0 builds.
Architecture / objective	RoBERTa-base initialized from UnixCoder. Three jointly-optimized heads on input [CLS] $e_1 \dots e_i$ [SEP] [SEP] $c_1 \dots c_j$ [SEP]: MLM, per-variable program-state classification (data_type, value_type, quantized_value), per-variable binary coverage.
Headline results	Static execution full-path accuracy 71.6% vs UnixCoder 63.7% (+12.4% relative); variable-value accuracy 89.2% vs 87.8%; POJ-104 clone-retrieval MAP@R 91.2 vs 89.5.
Distinctive contribution	An early demonstration that <i>quantized</i> variable-value prediction is a viable pretraining signal — concrete values lose to discretized bins.

7.3 Deep dive — NExT (Ni et al. 2024)

What it models	Program execution as inline-comment trace appended to source: each statement annotated # (k) varA=...; varB=... capturing changed variables in execution order.
----------------	---

Data	Mbpp-R (10,047 train / 1,468 dev repair tasks built from incorrect LLM outputs on MBPP) and HumanEvalFix-Plus. Base model: PaLM 2-L.
Training objective	STaR-style self-training. Sample 32 (rationale, fix) candidates at T=0.8 from current model; filter by unit-test pass; SFT on accepted. Always restart from initial checkpoint each iteration. 10 iterations.
Headline results	Mbpp-R pass@1 23.2 $\uparrow$ 49.3 (+26.1 abs); HumanEvalFix-Plus 32.2 $\uparrow$ 42.5 (+10.3). Crucially, generalizes when traces are removed at test time (40.8 vs 23.2).
Distinctive contribution	Inline-comment trace format fits ~95% of MBPP into 2K window (vs ~60% for Scratchpad), and the trained model retains benefit even without traces at inference.

7.4 Deep dive — SemCoder (Y. Ding et al. 2024)

What it models	Four jointly-trained semantics modalities: Approximate (NL docstring), Symbolic (source), Operational (execution effects), Abstract (input-invariant properties).
Data	PYX — synthetically curated, retried by generator until each sample executes; 4.3k decontaminated CodeContest problems for SemCoder-S. Base: DeepSeekCoder-6.7B-base.
Training objective	Standard NTP with loss on code + monologue tokens. Forward monologue verbalizes execution step-by-step (summarizing loop iterations rather than dumping every state). Backward monologue characterizes possible prior states given output (abstract constraints + concrete witness). Both rejection-sampled against ground-truth execution.
Headline results	HumanEval+ 79.3% / MBPP+ 79.9% / LCB-Lite 27.5% at 6.7B — beats GPT-3.5-turbo. CRUXEval-I 63.6 / CRUXEval-O 65.1 vs GPT-3.5-turbo 50.3 / 59.0. Monologue beats Scratchpad (48.8) and NExT (49.4) on CRUXEval-I: 61.8.

Distinctive contribution	Forward AND backward monologues (NExT is forward-only); abstract semantics constraints rather than concrete state at every step; entirely static at inference.
7.5 Deep dive — CWM (team et al. 2025)	
What it models	Python interpreter state at the granularity of an <i>interpreter stack frame</i> (one observation–action pair per executed line), plus agentic SWE trajectories where actions are shell tool calls and observations are environment responses.
Architecture	32B dense decoder-only Transformer with GQA, sliding-window blocks, RoPE — Llama-class. No separate dynamics head, no inverse model, no recurrent latent.
Data composition	Four-stage train: (i) 8T-token general pretraining at 8k ctx; (ii) 5T-token code-world-modeling mid-training at 131k ctx: 120M traced Python functions, 262k CodeContests traces, ~70k repo-level traced commits, 75M natural-language trace rewrites, and 3M ForagerAgent trajectories from 10.2k Docker images / 3.15k repos (55% issue-fix, 45% synthetic mutate-fix); (iii) 100B-token SFT at 32k; (iv) 172B-token joint RL at 131k.
Trace format	Per-line <code><\ frame_sep\ >{locals JSON}<\ action_sep\ >{source line}</code> so that next-token prediction is next-state prediction at line granularity.
Headline results	65.8% SWE-bench Verified with TTS (best@16 over 40 verifier-reranked samples); 68.6% LiveCodeBench-v5; 94.3% CRUXEval-Output; competitive with much larger closed models.
Distinctive contribution	To our knowledge the first open-weights model where line-level Python execution traces and large-scale (3M) agent–environment trajectories are mid-training data, not post-training data. Introduces <i>Activ</i> (using GitHub Actions CI for local image builds) to scale executable repository images.

Important caveat (developed further in §17): the 65.8% headline is *not* pure pass@1 but best-of-16

with verifier reranking. Pure pass@1 is approximately 53–55%. The trace-mid-training contribution is not causally isolated from the ForagerAgent-trajectory contribution and from the joint-RL contribution. Without an ablation removing one while holding the others fixed, the “world model” component’s causal role is unfalsifiable.

7.6 Direct descendants of CWM

- Debugging Code World Models (Rahmani 2026) — probes where CWM fails on long traces and string state; finds long-horizon failures are dominated by *action hallucination*, not state-propagation error.
- Learning Reasoning World Models for Parallel Code (Singh et al. 2026) — predicts race conditions and profiling artifacts from parallel source.
- Industrial CWM / InCoder-32B-Thinking (Jian Yang et al. 2026) — CWM recipe on Verilog and GPU execution traces.
- The Double Life of Code World Models (Sahoo 2025) — CWM trace predictions repurposed for malicious-behavior detection.
- Towards a Neural Debugger for Python (Beck et al. 2026) — neural debugger as forward/inverse world model.
- Neural Computers (Zhuge et al. 2026) — video-model-style WMs of CLI/GUI runtime from I/O traces.
- Generating Code World Models with LLMs Guided by MCTS (Dainese et al. 2024) — the WM is *the code itself*, synthesized by an LLM.
- General Agents Contain World Models (Richens et al. 2025) — proves that sufficiently competent goal-conditioned agents must contain extractable world models, under restrictive conditions discussed critically in §17.

7.7 Deep dive — GIF-MCTS / Generating Code World Models via MCTS (Dainese et al. 2024)

What it models	The world model itself is <i>Python code</i> — an <code>Environment.step(s,a) → (s', r, done)</code> class synthesized by an LLM to match a small set of pre-collected (s, a, r, s', d) transitions.
Data	CWMB benchmark — 18 RL environments (classic control, PyGame, MuJoCo) with NL descriptions and 5 random + 5 above-threshold demonstrations each. Plus APPS-Competition and RTFM.
Algorithm	Inference-time MCTS over partial programs with three action types: generate (append $L=2$ lines), improve (rewrite full program given failing transition), fix (repair runtime/syntax error). Reward = fraction of transitions correctly reproduced.

Headline results	APPS-Competition 28.3% pass@20 (Llama-3-70B), beating WorldCoder 25.1%. CWMB normalized return 0.76 vs WorldCoder 0.60. RTFM: GPT-4 reaches 1.00 accuracy.
Distinctive contribution	The world model <i>as code to be searched over</i> , not as a neural net to be trained. Once compiled and validated, runs 4–6 orders of magnitude faster than calling the LLM as WM.

8. World Models for Code Agents

Once an LLM is an *agent* taking actions in a non-trivial environment, the world-model question becomes whether the agent simulates the environment’s response. Three sub-environments dominate.

8.1 Web agents

Web Agents with World Models (Chae et al. 2024) systematizes the thread. DyMo / World Modeling Improves LM Agents (Guo et al. 2025) adds a next-state prediction head to function-calling agents and reports gains on BFCL-V2 — though with a caveat (§17): the WM head reaches 90–94% state-prediction accuracy while the underlying policy reaches only 72.8% task success, illustrating that WM-head accuracy and agent accuracy can decouple.

Deep dive — WebDreamer (Y. Gu et al. 2024)

What it models	Web environment as a POMDP where the LLM imagines natural-language state-change descriptions for each candidate action.
Data	3.1M synthesized web-interaction instances from Common Crawl URLs via biased random walking; state changes captured as before/after VLM screenshots described by Qwen2-VL-72B.
Training	Dreamer-7B trained from Qwen2-VL-7B on (initial state, action) → state-change description. Horizon H=1 empirically optimal.
Planning	Model Predictive Control — simulate each candidate action, GPT-4o scores trajectories on a 3-scale rubric, argmax executes.
Headline results	VisualWebArena 23.6 vs reactive 17.6 (+34.1% rel); Online-Mind2Web 37.0 vs 26.0 (+42.3% rel); Mind2Web-Live 25.0 vs 20.2 (+23.8% rel). 4-5× more efficient than tree search. Dreamer-7B specialist [7] GPT-4o on online benchmarks.

Distinctive contribution	An early open demonstration that LLM-as-WM + 1-step MPC beats tree search on web tasks where irreversibility makes backtracking infeasible.
--------------------------	---

8.2 OS / computer-use agents

Reinforcement World Model Learning for LLM-based Agents (Yu et al. 2026) and World Models as an Intermediary between Agents and the Real World (S. Yang 2026) generalize the lens: a learned WM mediates between LLM and expensive environment.

Deep dive — Dyna-Think (Yu et al. 2025)

What it models	A single Qwen2.5-32B that internalizes world-model simulation inside its <think> block — predicting next observation, action critique, or state-difference — for computer-use agents.
Training	DIT (imitation learning): few-shot prompt GPT-4o to reconstruct DeepSeek-R1’s CoT keeping only WM-simulation-related text; SFT on cleaned trace. DDT (Dyna-style RL): online rollouts feed three WM objectives — next-state, state-diff, critic — jointly with rejection-sampled policy training.
Headline results	OSWorld BoN All 43.1 (DDT, 32B) vs R1-685B 44.8 — matches 685B model at 5% of parameters and 2× fewer tokens. WindowsAgentArena 34.9 vs Qwen2.5-32B 23.9 / R1 26.9. World-model accuracy correlates with task success at $r=0.32$ across models.
Distinctive contribution	Among the first Dyna-Q-style integrations where a single LLM hosts both policy and world model with critique-prediction as the WM objective.

8.3 SWE agents

SWE-bench (Jimenez et al. 2023) and SWE-Gym (Pan et al. 2024) defined the eval and training environment respectively. CodeAct (X. Wang et al. 2024) made the Python interpreter the unified action space. Reflexion (Shinn et al. 2023) was the earliest entry with episodic verbal RL. Nanbeige SWE-World (S. Sun et al. 2026) trains a learned Docker-free execution surrogate. Understanding by Reconstruction (Zeng et al. 2026) reverses the development process to harvest agentic pretraining traces. SWE-TRACE (Han et al. 2026) provides process-level reward modeling over trajectories. Self-Play SWE-RL (Y. Wei, Sun, et al. 2025) introduces adversarial bug-injection/repair self-play. Bootstrapping Coding Agents — The Specification Is the Program (Monperrus 2026) reframes the SWE task itself as a programmatic spec.

The §16 empirical synthesis separates three regimes: execution-grounded open-weight model training (CWM, SWE-RL), execution-grounded agents with learned simulators (Nanbeige SWE-World), and scaffold evolution around closed-model executors (Darwin GM, Huxley GM). Under best-of-k or verifier-reranked protocols, several systems report 60–68% on SWE-bench Verified, but those numbers are not directly comparable to pass@1 model-training results, and the scaffold-evolved systems are not “32B open-weight world-model-trained” — they run frontier closed models inside an evolved harness.

9. RL with Execution as the World Signal

The model-based-RL framing — world model is what the policy plans over — has produced a clean lineage.

9.1 Deep dive — RLEF (Gehring et al. 2024)

What it models	Iterative code synthesis as a POMDP — actions are full code responses, observations are formatted public-test execution feedback, rewards come from held-out private tests.
Data	CodeContests train (13,328 problems, 669 discarded for missing tests). Initial policies Llama-3.0/3.1-Instruct 8B / 70B.
Objective	PPO with $R(s, a) = r(s, a) - \beta \cdot \log \frac{\pi(a \mid c)}{\rho(a \mid c)}$. $r = +1$ if all tests pass, -1 if any fail, -0.2 for malformed output. Turn limit 3. Geometric-mean response probability for KL bias correction.
Headline results	Llama-3.1-70B + RLEF: 37.5 / 40.1 valid/test pass@1 with budget 1@3 (vs 25.9 / 27.5 baseline). 54.5 / 54.5 at 10@100, surpassing AlphaCode 41B+clustering. Transfers to HumanEval+ (78.6 → 80.4) and MBPP+. Random-feedback ablation removes all gain.
Distinctive contribution	An early clean demonstration that standard PPO on multi-turn execution feedback beats both SFT and few-shot for code agents; isolates that the model learns to <i>use</i> feedback, not just sample more.

9.2 Deep dive — SWE-RL (Y. Wei, Duchenne, et al. 2025)

What it models	Bug-fix as single-shot search/replace edit conditioned on issue + full file context. <i>No execution at training time.</i>
Data	273k high-quality PR seeds extracted from a raw GitHub PR corpus. Trained on Llama-3.3-70B-Instruct.
Objective	GRPO with rule-based reward $R(o) = \text{compare}(\text{patch_pred}, \text{patch_gt})$ via <code>difflib.SequenceMatcher</code> (continuous 0..1). -1 for format violations. Continuous reward beats discrete in ablation (34.8 vs 29.0 oracle-repair).
Headline results	41.0% SWE-bench Verified pass@1 with Agentless Mini scaffold. OOD: HumanEval+ 76.2 79.9; CRUXEval-O 61.9 75.5; MATH 70.9 73.7 — <i>SFT degrades on these while RL improves.</i>
Distinctive contribution	Continuous-similarity reward on PR patches without execution induces emergent self-reflection, multi-approach exploration, and divide-and-conquer reasoning that transfers OOD. The “world model” is implicit in the patch-similarity reward.

9.3 Process Reward Models

ExecVerify (Tang et al. 2026), SWE-PRM (Gandhi et al. 2025), DataPRM (Qiu et al. 2026), ThinkPRM (Khalifa et al. 2025) form a cluster where the WM is a learned *evaluator* of partial trajectories. As §17 develops critically, this is not the same object as a forward world model — PRMs are critics with execution grounding. They cannot roll out, cannot simulate counterfactuals. Survey hygiene argues for keeping the distinction.

10. Planning and Search with Code World Models

10.1 Deep dive — RAP (Hao et al. 2023)

What it models	Generic reasoning MDP — state = textual configuration, action = step proposed by the same LLM, transition obtained by re-prompting the LLM as world model.
Method	MCTS-UCT over the reasoning tree. Rewards: action likelihood, state confidence (majority voting), self-evaluation (LLM “Is this correct?” probability), task heuristics.

Headline results	Blocksworld 4-step: RAP@10 = 0.86 (LLaMA-33B) vs CoT-pass@10 = 0.07 vs GPT-4+CoT = 0.63. LLaMA-33B+RAP beats GPT-4+CoT by 33% relative on plan generation. GSM8K: 51.6 (RAP+aggr) vs CoT+SC 46.8.
Distinctive contribution	Earliest clean formulation that repurposes the LLM as both policy and transition model under MCTS. The template every later “LLM-as-WM” paper extends.

Tree of Thoughts (Yao et al. 2023), AlphaZero-like Tree Search for LLM Decoding (Feng et al. 2023), Tree Search for LM Agents (Koh et al. 2024), and Mastering Board Games by External/Internal Planning with LMs (Schultz et al. 2024) develop the search frame; the last gives the clearest contemporary recipe for learned tree-search with LLM-as-WM, straightforwardly transferable to code.

10.2 Execution-conditioned generation

Execution Guided Line-by-Line Code Generation (Lavon, Katz, and Wolf 2025) uses classifier-free guidance to condition next-token prediction on candidate-runtime outcomes. Jupiter (S. Li et al. 2025) formulates notebook state as MCTS nodes. REPL-Plan (A. Z. Liu et al. 2024) reuses a REPL state pool across tasks. Substrate is well-developed for short-horizon code-gen; less so for long-horizon multi-file SWE.

11. JEPA, Dreamer, and the Latent-Action Gap

LeCun’s Joint Embedding Predictive Architecture (I-JEPA, (Assran et al. 2023)) predicts in embedding space rather than pixel space. The Dreamer family — Hafner et al.’s DreamerV1 (Hafner et al. 2019), DreamerV2 (Hafner et al. 2020), and DreamerV3 (Hafner et al. 2023), built around the Recurrent State-Space Model — has near-zero direct application to code. Two papers occupy the gap.

11.1 Deep dive — LLM-JEPA (Huang, LeCun, and Balestrierio 2025)

What it models	A <i>joint embedding</i> between two views of the same knowledge — Text (NL prompt) and Code (e.g., SQL). Not a temporal world model; an embedding-space abstraction objective.
Architecture	Predictor is tied-weights: a single [PRED] token (with K predictor tokens) is appended and the LLM re-runs to produce $\text{Pred}(\text{Enc}(\cdot))$. A custom block-causal attention mask lets both views go through a single forward pass.

Objective	$L = \sum L_NTP(\text{text}) + \lambda \cdot d(\text{Pred}(\text{Enc}(\text{Text})), \text{Enc}(\text{Code})).$ $d = \text{cosine.}$ Encoder reuses last-layer last-token hidden state.
Headline results	Llama-3.2-1B on NL-RX-SYNTH FT: 71.46% vs 57.29% NTP-FT (+14.2). Spider: ~50% vs ~47%. GSM8K: ~32% vs ~32%. Top-100 singular values of $\text{Enc}(\text{Text}) - \text{Enc}(\text{Code})$ collapse by orders of magnitude.
Distinctive contribution	To our knowledge the first JEPA-style embedding-space objective added to a generative LLM that preserves NTP loss while inducing low-rank $\text{Text} \rightarrow \text{Code}$ mapping. Critical question (§17): is this really JEPA in the LeCun sense, or a regularizer on LM training?

11.2 Deep dive — CoLA (Jia et al. 2025)

What it models	An MDP over text where the LLM is the transition model and actions are discrete <i>latent</i> tokens from a learned codebook, not vocabulary tokens.
Data	Llama-3.1-8B base, continued-pretrained on 200GB from SlimPajama / StarCoder / Proof-Pile-2 / WuDao; policy trained on 100GB subset.
Three modules	Inverse-Dynamics Model $f_inverse(x_1:t, x_{t+1}) \rightarrow a_t$ implemented as VQ-VAE-style encoder with codebook C. Language World Model inserts the chosen latent action into the LLM embedding stream and decodes the next token. Policy $\pi(a_t)$
Planning	Action-level MCTS over latent-action subtrees. MCTS-Q variant uses Double-DQN over (prompt, response, reward) tuples.
Headline results	Math-500: 42.4 (CoLA+RL) vs 38.2 baseline. MCTS-Q on Math-500: 68.2 vs 63.0 baseline. MCTS-Q. +11% averaged on math reasoning; 64% win rate on alignment.
Distinctive contribution	Among the first to <i>replace</i> the 128k token-level action space of an LLM with a small learned latent-action codebook for RL — making the action space tractable for tree search.

11.3 The gap

Despite CWM and dozens of LLM-as-world-model papers, *no public Dreamer/RSSM-style latent-imagination world model has been trained for SWE agents*. CWM rolls out in token space. CoLA is the closest concrete instance. UniZero (Pu et al. 2024) generalizes MuZero with transformers but is rarely instantiated on code. Genie (Bruce et al. 2024) gives the vision-side template. JEPA for RL (Kenneweg, Kenneweg, and Hammer 2025) extends the energy-based objective to RL.

Whether the gap matters is a live question developed critically in §17. The vision-domain pressure that motivated Dreamer’s RSSM design (pixel-space rollout cost) does not exist for code, where state is small and the simulator is available. The argument *for* latent imagination rests on inference-time speed and the action-space compression CoLA demonstrates, not on rollout cost per se.

12. Specialized Domains

Diffusion code models. DiffuCoder (Gong et al. 2025), Dream-Coder 7B (Xie et al. 2025). Iterative denoising naturally accommodates plan-then-refine generation.

Decompilation and cross-language. SK2Decompile (Tan et al. 2025), SALT4Decompile (Yongpan Wang et al. 2025). Translation as semantic-simulation task. EquiBench (A. Wei, Cao, et al. 2025) supplies the equivalence eval.

Hardware / RTL. VeriRL (Teng et al. 2025), ChipSeek (Z. Chen et al. 2025), VeriCoder (A. Wei, Tan, et al. 2025) form a cluster where the simulator is the world model. Hardware is an attractive domain because simulators are precise, fast, and deterministic — closer to Atari than to Python.

ARC and abstract synthesis. Executable World Models for ARC-AGI-3 (Rodionov 2026) instantiates literal-WM-per-task: synthesize a Python world model verified against observations. SOAR (Pourcel, Colas, and Oudeyer 2025) evolves programs over ARC. Darwin / Huxley Godel Machines ((J. Zhang et al. 2025), (Wenyi Wang et al. 2025)) close the self-improvement loop.

13. Reasoning, Process Rewards, Memory

Long-CoT reasoning for code. o1-Coder (Yuxiang Zhang et al. 2024) replicates o1 with MCTS+RL. R1-Code-Interpreter (Y. Chen et al. 2025) supplies the open SFT+RL recipe across 144 tasks. Scaling Test-Time Compute to Achieve IOI Gold Medal (Samadi et al. 2025) shows open-weight gpt-oss-120b matching closed reasoning models via inference-time scaling.

Long-CoT reasoning is *mental execution* — the chain-of-thought simulates the world model the network never explicitly trained. CWM-style explicit world modeling and R1-style reasoning are partial substitutes; whether they compose multiplicatively is open.

Memory. Episodic Memory is the Missing Piece for Long-Term LLM Agents (Pink et al. 2025) frames the gap. Memory as Action (Yuxiang Zhang et al. 2025) treats memory operations as RL-learnable actions. RepoGraph (Ouyang et al. 2024) provides a durable repo-level dependency graph.

14. Verification, Probing, Safety

14.1 Formal verification: the leading edge

The verifier-grounded lineage is the only research direction in the corpus that does not rely on LLM self-report for correctness — the verifier provides ground truth.

Deep dive — ATLAS (Baksys et al. 2025)

What it models	Verifier-grounded synthesis of Dafny programs (specification + implementation + proof annotations) from NL + Python reference + tests.
Data	TACO-verified yields 2,751 verified Dafny programs decomposed into 19,385 training examples across 6 tasks: NL-to-Code, NL-to-Spec, Spec-to-Code, Spec-Repair, Impl-Repair, Proof-Infilling. Base: Qwen-2.5-Coder-7B + LoRA.
Spec quality	Three lemma types: Soundness (contract holds on test inputs), Completeness-Contradiction (negated output $\neg$ false), Completeness-Perturbation (contract rejects structurally perturbed outputs).
Headline results	DafnyBench Pass@1: 32.4 $\rightarrow$ 55.8 (+23.4); Pass@10 $\rightarrow$ 56.9. DafnySynthesis Pass@5: 15.8 $\rightarrow$ 65.8 (+50), surpassing GPT-4 (53.4) by 12.4 points at 7B.
Distinctive contribution	An early end-to-end pipeline for scaling verified-code data the way auto-formalization scaled Lean theorems. Operational soundness/completeness criteria via SMT-discharged lemmas filter degenerate specs.

The cluster also includes Re:Form ((C. Yan et al. 2025), Dafny+RL), CLEVER ((Thakur et al. 2025), Lean), VeriStruct ((C. Sun et al. 2025), Verus), AutoRocq ((Tu et al. 2025), Rocq), and Semantic Equivalence Self-Play with Formal Verification ((Barone and Nok 2026), Liquid Haskell).

CLEVER’s $\frac{1}{161}$ end-to-end Lean result is the most sobering number in the corpus: frontier models, with Lean type-checker access for self-verification, still fail on >99% of HumanEval-derived problems requiring joint spec + implementation verification. Understanding Formal Reasoning Failures in LLMs as Abstract Interpreters (Mitchell et al. 2025) is the diagnostic: when asked to reason in the style of formal abstract interpretation over 22 SV-COMP programs, all frontier reasoning models make systematic errors in widening, fixpoint termination, and join operations.

14.2 Symbolic execution and LLMs

AutoBug (Yihe Li, Meng, and Duck 2025), SESpec (F. Yang et al. 2025), LLM-Sym (Wenhan Wang et al. 2024), Loop Invariant Generation via Reasoning LLMs + SMT (Bharti et al. 2025) combine LLMs with concrete or symbolic engines. The unifying pattern: the LLM hypothesizes the world model; symbolic execution verifies or extends it.

14.3 Probing and mechanistic interpretability

Mechanistic Interpretability of Code Correctness via SAEs (Tahimic and Cheng 2025) and On LLMs’ Internal Representation of Code Correctness (Ribeiro et al. 2025) ask what code LLMs actually represent. Findings: *partial, brittle* internal execution representations — a vindication of explicit trace pretraining.

14.4 Repair and debugging as world-model probing

Self-Debug (X. Chen et al. 2023), InspectCoder (Yunkun Wang et al. 2025), Agent That Debugs — Dynamic State-Guided Vulnerability Repair (Z. Liu et al. 2025), Agentic Code Reasoning ((Ugare and Chandra 2026), semi-formal execution-path reasoning without running code). Shared pattern: maintain a belief over program state, query the runtime to update the belief, act on the posterior — Bayesian world-modeling in everything but name.

14.5 Safety and malicious code

The Double Life of Code World Models (Sahoo 2025) repurposes CWM-style trace predictions for malicious-behavior detection. CodeBreaker (S. Yan et al. 2024) is the offensive analogue. Concolic Execution + LLM for Zero-Day Malware Detection (Edwards and Eslamimehr 2026) pairs path-prioritization with concrete execution.

15. Benchmarks and the Evaluation Gap

Benchmarks split cleanly by what they measure.

- Static code quality — HumanEval, MBPP, LiveCodeBench (Jain et al. 2024) measure code-LLM output without exercising the world-model claim.
- Execution reasoning — CRUXEval (A. Gu et al. 2024), REval (J. Chen et al. 2024), CRUXEval-X (Xu et al. 2024), TraceEval (Yikun Li et al. 2026), PLSemanticsBench (Thimmaiah et al. 2025). The canonical WM evals.
- Semantic equivalence — EquiBench (A. Wei, Cao, et al. 2025), CodeARC (A. Wei, Suresh, et al. 2025).
- Agentic — SWE-bench, SWE-Gym, WebArena (Zhou et al. 2023), Mind2Web (Deng et al. 2023), PyBench (Yaolun Zhang et al. 2024), OSWorld, WindowsAgentArena.

The eval gap. No widely adopted benchmark directly measures world-model fidelity by holding the policy fixed and varying the WM. Every measurement of WM capability is mediated through downstream task performance, so we cannot distinguish *the model has internalized program semantics* from *the model is exploiting trace-token shortcuts in the test distribution*. Demystifying Errors in LLM Reasoning Traces (Abdollahi et al. 2025) supplies the diagnostic: even DeepSeek-R1, o4-mini, Gemini 2.5 Flash, and Claude 4, when asked to simulate execution and explain their reasoning, produce traces with

errors clustering into nine categories (Computation, Indexing, Control Flow, Skip Statements, Misvaluation of Native API, Hallucination, Input Misread, etc.). Models with 85–98% final-answer accuracy on output prediction produce traces with systematic errors throughout. The decoupling — high outcome accuracy, low process fidelity — is the signature of a system that has learned to predict outcomes without faithfully simulating dynamics.

16. Empirical Landscape

16.1 SWE-bench scoreboard

Protocols mix in SWE-bench reporting and are non-comparable across the obvious axes. We split the table into three protocol classes and ask readers to compare within, not across, classes. Training / Evaluation regime: *Frozen-model scaffold* = no model training (agent loop around a frozen closed model); *EG-LLM SFT* = supervised fine-tune on agent trajectories; *EG-LLM + RL* = RL with execution rewards; *EG-agent + learned simulator* = SFT+RL plus a learned execution surrogate; *EG-LLM + trace mid-train + RL* = the CWM lineage; *Scaffold-evolved* = recursive scaffold/agent self-improvement around a frozen closed model. WM-Arch is reserved for systems with explicit transition/rollout machinery (N1/N2/N3 from §3); no row in this table qualifies. Tier is the evidence grading from §2 (A peer-reviewed, B major arxiv preprint with code, C single-author/institution, D blog/leaderboard).

Table 16.1a — Single-sample resolution (pass@1 on Verified unless noted)

System	Base model	Bench	Pass@1	Regime	Tier	Source
SWE-agent	GPT-4 Turbo	full	12.5%	Frozen-model scaffold	A	(John Yang et al. 2024)
AutoCodeRover	GPT-4	Lite	19.0%	Frozen-model scaffold	A	(Yuntong Zhang et al. 2024)
Agentless	GPT-4o	Lite	32.0%	Frozen-model scaffold	B	(Xia et al. 2024)
SWE-Gym	Qwen-2.5-Coder-32B	Verified	20.6%	EG-LLM SFT	B	(Pan et al. 2024)
Agent-RLVR (no RM)	Qwen-2.5-72B	Verified	22.4%	EG-LLM + RL	B	(Da et al. 2025)
Long-Context Multi-Turn RL	Qwen-2.5-72B	Verified	39.0%	EG-LLM + RL	B	(Golubev et al. 2025)
SWE-RL (Llama3-SWE-RL-70B)	Llama-3.3-70B	Verified	41.0%	EG-LLM + RL	B	(Y. Wei, Duchenne, et al. 2025)

System	Base model	Bench	Pass@1	Regime	Tier	Source
Nanbeige SWE-World (RL)	Qwen-2.5- Coder-32B	Verified	55.0%	EG-agent + learned simulator	B	(S. Sun et al. 2026)
CWM	CWM-32B	Verified	<i>not directly reported</i> †	EG-LLM + trace mid-train + RL	B	(team et al. 2025)

† The CWM paper does not report a pure pass@1 number under the standard protocol; the headline 65.8% in Table 16.1b is best@16 with verifier reranking. Inferred pass@1 from the paper’s reported figures is approximately 53–55%, but this is *our estimate, not a directly reported number*, and we do not place it in the main row.

Table 16.1b — Best-of-k with verifier reranking / TTS

System	Base	Bench	Score	Sample budget	Tier	Source
Agent- RLVR + RM rerank	Qwen-2.5- 72B	Verified	27.8%	rerank over multi- sample	B	(Da et al. 2025)
Nanbeige SWE-World (TTS@8)	Qwen-2.5- Coder-32B	Verified	68.2%	best@8	B	(S. Sun et al. 2026)
CWM (TTS)	CWM-32B	Verified	65.8%	best@16 over 40 reranked	B	(team et al. 2025)

Table 16.1c — Scaffold-evolution and self-play around closed-model executors

System	Backbone executor	Bench	Score (start [?] end)	Note	Tier	Source
Darwin Godel Machine	Claude-3.5- Sonnet	Verified	20% [?] 50%	Scaffold evolves over 80 iterations	B	(J. Zhang et al. 2025)
Huxley Godel Machine	GPT-5-mini	Verified (500)	— [?] 61.4%	Scaffold optimized for Verified-60	B	(Wenyi Wang et al. 2025)

System	Backbone executor	Bench	Score (start [?] end)	Note	Tier	Source
SICA	Claude-3.5-Sonnet + o3-mini	Verified (subset)	17% [?] 53%	Scaffold evolves; not a model-training method	C	(Robeyns, Szummer, and Aitchison 2025)

Citations that quote “CWM 65.8% SWE-bench Verified” without disclosing the best@16/TTS protocol should be read as misreporting, not as direct pass@1 numbers.

What can and cannot be compared. Within Table 16.1a the comparable axis is pass@1. SWE-RL’s 41.0%, Long-Context-MT-RL’s 39.0%, and Nanbeige SWE-World’s 55.0% are roughly comparable as “open-weight WM-trained pass@1 on Verified.” Table 16.1c sits in a different category — these are scaffold-evolution methods over closed models, and treating them as evidence for “world-model training works” conflates scaffold search with model improvement. Reading them inside Table 16.1a, as some online discussion does, is apples-to-oranges.

Two empirical claims, separately defensible. (i) On *pass@1*, open-weight WM-trained 32B systems (Nanbeige 55.0%, Long-Context-MT-RL 39.0%, SWE-RL 41.0%) outperform frozen open-weight baselines (Qwen-2.5-Coder-32B raw 6.2%) by 30–50 absolute points; this gain is the strongest case for training on agent trajectories with execution feedback. (ii) On *best@k with verifier reranking* (Table 16.1b), the same models reach 65–68%, but the additional 13–15 points are attributable to TTS reranking infrastructure, not to the WM training. Conflating (i) and (ii) by quoting CWM’s 65.8% alongside SWE-RL’s 41.0% as both “world-model-trained pass@1 SOTA” is exactly the misreporting the protocol split is designed to prevent.

16.2 Trace pretraining gains on execution-reasoning

Paper	Backbone	Baseline	After trace pretrain/FT	Delta	Benchmark	Tier
TRACED	UnixCoder	—	+12.4% rel branch-coverage; +25.2% rel variable-value	rel	CodeNet exec	A
NExT	PaLM 2-L	23.2	49.3	+26.1 abs	MBPP-R	A
NExT	PaLM 2-L	32.2	42.5	+10.3 abs	HumanEvalFixA Plus	A
SemCoder (1.3B)	DS-Coder 1.3B	base	63.6 / 63.9	+23 abs	CRUXEval-I / O	B
“What I cannot execute”	Llama-3.1-8B	37.8%	~80%	+42 abs	CRUXEval-O	B

Paper	Backbone	Baseline	After trace pretrain/FT	Delta	Benchmark	Tier
Do Code Semantics Help?	DSCoder, Llama-3, Gemma-2	various	⌈ a few abs; some regressions	mixed	comprehensive	B

For under-trained ⌈8B open-weights, trace pretraining delivers +15 to +42 absolute on CRUXEval-O. The “Do Code Semantics Help?” ablation (J. Wang et al. 2025) is the disconfirming evidence: across multiple backbones and five trace representations, no single representation consistently outperforms others, and several downstream tasks regress under trace augmentation. The gain shrinks rapidly with base-model quality. Trace pretraining is a remedial intervention for weak code models; whether frontier models still benefit is unsettled.

16.3 Web/OS agents from WMs

Paper	Benchmark	Base	With WM	Delta	Mechanism	Tier
WebDreamer (GPT-4o WM)	VisualWebArena	17.6%	23.6%	+34.1% rel (⌈+6 abs)	LLM-as-WM + MPC	B
WebDreamer	Online-Mind2Web	26.0%	37.0%	+42.3% rel (⌈+11 abs)	same	B
Dreamer-7B (trained WM)	VisualWebArena	base	+4.7 abs	—	trained WM	B
WMA (Chae et al. 2024)	WebArena	base	action-selection 52⌈70%	—	trained transition WM	B
Dyna-Think DDT (32B)	OSWorld BoN	RFT~28%	43.1%	⌈+15 abs	Dyna-Q + WM head	B
Dyna-Think DDT	WindowsAgentArena	28.4%	34.9%	+6.5 abs	same	B

WM gains on web/OS are real but quantitatively small (⌈+5–10 absolute task success rate on most benchmarks), and partly confounded with the extra synthetic data the WM generates. DyMo’s 90%+ state-prediction accuracy versus 72.8% task success rate exemplifies the decoupling: WM heads can be accurate without the agent being accurate.

16.4 Formal verification vs LLM-only

System	Language	Benchmark	Baseline	With system	Tier	Source
ATLAS	Dafny	DafnyBench Pass@1	32.4%	55.8%	B	(Baksys et al. 2025)
ATLAS	Dafny	DafnySynthesis Pass@5	15.8%	65.8% (>GPT-4 53.4)	B	(Baksys et al. 2025)
CLEVER	Lean 4	161 problems end-to-end	best frontier	21/161	B	(Thakur et al. 2025)
VeriStruct	Verus / Rust	11 modules	—	99.2% (128/129 fns)	C	(C. Sun et al. 2025)
AutoRocq	Rocq	math + verif lemmas	5 baselines	48.0% math / 30.9% verif	B	(Tu et al. 2025)
Semantic Equiv Self-Play	Liquid Haskell	EquiBench	base	+13.3 pp	B	(Barone and Nok 2026)

Verified codegen has the steepest training-data sensitivity in the survey: small synthetic datasets (2.7K verified Dafny programs in ATLAS) produce +25–50 absolute gains because LLM-only baselines start near zero. CLEVER’s 21/161 shows that without explicit data/scaffolding, frontier models cannot reliably produce verified code. VeriStruct shows that on curated targets near-perfect is reachable.

16.5 Reasoning-model competitive programming

Model	Codeforces	IOI / ICPC	Tier
gpt-4o	808 (11th pct)	—	D
o1-preview	1258 (62nd pct)	—	D
o1	1673 (89th pct)	—	D
o1-ioi	1807 (~93rd pct)	IOI 2024 49th pct live	B
o3	elite-human-class	IOI 2024 gold	D
gpt-oss-120b + GenCluster	—	IOI 2025 gold (open-weight)	B
Gemini 2.5 Pro Exp	—	ICPC-Eval Pass@1 22.0%	B
DeepSeek-R1	—	ICPC-Eval Pass@1 14.4%	B
Claude 3.7 Sonnet	—	ICPC-Eval Pass@1 11.8%	B
GPT-4o	—	ICPC-Eval Pass@1 5.9%	B

Codeforces +999 rating points in 14 months on the o-series. The same line shows that frontier reasoning models gain almost everything from RL scaling, not from domain-specialized scaffolds — which complicates the case that “trace-style world models” are doing causal work for frontier systems.

17. Critical Perspectives

This section names where the field overclaims, where the consensus is fragile, and where vocabulary is doing more work than evidence. We develop seven theses.

17.1 The “world model” label has become marketing for any code LLM trained on something other than raw source

This is a terminology argument, not a contradiction with §3. §3 deliberately admits “implicit WMs in token policies” as a fourth flavor — including CWM, TRACED, SemCoder under that flavor — because that is how the field currently uses the term. The claim of §17.1 is that this permissive definition is what allows the rhetorical sleight of hand, and that a stricter definition would be more useful going forward.

Concretely: read CWM (team et al. 2025) carefully and the architecture it ships is a 32B decoder-only Transformer with GQA, sliding-window blocks, RoPE, AdamW — a Llama-class model. What earns it the “world model” badge is the mid-training datamix: 5T tokens of Python observation-action traces plus ForagerAgent SWE trajectories. There is no separate dynamics head, no inverse model, no recurrent latent. The same observation lands on LLM-JEPA, DyMo, and most “world model” papers from 2024–2026: the artifact is a standard LLM with an enriched objective.

A useful purity test: *can we ablate the supposed world-modeling component without changing the architecture?* If yes, the system is a trace-trained LLM. If no, there is a genuine architectural commitment. By that test, CoLA (Jia et al. 2025) and the Dreamer-for-LLMs gestures pass. CWM, SemCoder, NExT, and most of the “explicit WM” cluster fail. We propose that “world model” be reserved for systems with the architectural commitment, and that *execution-grounded code LLM* serve for the rest. The §3 permissive definition can stay as descriptive — what the field currently calls a code WM — but a normative tightening is warranted, and the survey from §17 onward uses the stricter sense.

17.2 Trace pretraining has a causal-isolation problem the surface numbers obscure

“Do Code Semantics Help?” (J. Wang et al. 2025) is the most damaging paper for the prevailing optimism. It runs a comprehensive ablation across DeepSeek-Coder, LLaMA-3, and Gemma-2 with five representations (Scratchpad, NExT, CodeExecutor, Concise, SemCoder) on program repair, code synthesis, BigCodeBench, LiveCodeBench, and CRUXEval. Its headline: integrating trace-based semantic information into SFT *cannot significantly enhance* code-generation ability. In 7 of 9 synthesis settings the no-trace baseline wins or ties. At inference, in 36 of 56 test-scaling configurations, trace prompts hurt.

“What I cannot execute, I do not understand” (Armengol-Estapé et al. 2025) is gentler — Execution Tuning reaches ~80% CRUXEval-O — but the same paper’s downstream evaluations on HumanEval, MBPP, and GSM8K show *negligible* gains from trace data in the SFT mix.

The strongest counterexample to “barely transfers” is NExT (Ni et al. 2024), which pushes PaLM 2-L on Mbbpp-R from 23.2% to 40.8% *with traces removed at test time* — a +17.6 absolute gain on the no-trace-at-inference setting, which is exactly the transfer pattern the skeptical reading says shouldn’t happen. The honest version of the thesis is therefore narrower: trace pretraining transfers to *program-repair* tasks where the model needs to reason about a buggy program’s behavior (where NExT and TraceFixer-style work shines), and barely transfers to *fresh code synthesis* on benchmarks like HumanEval and MBPP that are closer to the model’s prior. The “Do Code Semantics Help?” disconfirmation is on synthesis; NExT’s transfer is on repair. Both can be true.

The honest reading: trace pretraining helps execution prediction (the thing trained on), transfers to runtime-reasoning-heavy downstream tasks like repair, and barely transfers to fresh code synthesis, exactly the pattern you would expect if dense execution supervision is teaching a runtime-tracking skill rather than a generative model of program intent.

CWM’s 65.8% on SWE-bench is the apparent counterexample, but the Meta team reports it after trace mid-training *and* 3M ForagerAgent SWE trajectories *and* multi-task RL with verifiable rewards — three interventions stacked. The “world model” component is not causally isolated from the SWE-trajectory and RL components. Without an ablation that removes trace mid-training while holding ForagerAgent and RL fixed, the headline is unfalsifiable. The field has agreed to call CWM a world-model success because the name is on the model card, not because the experimental design demonstrates it.

17.3 The Dreamer-for-code gap may be a non-problem

The conventional framing treats the absence of latent-imagination world models for SWE as the field’s largest architectural gap. The empirical record argues the opposite. CWM in token space reaches 65.8% SWE-bench. CoLA produces respectable but not field-shifting results, and even there the WM is fine-tuned on top of a standard LLM rather than replacing it.

Vision world models needed latent rollouts because pixel-space rollouts were too expensive — a frame is $\sim 10^6$ dimensions, dynamics are partially observed. Program execution is the opposite: a Python frame is small, dynamics are observable, and the simulator (CPython) is available for free at training time. The pressure that drove Dreamer’s RSSM design does not exist for code.

Debugging Code World Models (Rahmani 2026) shows CWM’s long-horizon failures are dominated by *action hallucination*, not state-propagation error — under teacher forcing CWM tracks state correctly for 128 steps. A latent rollout would compress states but not fix the action policy, which is the actual bottleneck.

The counterargument is fair: latent rollouts permit faster planning at inference, and for multi-agent or population-scale search the speedup is asymptotically meaningful. But the survey should retire “single largest architectural gap” framing and replace it with “an interesting open question whose payoff is not yet demonstrated.”

17.4 PRMs are critics, not world models

Process reward models — ExecVerify (Tang et al. 2026), SWE-PRM (Gandhi et al. 2025), ThinkPRM (Khalifa et al. 2025), FunPRM (R. Zhang et al. 2026), DataPRM (Qiu et al. 2026) — are often grouped under the world-modeling umbrella. This is wrong in a way that matters. A world model is, by every definition the survey uses, a *forward* predictor of $(\text{state}, \text{action}) \rightarrow \text{next_state}$. A PRM is a *backward-looking evaluator*: given a partial trajectory, score it. In classical model-based RL these are different objects — Dreamer has both a world model (RSSM) and a critic (value function).

PRMs cannot roll out. Cannot simulate counterfactuals. Cannot be used by a planner that wants to score a hypothesized future. Conflating them dilutes the world-model concept until it means “any neural network trained on execution-related signals” — at which point the term is useless. Vocabulary discipline is cheap and the field would benefit from it.

The same critique applies, less severely, to verifier-grounded systems: a Lean proof checker is not a world model, it is a deterministic verifier of a candidate output. Calling it “the world model” when

ATLAS, Re:Form, or AutoRocq use it makes for a tidy survey arc but blurs the actual computational structure.

17.5 “General Agents Contain World Models” is much weaker than its title suggests

Richens et al. (Richens et al. 2025) prove that any goal-conditioned policy satisfying a regret bound δ for sufficiently deep composite goals (depth $n \geq 1$) must encode an extractable approximation of the transition function with bounded error. Genuine and elegant.

But read the assumptions: fully observed environment, finite communicating stationary controlled MDP, goal-conditioned policy satisfying a regret bound for a specific class of LTL composite goals of depth n . Theorem 2 of the same paper explicitly shows that for myopic agents (depth-1 goals), *no world model is needed*. Real SWE agents are myopic-ish over short turns and approximately competent over longer ones; their environments are partially observed (rarely full filesystem state); they violate stationarity (the repository changes under their actions); and their regret bound for arbitrary composite goals is unknown and almost certainly not satisfied. The authors caveat this in §6 (“Limitations”) of their paper.

The theorem is a beautiful existence proof for an idealized agent class. It is *not* an empirical statement that SWE coding agents have learned world models, and it provides no guidance about the fidelity of any world model they may have learned.

17.6 The verifier-grounded lineage is the actual leading edge

ATLAS, Re:Form, CLEVER, VeriStruct, AutoRocq, and the Liquid Haskell self-play paper (Barone and Nok 2026) share a property no LLM-only system possesses: code whose correctness is *machine-checked* against a formal specification. Compare to the SWE-bench paradigm, where “correctness” means “hidden unit tests pass” — a weaker guarantee, since unit tests cover specific inputs and the system can pass them while being wrong on adjacent inputs.

The abstract-interpreter paper (Mitchell et al. 2025) is the diagnostic: when frontier reasoning LLMs are asked to reason in the style of formal abstract interpretation over 22 SV-COMP programs, they make systematic errors in widening, fixpoint termination, control-flow propagation, and meet/join operations. They generated unsound invariants on programs as small as `count_by_2.c`. If LLMs cannot reliably perform interval-domain abstract interpretation on toy C programs, claims that they have learned faithful internal world models of program semantics are doing a lot of inferential work.

The verifier-grounded line is the only research direction that does not rely on LLM self-report for correctness. Scoped carefully, it leads on *synthesis-from-spec* (ATLAS reaches 65.8% Pass@5 on DafnySynthesis at 7B, surpassing GPT-4) and on *near-perfect verified completion of curated targets* (VeriStruct 99.2% on 11 Verus modules). It does *not* lead on end-to-end verified codegen from natural language: CLEVER reports 1/161 on Lean problems requiring joint spec + implementation verification. The future of *correct* code is plausibly hybrid — neural proposal, symbolic verification, with the verifier providing ground truth that learned world models cannot — but the hybrid is not yet a finished pillar. We catalog the verifier line in §14.1 alongside symbolic execution and abstract-interpretation probing; readers who accept this thesis should weight that subsection more heavily than its sequence position suggests.

17.7 The evaluation gap is the structural reason the field looks confused

Across CRUXEval, REval, CRUXEval-X, PLSemanticsBench, TraceEval, and EquiBench, no benchmark holds policy fixed and varies world-model quality. Every measurement of “world-modeling capability” is mediated through downstream task performance, so we cannot distinguish *the model has internalized program semantics* from *the model is exploiting trace-token shortcuts in the test distribution*.

“Demystifying Errors in LLM Reasoning Traces” (Abdollahi et al. 2025) is the diagnostic: even DeepSeek-R1, o4-mini, Gemini 2.5 Flash, and Claude 4, when asked to simulate execution, produce traces with errors in nine systematic categories. Models with 85–98% final-answer accuracy on output prediction produce traces with systematic errors throughout — high outcome accuracy, low process fidelity, the signature of a system that predicts outcomes without faithfully simulating dynamics.

Self-repair literature exhibits the same pathology: Olausson et al. (Olausson et al. 2023) showed that GPT-4 self-repair on APPS and HumanEval, normalized by compute, often performs *worse* than i.i.d. resampling. The bottleneck is the model’s feedback quality, not its repair capability — human-written feedback boosts repair success by 1.58×. The model can generate code, can sometimes recognize bugs, but cannot reliably simulate why its code is wrong — which is exactly what a faithful world model would let it do. The empirical bound on LLM self-repair is, in effect, an empirical bound on the fidelity of the implicit world model the LLM is running. Calling that world model internal is fine; calling it good is not.

Until benchmarks measure process fidelity independently of outcome, “is this system actually building a world model?” remains scientifically undecidable.

18. Open Problems

The critical perspectives of §17 reshape the conventional open-problems list. We propose six problems where the literature is thinnest *and* the upside is largest.

1. Causal isolation of trace-pretraining contributions. Every claim of the form “this WM-trained model achieves X” should be paired with an ablation removing the WM component while holding training data and RL fixed. CWM in particular needs this. Without it, the headline numbers underdetermine whether the WM did the work.
2. World-model fidelity as a first-class metric. §15’s eval gap is concrete: build a benchmark where holding policy fixed and varying WM quality causes measurable variation in planning quality, independent of downstream task. This benchmark would clarify the field more than any single new model.
3. Hybrid neural-symbolic systems. §17.6 argues the verifier-grounded line is the leading edge. The natural integration is *neural proposal, symbolic verification*, with the verifier providing gradient-free correctness signal and the neural component providing proposals at scale. Differentiable surrogates of symbolic verifiers (Lean / Dafny / Rocq) that pass verifier-style gradients during training are open.
4. Multi-modal WMs for coding. GUI agents need pixel-level WMs (Neural Computers, (Zhuge et al. 2026), is a first attempt). Tying pixel WMs to code-state WMs through a shared latent is essentially unsolved.
5. Long-horizon credit assignment with execution-grounded rewards. PRMs (§9.3) are early, and §17.4 argues they should be conceptually separated from world models. The right structure for rewarding an agent across hundreds of execution-grounded steps is a live question.

6. World models of the developer, not just the program. All current WMs model the *machine*. Few model the *developer intent* with comparable fidelity. ATLAS and Re:Form gesture in this direction by treating the spec as the WM. A full developer-intent WM would close the agentic loop.

We do not list “Dreamer-for-SWE-agents” as the field’s largest gap, contrary to common framing. §17.3 argues the pressure motivating that direction in vision does not transfer to code. It remains an interesting research question, not the highest-leverage one.

Three under-explored representations. The §5.3 taxonomy table flags three classes of representation, mature in vision-WM, that have no code-WM exemplar:

- *Global Latent Vector (Dreamer-style RSSM)* — discussed and contested above. The Hafner et al. DreamerV1–V3 line proves the design space exists; whether it pays off for code is the question §17.3 leaves open.
- *Spatial / Structural Grid* — the analog of OccWorld / BEV for code would be a learned predictive grid over AST nodes, call-graph edges, or CFG states. RepoGraph (Ouyang et al. 2024) shows the static version is useful as agent state; the predictive version is unexplored.
- *Decomposed Object / Slot (object-centric WMs)* — the analog for code would model variables, scopes, or classes as discrete persistent slots whose state propagates independently. No paper in the corpus instantiates this, despite obvious mappings (each variable is an object, each frame is a scene).

The object-centric and structural-grid gaps look more genuine than the Dreamer one, in our reading, because they exploit structure that code *already has* (objects = variables, grids = AST/CFG) rather than borrowing pressure from a domain (vision) where the structural assumptions differ.

19. Conclusion

Across the literature surveyed here, a single trajectory is visible: from neural execution (modeling the machine), through trace pretraining (modeling execution implicitly), to CWM and its descendants (modeling execution explicitly with a named artifact), to agentic SWE and RL (modeling the environment), to JEPA and latent-action models (modeling in compressed space), and on toward formal verification, probing, and safety (modeling reliably). What was a scattered set of insights in 2014 has by 2026 cohered into a recognizable research program with a recognizable artifact — the code world model.

The remaining work splits into two halves. The first is empirical: close the eval gap, isolate the causal contribution of WM-training, build hybrid neural-symbolic systems whose correctness is verifier-checkable rather than test-checkable. The second is rhetorical: hold the term “world model” to a strict definition so the literature can distinguish architectural commitments from training-data choices, and resist the temptation to oversell extractability theorems and latent-imagination analogies whose premises do not transfer to code.

The opportunity is large precisely because the framework is now clear enough to identify what is missing. The work to do is the work this survey has tried to make visible.

Appendix · Glossary

- CWM — Code World Model ((team et al. 2025) and lineage).
 - JEPA — Joint Embedding Predictive Architecture (LeCun et al.).
 - RSSM — Recurrent State-Space Model (Dreamer family).
 - PRM — Process Reward Model.
 - SWE agent — Software-engineering agent operating on real repositories.
 - Trace pretraining — Pretraining where execution traces appear in input or target.
 - Execution-grounded RL — RL whose reward derives from program execution.
 - Latent-imagination rollout — Forward simulation in compressed latent space rather than token space.
 - TTS — Test-time scaling (sampling many candidates + verifier reranking at inference).
 - GRPO — Group Relative Policy Optimization (R1-family RL algorithm).
-

References

- Abdollahi, Mohammad, Khandaker Rifah Tasnia, Soumit Kanti Saha, Jinqiu Yang, Song Wang, and Hadi Hemmati. 2025. “Demystifying Errors in LLM Reasoning Traces: An Empirical Study of Code Execution Simulation.” <https://arxiv.org/abs/2512.00215>.
- Armengol-Estapé, Jordi, Quentin Carbonneaux, Tianjun Zhang, Aram H. Markosyan, Volker Seeker, Chris Cummins, Melanie Kambadur, et al. 2025. “What i Cannot Execute, i Do Not Understand: Training and Evaluating LLMs on Program Execution Traces.” <https://arxiv.org/abs/2503.05703>.
- Assran, Mahmoud, Quentin Duval, Ishan Misra, Piotr Bojanowski, Pascal Vincent, Michael Rabat, Yann LeCun, and Nicolas Ballas. 2023. “Self-Supervised Learning from Images with a Joint-Embedding Predictive Architecture.” <https://arxiv.org/abs/2301.08243>.
- Austin, Jacob, Augustus Odena, Maxwell Nye, Maarten Bosma, Henryk Michalewski, David Dohan, Ellen Jiang, et al. 2021. “Program Synthesis with Large Language Models.” <https://arxiv.org/abs/2108.07732>.
- Baksys, Mantas, Stefan Zetsche, Olivier Bouissou, Remi Delmas, Soonho Kong, and Sean B. Holden. 2025. “ATLAS: Automated Toolkit for Large-Scale Verified Code Synthesis.” <https://arxiv.org/abs/2512.10173>.
- Barone, Antonio Valerio Miceli, and Poon Tsz Nok. 2026. “Improving LLM Code Reasoning via Semantic Equivalence Self-Play with Formal Verification.” <https://arxiv.org/abs/2604.17010>.
- Beck, Maximilian, Jonas Gehring, Jannik Kossen, and Gabriel Synnaeve. 2026. “Towards a Neural Debugger for Python.” <https://arxiv.org/abs/2603.09951>.
- Bharti, Varun, Shashwat Jha, Dhruv Kumar, and Pankaj Jalote. 2025. “Loop Invariant Generation: A Hybrid Framework of Reasoning Optimised LLMs and SMT Solvers.” <https://arxiv.org/abs/2508.00419>.
- Bieber, David, Charles Sutton, Hugo Larochelle, and Daniel Tarlow. 2020. “Learning to Execute Programs with Instruction Pointer Attention Graph Neural Networks.” <https://arxiv.org/abs/2010.12621>.

- Bruce, Jake, Michael Dennis, Ashley Edwards, Jack Parker-Holder, Yuge Shi, Edward Hughes, Matthew Lai, et al. 2024. "Genie: Generative Interactive Environments." <https://arxiv.org/abs/2402.15391>.
- Chae, Hyungjoo, Namyoun Kim, Kai Tzu-iunn Ong, Minju Gwak, Gwanwoo Song, Jihoon Kim, Sunghwan Kim, Dongha Lee, and Jinyoung Yeo. 2024. "Web Agents with World Models: Learning and Leveraging Environment Dynamics in Web Navigation." <https://arxiv.org/abs/2410.13232>.
- Chen, Junkai, Zhiyuan Pan, Xing Hu, Zhenhao Li, Ge Li, and Xin Xia. 2024. "Reasoning Runtime Behavior of a Program with LLM: How Far Are We?" <https://arxiv.org/abs/2403.16437>.
- Chen, Mark, Jerry Tworek, Heewoo Jun, Qiming Yuan, Henrique Ponde de Oliveira Pinto, Jared Kaplan, Harri Edwards, et al. 2021. "Evaluating Large Language Models Trained on Code." <https://arxiv.org/abs/2107.03374>.
- Chen, Xinyun, Maxwell Lin, Nathanael Schärli, and Denny Zhou. 2023. "Teaching Large Language Models to Self-Debug." <https://arxiv.org/abs/2304.05128>.
- Chen, Yongchao, Yueying Liu, Junwei Zhou, Yilun Hao, Jingquan Wang, Yang Zhang, Na Li, and Chuchu Fan. 2025. "R1-Code-Interpreter: LLMs Reason with Code via Supervised and Multi-Stage Reinforcement Learning." <https://arxiv.org/abs/2505.21668>.
- Chen, Zhirong, Kaiyan Chang, Zhuolin Li, Cangyuan Li, Xinyang He, Chujie Chen, Mengdi Wang, et al. 2025. "ChipSeek: Optimizing Verilog Generation via EDA-Integrated Reinforcement Learning." <https://arxiv.org/abs/2507.04736>.
- Da, Jeff, Clinton Wang, Xiang Deng, Yuntao Ma, Nikhil Barhate, and Sean Hendryx. 2025. "Agent-RLVR: Training Software Engineering Agents via Guidance and Environment Rewards." <https://arxiv.org/abs/2506.11425>.
- Dainese, Nicola, Matteo Merler, Minttu Alakuijala, and Pekka Marttinen. 2024. "Generating Code World Models with Large Language Models Guided by Monte Carlo Tree Search." <https://arxiv.org/abs/2405.15383>.
- DeepSeek-AI, Daya Guo, Dejian Yang, Haowei Zhang, Junxiao Song, Peiyi Wang, Qihao Zhu, et al. 2025. "DeepSeek-R1: Incentivizing Reasoning Capability in LLMs via Reinforcement Learning." <https://arxiv.org/abs/2501.12948>.
- Deng, Xiang, Yu Gu, Boyuan Zheng, Shijie Chen, Samuel Stevens, Boshi Wang, Huan Sun, and Yu Su. 2023. "Mind2Web: Towards a Generalist Agent for the Web." <https://arxiv.org/abs/2306.06070>.
- Ding, Jingtao, Yunke Zhang, Yu Shang, Jie Feng, Yuheng Zhang, Zefang Zong, Yuan Yuan, et al. 2024. "Understanding World or Predicting Future? A Comprehensive Survey of World Models." <https://arxiv.org/abs/2411.14499>.
- Ding, Yangruibo, Jinjun Peng, Marcus J. Min, Gail Kaiser, Junfeng Yang, and Baishakhi Ray. 2024. "SemCoder: Training Code Language Models with Comprehensive Semantics Reasoning." <https://arxiv.org/abs/2406.01006>.
- Ding, Yangruibo, Ben Steenhoeck, Kexin Pei, Gail Kaiser, Wei Le, and Baishakhi Ray. 2023. "TRACED: Execution-Aware Pre-Training for Source Code." <https://arxiv.org/abs/2306.07487>.
- Edwards, George, and Mahdi Eslamimehr. 2026. "Synergistic Directed Execution and LLM-Driven Analysis for Zero-Day AI-Generated Malware Detection." <https://arxiv.org/abs/2603.09044>.

- Feng, Xidong, Ziyu Wan, Muning Wen, Stephen Marcus McAleer, Ying Wen, Weinan Zhang, and Jun Wang. 2023. “Alphazero-Like Tree-Search Can Guide Large Language Model Decoding and Training.” <https://arxiv.org/abs/2309.17179>.
- Gandhi, Shubham, Jason Tsay, Jatin Ganhotra, Kiran Kate, and Yara Rizk. 2025. “When Agents Go Astray: Course-Correcting SWE Agents with PRMs.” <https://arxiv.org/abs/2509.02360>.
- Gehring, Jonas, Kunhao Zheng, Jade Copet, Vegard Mella, Quentin Carbonneaux, Taco Cohen, and Gabriel Synnaeve. 2024. “RLEF: Grounding Code LLMs in Execution Feedback with Reinforcement Learning.” <https://arxiv.org/abs/2410.02089>.
- Golubev, Alexander, Maria Trofimova, Sergei Polezhaev, Ibragim Badertdinov, Maksim Nekrashevich, Anton Shevtsov, Simon Karasik, et al. 2025. “Training Long-Context, Multi-Turn Software Engineering Agents with Reinforcement Learning.” <https://arxiv.org/abs/2508.03501>.
- Gong, Shansan, Ruixiang Zhang, Huangjie Zheng, Jiatao Gu, Navdeep Jaitly, Lingpeng Kong, and Yizhe Zhang. 2025. “DiffuCoder: Understanding and Improving Masked Diffusion Models for Code Generation.” <https://arxiv.org/abs/2506.20639>.
- Gu, Alex, Baptiste Rozière, Hugh Leather, Armando Solar-Lezama, Gabriel Synnaeve, and Sida I. Wang. 2024. “CRUXEval: A Benchmark for Code Reasoning, Understanding and Execution.” <https://arxiv.org/abs/2401.03065>.
- Gu, Yu, Kai Zhang, Yuting Ning, Boyuan Zheng, Boyu Gou, Tianci Xue, Cheng Chang, et al. 2024. “Is Your LLM Secretly a World Model of the Internet? Model-Based Planning for Web Agents.” <https://arxiv.org/abs/2411.06559>.
- Guo, Shangmin, Omar Darwiche Domingues, Raphaël Avalos, Aaron Courville, and Florian Strub. 2025. “World Modelling Improves Language Model Agents.” <https://arxiv.org/abs/2506.02918>.
- Ha, David, and Jürgen Schmidhuber. 2018. “World Models.” <https://arxiv.org/abs/1803.10122>.
- Hafner, Danijar, Timothy Lillicrap, Jimmy Ba, and Mohammad Norouzi. 2019. “Dream to Control: Learning Behaviors by Latent Imagination.” <https://arxiv.org/abs/1912.01603>.
- Hafner, Danijar, Timothy Lillicrap, Mohammad Norouzi, and Jimmy Ba. 2020. “Mastering Atari with Discrete World Models.” <https://arxiv.org/abs/2010.02193>.
- Hafner, Danijar, Jurgis Pasukonis, Jimmy Ba, and Timothy Lillicrap. 2023. “Mastering Diverse Domains Through World Models.” <https://arxiv.org/abs/2301.04104>.
- Han, Hao, Jin Xie, Xuehao Ma, Weiquan Zhu, Ziyao Zhang, ZhiLiang Long, Hongkai Chen, and Qingwen Ye. 2026. “SWE-TRACE: Optimizing Long-Horizon SWE Agents Through Rubric Process Reward Models and Heuristic Test-Time Scaling.” <https://arxiv.org/abs/2604.14820>.
- Hao, Shibo, Yi Gu, Haodi Ma, Joshua Jiahua Hong, Zhen Wang, Daisy Zhe Wang, and Zhiting Hu. 2023. “Reasoning with Language Model Is Planning with World Model.” <https://arxiv.org/abs/2305.14992>.
- Huang, Hai, Yann LeCun, and Randall Balestriero. 2025. “LLM-JEPA: Large Language Models Meet Joint Embedding Predictive Architectures.” <https://arxiv.org/abs/2509.14252>.
- Jain, Naman, King Han, Alex Gu, Wen-Ding Li, Fanjia Yan, Tianjun Zhang, Sida Wang, Armando Solar-Lezama, Koushik Sen, and Ion Stoica. 2024. “LiveCodeBench: Holistic and Contamination Free Evaluation of Large Language Models for Code.” <https://arxiv.org/abs/2403.07974>.

- Jia, Chengxing, Ziniu Li, Pengyuan Wang, Yi-Chen Li, Zhenyu Hou, Yuxiao Dong, and Yang Yu. 2025. “Controlling Large Language Model with Latent Actions.” <https://arxiv.org/abs/2503.21383>.
- Jiang, Juyong, Fan Wang, Jiasi Shen, Sungju Kim, and Sunghun Kim. 2024. “A Survey on Large Language Models for Code Generation.” <https://arxiv.org/abs/2406.00515>.
- Jimenez, Carlos E., John Yang, Alexander Wettig, Shunyu Yao, Kexin Pei, Ofir Press, and Karthik Narasimhan. 2023. “SWE-Bench: Can Language Models Resolve Real-World GitHub Issues?” <https://arxiv.org/abs/2310.06770>.
- Jung, Dongwon, Wenxuan Zhou, and Muhao Chen. 2025. “Code Execution as Grounded Supervision for LLM Reasoning.” <https://arxiv.org/abs/2506.10343>.
- Kenneweg, Tristan, Philip Kenneweg, and Barbara Hammer. 2025. “JEPA for RL: Investigating Joint-Embedding Predictive Architectures for Reinforcement Learning.” <https://arxiv.org/abs/2504.16591>.
- Khalifa, Muhammad, Rishabh Agarwal, Lajanugen Logeswaran, Jaekyeom Kim, Hao Peng, Moontae Lee, Honglak Lee, and Lu Wang. 2025. “Process Reward Models That Think.” <https://arxiv.org/abs/2504.16828>.
- Koh, Jing Yu, Stephen McAleer, Daniel Fried, and Ruslan Salakhutdinov. 2024. “Tree Search for Language Model Agents.” <https://arxiv.org/abs/2407.01476>.
- Lavon, Boaz, Shahar Katz, and Lior Wolf. 2025. “Execution Guided Line-by-Line Code Generation.” <https://arxiv.org/abs/2506.10948>.
- Li, Shuocheng, Yihao Liu, Silin Du, Wenxuan Zeng, Zhe Xu, Mengyu Zhou, Yeye He, Haoyu Dong, Shi Han, and Dongmei Zhang. 2025. “Jupiter: Enhancing LLM Data Analysis Capabilities via Notebook and Inference-Time Value-Guided Search.” <https://arxiv.org/abs/2509.09245>.
- Li, Xinqing, Xin He, Le Zhang, Min Wu, Xiaoli Li, and Yun Liu. 2025. “A Comprehensive Survey on World Models for Embodied AI.” <https://arxiv.org/abs/2510.16732>.
- Li, Yihe, Ruijie Meng, and Gregory J. Duck. 2025. “Large Language Model Powered Symbolic Execution.” <https://arxiv.org/abs/2505.13452>.
- Li, Yikun, Jinfeng Jiang, Ting Zhang, Chengran Yang, Chenxing Zhong, Yin Yide, Leow Wen Bin, Eng Lieh Ouh, Lwin Khin Shar, and David Lo. 2026. “An Execution-Verified Multi-Language Benchmark for Code Semantic Reasoning.” <https://arxiv.org/abs/2605.11006>.
- Liu, Anthony Z., Xinhe Wang, Jacob Sansom, Yao Fu, Jongwook Choi, Sungryull Sohn, Jaekyeom Kim, and Honglak Lee. 2024. “Interactive and Expressive Code-Augmented Planning with Large Language Models.” <https://arxiv.org/abs/2411.13826>.
- Liu, Chenxiao, Shuai Lu, Weizhu Chen, Daxin Jiang, Alexey Svyatkovskiy, Shengyu Fu, Neel Sundaresan, and Nan Duan. 2023. “Code Execution with Pre-Trained Language Models.” <https://arxiv.org/abs/2305.05383>.
- Liu, Zhengyao, Yunlong Ma, Jingxuan Xu, Junchen Ai, Xiang Gao, Hailong Sun, and Abhik Roychoudhury. 2025. “Agent That Debugs: Dynamic State-Guided Vulnerability Repair.” <https://arxiv.org/abs/2504.07634>.
- Maimon, Gallil, Ori Yoran, Felix Kreuk, Michael Hassid, Gal Cohen, Pierre Chambon, and Yossi Adi. 2026. “Self-Execution Simulation Improves Coding Models.” <https://arxiv.org/abs/2604.03253>.

- Mitchell, Jacqueline L., Brian Hyeongseok Kim, Chenyu Zhou, and Chao Wang. 2025. “Understanding Formal Reasoning Failures in LLMs as Abstract Interpreters.” <https://arxiv.org/abs/2503.12686>.
- Monperrus, Martin. 2026. “Bootstrapping Coding Agents: The Specification Is the Program.” <https://arxiv.org/abs/2603.17399>.
- Ni, Ansong, Miltiadis Allamanis, Arman Cohan, Yinlin Deng, Kensen Shi, Charles Sutton, and Pengcheng Yin. 2024. “NExT: Teaching Large Language Models to Reason about Code Execution.” <https://arxiv.org/abs/2404.14662>.
- Ni, Ansong, Srini Iyer, Dragomir Radev, Ves Stoyanov, Wen-tau Yih, Sida I. Wang, and Xi Victoria Lin. 2023. “LEVER: Learning to Verify Language-to-Code Generation with Execution.” <https://arxiv.org/abs/2302.08468>.
- Nye, Maxwell, Anders Johan Andreassen, Guy Gur-Ari, Henryk Michalewski, Jacob Austin, David Bieber, David Dohan, et al. 2021. “Show Your Work: Scratchpads for Intermediate Computation with Language Models.” <https://arxiv.org/abs/2112.00114>.
- Olausson, Theo X., Jeevana Priya Inala, Chenglong Wang, Jianfeng Gao, and Armando Solar-Lezama. 2023. “Is Self-Repair a Silver Bullet for Code Generation?” <https://arxiv.org/abs/2306.09896>.
- Ouyang, Siru, Wenhao Yu, Kaixin Ma, Zilin Xiao, Zhihan Zhang, Mengzhao Jia, Jiawei Han, Hongming Zhang, and Dong Yu. 2024. “RepoGraph: Enhancing AI Software Engineering with Repository-Level Code Graph.” <https://arxiv.org/abs/2410.14684>.
- Pan, Jiayi, Xingyao Wang, Graham Neubig, Navdeep Jaitly, Heng Ji, Alane Suhr, and Yizhe Zhang. 2024. “Training Software Engineering Agents and Verifiers with SWE-Gym.” <https://arxiv.org/abs/2412.21139>.
- Pink, Mathis, Qinyuan Wu, Vy Ai Vo, Javier Turek, Jianing Mu, Alexander Huth, and Mariya Toneva. 2025. “Position: Episodic Memory Is the Missing Piece for Long-Term LLM Agents.” <https://arxiv.org/abs/2502.06975>.
- Pourcel, Julien, Cédric Colas, and Pierre-Yves Oudeyer. 2025. “Self-Improving Language Models for Evolutionary Program Synthesis: A Case Study on ARC-AGI.” <https://arxiv.org/abs/2507.14172>.
- Pu, Yuan, Yazhe Niu, Zhenjie Yang, Jiyuan Ren, Hongsheng Li, and Yu Liu. 2024. “UniZero: Generalized and Efficient Planning with Scalable Latent World Models.” <https://arxiv.org/abs/2406.10667>.
- Qiu, Zhisong, Shuofei Qiao, Kewei Xu, Yuqi Zhu, Lun Du, Ningyu Zhang, and Huajun Chen. 2026. “Rewarding the Scientific Process: Process-Level Reward Modeling for Agentic Data Analysis.” <https://arxiv.org/abs/2604.24198>.
- Rahmani, Babak. 2026. “Debugging Code World Models.” <https://arxiv.org/abs/2602.07672>.
- Reed, Scott, and Nando de Freitas. 2015. “Neural Programmer-Interpreters.” <https://arxiv.org/abs/1511.06279>.
- Ribeiro, Francisco, Claudio Spiess, Prem Devanbu, and Sarah Nadi. 2025. “On LLMs’ Internal Representation of Code Correctness.” <https://arxiv.org/abs/2512.07404>.
- Richens, Jonathan, David Abel, Alexis Bellot, and Tom Everitt. 2025. “General Agents Contain World Models.” <https://arxiv.org/abs/2506.01622>.
- Robeys, Maxime, Martin Szummer, and Laurence Aitchison. 2025. “A Self-Improving Coding Agent.” <https://arxiv.org/abs/2504.15228>.

- Rodionov, Sergey. 2026. “Executable World Models for ARC-AGI-3 in the Era of Coding Agents.” <https://arxiv.org/abs/2605.05138>.
- Sahoo, Subramanyam. 2025. “The Double Life of Code World Models: Provably Unmasking Malicious Behavior Through Execution Traces.” <https://arxiv.org/abs/2512.13821>.
- Samadi, Mehrzad, Aleksander Ficek, Sean Narenthiran, Siddhartha Jain, Wasi Uddin Ahmad, Somshubra Majumdar, Vahid Noroozi, and Boris Ginsburg. 2025. “Scaling Test-Time Compute to Achieve IOI Gold Medal with Open-Weight Models.” <https://arxiv.org/abs/2510.14232>.
- Schultz, John, Jakub Adamek, Matej Jusup, Marc Lanctot, Michael Kaisers, Sarah Perrin, Daniel Hennes, et al. 2024. “Mastering Board Games by External and Internal Planning with Language Models.” <https://arxiv.org/abs/2412.12119>.
- Shi, Zhan, Kevin Swersky, Daniel Tarlow, Parthasarathy Ranganathan, and Milad Hashemi. 2019. “Learning Execution Through Neural Code Fusion.” <https://arxiv.org/abs/1906.07181>.
- Shinn, Noah, Federico Cassano, Edward Berman, Ashwin Gopinath, Karthik Narasimhan, and Shunyu Yao. 2023. “Reflexion: Language Agents with Verbal Reinforcement Learning.” <https://arxiv.org/abs/2303.11366>.
- Singh, Gautam, Arjun Guha, Bhavya Kaikhura, and Harshitha Menon. 2026. “Learning Reasoning World Models for Parallel Code.” <https://arxiv.org/abs/2604.20926>.
- Sun, Chuyue, Yican Sun, Daneshvar Amrollahi, Ethan Zhang, Shuvendu Lahiri, Shan Lu, David Dill, and Clark Barrett. 2025. “VeriStruct: AI-Assisted Automated Verification of Data-Structure Modules in Verus.” <https://arxiv.org/abs/2510.25015>.
- Sun, Shuang, Huatong Song, Lisheng Huang, Jinhao Jiang, Ran Le, Zhihao Lv, Zongchao Chen, et al. 2026. “SWE-World: Building Software Engineering Agents in Docker-Free Environments.” <https://arxiv.org/abs/2602.03419>.
- Tahimic, Kriz, and Charibeth Cheng. 2025. “Mechanistic Interpretability of Code Correctness in LLMs via Sparse Autoencoders.” <https://arxiv.org/abs/2510.02917>.
- Tan, Hanzhuo, Weihao Li, Xiaolong Tian, Siyi Wang, Jiaming Liu, Jing Li, and Yuqun Zhang. 2025. “SK2Decompile: LLM-Based Two-Phase Binary Decompilation from Skeleton to Skin.” <https://arxiv.org/abs/2509.22114>.
- Tang, Lingxiao, He Ye, Zhaoyang Chu, Muyang Ye, Zhongxin Liu, Xiaoxue Ren, and Lingfeng Bao. 2026. “ExecVerify: White-Box RL with Verifiable Stepwise Rewards for Code Execution Reasoning.” <https://arxiv.org/abs/2603.11226>.
- team, FAIR CodeGen, Jade Copet, Quentin Carbonneaux, Gal Cohen, Jonas Gehring, Jacob Kahn, Jan-nik Kossen, et al. 2025. “CWM: An Open-Weights LLM for Research on Code Generation with World Models.” <https://arxiv.org/abs/2510.02387>.
- Teng, Fu, Miao Pan, Xuhong Zhang, Zhezhi He, Yiyao Yang, Xinyi Chai, Mengnan Qi, Liqiang Lu, and Jianwei Yin. 2025. “VERIRL: Boosting the LLM-Based Verilog Code Generation via Reinforcement Learning.” <https://arxiv.org/abs/2508.18462>.
- Thakur, Amitayush, Jasper Lee, George Tsoukalas, Meghana Sistla, Matthew Zhao, Stefan Zetsche, Greg Durrett, Yisong Yue, and Swarat Chaudhuri. 2025. “CLEVER: A Curated Benchmark for Formally Verified Code Generation.” <https://arxiv.org/abs/2505.13938>.

- Thimmaiah, Aditya, Jiyang Zhang, Jayanth Srinivasa, Junyi Jessy Li, and Milos Gligoric. 2025. “PLSemanticsBench: Large Language Models as Programming Language Interpreters.” <https://arxiv.org/abs/2510.03415>.
- Tu, Haoxin, Huan Zhao, Yahui Song, Mehtab Zafar, Ruijie Meng, and Abhik Roychoudhury. 2025. “Agentic Verification of Software Systems.” <https://arxiv.org/abs/2511.17330>.
- Ugare, Shubham, and Satish Chandra. 2026. “Agentic Code Reasoning.” <https://arxiv.org/abs/2603.01896>.
- Wang, Jian, Xiaofei Xie, Qiang Hu, Shangqing Liu, and Yi Li. 2025. “Do Code Semantics Help? A Comprehensive Study on Execution Trace-Based Information for Code Large Language Models.” <https://arxiv.org/abs/2509.11686>.
- Wang, Ke, Rishabh Singh, and Zhendong Su. 2017. “Dynamic Neural Program Embedding for Program Repair.” <https://arxiv.org/abs/1711.07163>.
- Wang, Wenhan, Kaibo Liu, An Ran Chen, Ge Li, Zhi Jin, Gang Huang, and Lei Ma. 2024. “Python Symbolic Execution with LLM-Powered Code Generation.” <https://arxiv.org/abs/2409.09271>.
- Wang, Wenyi, Piotr Piękos, Li Nanbo, Firas Laakom, Yimeng Chen, Mateusz Ostaszewski, Mingchen Zhuge, and Jürgen Schmidhuber. 2025. “Huxley-Gödel Machine: Human-Level Coding Agent Development by an Approximation of the Optimal Self-Improving Machine.” <https://arxiv.org/abs/2510.21614>.
- Wang, Xingyao, Yangyi Chen, Lifan Yuan, Yizhe Zhang, Yunzhu Li, Hao Peng, and Heng Ji. 2024. “Executable Code Actions Elicit Better LLM Agents.” <https://arxiv.org/abs/2402.01030>.
- Wang, Yongpan, Xin Xu, Xiaojie Zhu, Xiaodong Gu, and Beijun Shen. 2025. “SALT4Decompile: Inferring Source-Level Abstract Logic Tree for LLM-Based Binary Decompilation.” <https://arxiv.org/abs/2509.14646>.
- Wang, Yunkun, Yue Zhang, Guochang Li, Chen Zhi, Binhua Li, Fei Huang, Yongbin Li, and Shuiguang Deng. 2025. “InspectCoder: Dynamic Analysis-Enabled Self Repair Through Interactive LLM-Debugger Collaboration.” <https://arxiv.org/abs/2510.18327>.
- Wei, Anjiang, Jiannan Cao, Ran Li, Hongyu Chen, Yuhui Zhang, Ziheng Wang, Yuan Liu, et al. 2025. “EquiBench: Benchmarking Large Language Models’ Reasoning about Program Semantics via Equivalence Checking.” <https://arxiv.org/abs/2502.12466>.
- Wei, Anjiang, Tarun Suresh, Jiannan Cao, Naveen Kannan, Yuheng Wu, Kai Yan, Thiago S. F. X. Teixeira, Ke Wang, and Alex Aiken. 2025. “CodeARC: Benchmarking Reasoning Capabilities of LLM Agents for Inductive Program Synthesis.” <https://arxiv.org/abs/2503.23145>.
- Wei, Anjiang, Huanmi Tan, Tarun Suresh, Daniel Mendoza, Thiago S. F. X. Teixeira, Ke Wang, Caroline Trippel, and Alex Aiken. 2025. “VeriCoder: Enhancing LLM-Based RTL Code Generation Through Functional Correctness Validation.” <https://arxiv.org/abs/2504.15659>.
- Wei, Yuxiang, Olivier Duchenne, Jade Copet, Quentin Carbonneaux, Lingming Zhang, Daniel Fried, Gabriel Synnaeve, Rishabh Singh, and Sida I. Wang. 2025. “SWE-RL: Advancing LLM Reasoning via Reinforcement Learning on Open Software Evolution.” <https://arxiv.org/abs/2502.18449>.
- Wei, Yuxiang, Zhiqing Sun, Emily McMilin, Jonas Gehring, David Zhang, Gabriel Synnaeve, Daniel Fried, Lingming Zhang, and Sida Wang. 2025. “Toward Training Superintelligent Software Agents Through Self-Play SWE-RL.” <https://arxiv.org/abs/2512.18552>.

- Xia, Chunqiu Steven, Yinlin Deng, Soren Dunn, and Lingming Zhang. 2024. "Agentless: Demystifying LLM-Based Software Engineering Agents." <https://arxiv.org/abs/2407.01489>.
- Xie, Zhihui, Jiacheng Ye, Lin Zheng, Jiahui Gao, Jingwei Dong, Zirui Wu, Xueliang Zhao, et al. 2025. "Dream-Coder 7b: An Open Diffusion Language Model for Code." <https://arxiv.org/abs/2509.01142>.
- Xu, Ruiyang, Jialun Cao, Yaojie Lu, Ming Wen, Hongyu Lin, Xianpei Han, Ben He, Shing-Chi Cheung, and Le Sun. 2024. "CRUXEval-x: A Benchmark for Multilingual Code Reasoning, Understanding and Execution." <https://arxiv.org/abs/2408.13001>.
- Yan, Chuanhao, Fengdi Che, Xuhan Huang, Xu Xu, Xin Li, Yizhi Li, Xingwei Qu, et al. 2025. "Re-form – Reducing Human Priors in Scalable Formal Software Verification with RL in LLMs: A Preliminary Study on Dafny." <https://arxiv.org/abs/2507.16331>.
- Yan, Shenao, Shen Wang, Yue Duan, Hanbin Hong, Kiho Lee, Doowon Kim, and Yuan Hong. 2024. "An LLM-Assisted Easy-to-Trigger Backdoor Attack on Code Completion Models: Injecting Disguised Vulnerabilities Against Strong Detection." <https://arxiv.org/abs/2406.06822>.
- Yang, Fanpeng, Xu Ma, Shuling Wang, Xiong Xu, Qinxiong Cao, Naijun Zhan, Xiaofeng Li, and Bin Gu. 2025. "Integrating Symbolic Execution with LLMs for Automated Generation of Program Specifications." <https://arxiv.org/abs/2506.09550>.
- Yang, Jian, Wei Zhang, Jiajun Wu, Junhang Cheng, Tuney Zheng, Fanglin Xu, Weicheng Gu, et al. 2026. "InCoder-32b-Thinking: Industrial Code World Model for Thinking." <https://arxiv.org/abs/2604.03144>.
- Yang, John, Carlos E. Jimenez, Alexander Wettig, Kilian Lieret, Shunyu Yao, Karthik Narasimhan, and Ofir Press. 2024. "SWE-Agent: Agent-Computer Interfaces Enable Automated Software Engineering." <https://arxiv.org/abs/2405.15793>.
- Yang, Sherry. 2026. "World Models as an Intermediary Between Agents and the Real World." <https://arxiv.org/abs/2602.00785>.
- Yao, Shunyu, Dian Yu, Jeffrey Zhao, Izhak Shafran, Thomas L. Griffiths, Yuan Cao, and Karthik Narasimhan. 2023. "Tree of Thoughts: Deliberate Problem Solving with Large Language Models." <https://arxiv.org/abs/2305.10601>.
- Yu, Xiao, Baolin Peng, Ruize Xu, Michel Galley, Hao Cheng, Suman Nath, Jianfeng Gao, and Zhou Yu. 2025. "Dyna-Think: Synergizing Reasoning, Acting, and World Model Simulation in AI Agents." <https://arxiv.org/abs/2506.00320>.
- Yu, Xiao, Baolin Peng, Ruize Xu, Yelong Shen, Pengcheng He, Suman Nath, Nikhil Singh, Jiangfeng Gao, and Zhou Yu. 2026. "Reinforcement World Model Learning for LLM-Based Agents." <https://arxiv.org/abs/2602.05842>.
- Zaremba, Wojciech, and Ilya Sutskever. 2014. "Learning to Execute." <https://arxiv.org/abs/1410.4615>.
- Zeng, Zhiyuan, Yichi Zhang, Yong Shan, Kai Hua, Siyuan Fang, Zhaiyu Liu, Jiaheng Liu, et al. 2026. "Understanding by Reconstruction: Reversing the Software Development Process for LLM Pretraining." <https://arxiv.org/abs/2603.11103>.
- Zhang, Jenny, Shengran Hu, Cong Lu, Robert Lange, and Jeff Clune. 2025. "Darwin Godel Machine: Open-Ended Evolution of Self-Improving Agents." <https://arxiv.org/abs/2505.22954>.

- Zhang, Ruiyi, Peijia Qin, Qi Cao, Eric Xue, and Pengtao Xie. 2026. “FunPRM: Function-as-Step Process Reward Model with Meta Reward Correction for Code Generation.” <https://arxiv.org/abs/2601.22249>.
- Zhang, Yaolun, Yinxu Pan, Yudong Wang, and Jie Cai. 2024. “PyBench: Evaluating LLM Agent on Various Real-World Coding Tasks.” <https://arxiv.org/abs/2407.16732>.
- Zhang, Yuntong, Haifeng Ruan, Zhiyu Fan, and Abhik Roychoudhury. 2024. “AutoCodeRover: Autonomous Program Improvement.” <https://arxiv.org/abs/2404.05427>.
- Zhang, Yuxiang, Jiangming Shu, Ye Ma, Xueyuan Lin, Shangxi Wu, and Jitao Sang. 2025. “Memory as Action: Autonomous Context Curation for Long-Horizon Agentic Tasks.” <https://arxiv.org/abs/2510.12635>.
- Zhang, Yuxiang, Shangxi Wu, Yuqi Yang, Jiangming Shu, Jinlin Xiao, Chao Kong, and Jitao Sang. 2024. “O1-Coder: An O1 Replication for Coding.” <https://arxiv.org/abs/2412.00154>.
- Zhou, Shuyan, Frank F. Xu, Hao Zhu, Xuhui Zhou, Robert Lo, Abishek Sridhar, Xianyi Cheng, et al. 2023. “WebArena: A Realistic Web Environment for Building Autonomous Agents.” <https://arxiv.org/abs/2307.13854>.
- Zhuge, Mingchen, Changsheng Zhao, Haozhe Liu, Zijian Zhou, Shuming Liu, Wenyi Wang, Ernie Chang, et al. 2026. “Neural Computers.” <https://arxiv.org/abs/2604.06425>.