
A Comprehensive Survey of World Models for Coding

Keon Kim¹

¹Independent Researcher



Code & Survey



arXiv




Figure 1: Twelve-year arc of code world models. Three parallel lineages — neural execution, world models and RL, and code LLMs — converge in 2025–2026 around the named artifact CWM. Single-diamond markers in the figure tag the first paper to use “world model” or “Dreamer” in a contribution name; the double diamond marks CWM as the named open-weights artifact.

0.1 Abstract

A *world model* is an internal predictor over environment dynamics, used to imagine the consequences of actions. In coding, the environment is the program: its runtime state, execution trace, filesystem, tests, and developer task. Zaremba and Sutskever (Zaremba and Sutskever, 2014) trained an LSTM to predict Python execution output in 2014. Meta FAIR released CWM (FAIR CodeGen Team et al., 2025), the first open-weights LLM branded a Code World Model, in October 2025. Internal models of execution are learnable. Whether they are necessary, architecturally separable, or causally responsible for coding-agent gains is not yet established.

This survey synthesizes 184 papers under two definitions — a permissive *descriptive* one matching field usage, and a strict *normative* one requiring forward-prediction machinery. It traces a twelve-year arc; builds a three-axis taxonomy (functionality $\times$ temporal $\times$ representation); produces system cards for thirteen representative systems; assembles protocol-stratified tables for SWE-bench, CRUXEval, web agents, and formal verification; develops seven critical theses; and lists open problems. The defensible empirical claim is that execution-grounded supervision improves code agents on runtime-reasoning benchmarks. The defensible critical claim is that broader inferences to “world models for coding” remain underdetermined.

0.2 Theses Summary

The seven critical theses developed in §17, each grounded in a specific disconfirmation from the corpus:

1. The “world model” label has become overloaded. Under the field’s permissive descriptive definition (D, §3.1), CWM, TRACED, SemCoder, LLM-JEPA, and DyMo are all “world models” despite being architecturally standard LLMs with enriched training. A strict normative definition (N1/N2/N3, §3.1) admits only systems with explicit forward-prediction machinery: neural dynamics models (N1, no code exemplar), latent-action rollout models such as CoLA (N2), and synthesized executable simulators such as GIF-MCTS and ARC executable WMs (N3).
2. Trace pretraining has a causal-isolation problem. “Do Code Semantics Help?” (Wang et al., 2025a) finds no trace representation consistently improves code synthesis across multiple backbones; CWM’s 65.8% SWE-bench Verified is confounded between trace mid-training, ForagerAgent trajectories, and RL.
3. The Dreamer-for-code gap may be a non-problem. Vision needed latent rollouts because pixels are expensive; Python state is small and CPython is free, so the architectural pressure does not transfer. CWM in token space reaches 65.8% on SWE-bench Verified at 32B parameters.
4. PRMs are critics, not world models. A learned evaluator of partial trajectories cannot roll out future states; the conceptual conflation dilutes the world-model vocabulary.
5. “General Agents Contain World Models” (Richens et al., 2025) is weaker than its title. The theorem applies under restrictive assumptions (full observability, stationarity, deep-composite-goal regret bounds) that SWE agents demonstrably violate.
6. The verifier-grounded lineage is the actual leading edge, but scoped. ATLAS, Re:Form, CLEVER, VeriStruct, AutoRocq produce machine-checkable correctness. Scoped: they lead on synthesis-from-spec (e.g., ATLAS DafnySynthesis 65.8% pass@5); they do *not* lead on end-to-end verified codegen from NL (CLEVER 71/161 Lean).
7. The evaluation gap is structural. No benchmark holds policy fixed and varies WM quality. Models with 85–98% output-prediction accuracy still produce traces with systematic errors throughout — outcome accuracy decouples from process fidelity.

0.3 1. Introduction

Autoregressive code LLMs generate tokens conditioned on syntactic context. Correct programs are objects in two domains: source tokens, and the values, control flow, side effects, and developer intent those tokens encode. The world-model framing comes from model-based reinforcement learning, where it names an internal predictor of environment dynamics. The hypothesis examined in this survey is that training on execution rather than only on source-token prediction reduces the gap between code that compiles and code that runs.

Two adjacent surveys cover non-overlapping ground. A Survey on LLMs for Code Generation (Jiang et al., 2024) maps the code-LLM space without the world-model lens. Understanding World or Predicting Future (Ding et al., 2024a) maps world models in general without the code lens. A Comprehensive Survey on World Models for Embodied AI (Li et al., 2025b) maps embodied world models but excludes the coding domain. The intersection — the subject of this document — has cohered only recently into a recognizable program.

0.4 2. Methodology

The corpus is 184 arxiv preprints, assembled in four passes (March–May 2026) seeded from CWM (FAIR CodeGen Team et al., 2025) via citation BFS and targeted topic searches. A paper

enters the corpus if it intersects both world-model/state-tracking architectures and code generation, debugging, repair, or agentic coding. Pure vision world models (DreamerV1–V3, V-JEPA, Genie) and pure code-LLM papers without a world-model angle are excluded except as cited precedent. Date cutoff: 2026-05-15. One curator did the taxonomy coding; numerical claims in §§6–14 derive from abstracts, §16 from source PDFs. 60% of the corpus is 2025 or later.

0.5 3. Defining a World Model for Coding

The literature uses “world model” in two distinct senses that this survey will carefully separate.

0.5.1 3.1 Two definitions

Definition D (descriptive, permissive). A *world model for coding* is any code LLM whose training objective concretely encodes program semantics — execution traces, runtime state, environment feedback, or simulated outcomes — in addition to or instead of source-token prediction. This is the field’s current usage. Under D, CWM, TRACED, SemCoder, LLM-JEPA, DyMo, RLEF, and most papers in the corpus are world models.

Definition N (normative, strict). A *world model for coding* is a system with an explicit, separable forward-prediction mechanism. We split N into three sub-types:

- N1 — Neural dynamics model. A learned $W : (\text{state}, \text{action}) \rightarrow \text{next_state}$ instantiated as a distinct architectural component with latent recurrent state, separate dynamics head, or inverse model. Dreamer-class. No code exemplar in the corpus.
- N2 — Latent-action model. A learned action abstraction over a base LLM, with the LLM serving as transition model in compressed action space, supporting rollout or tree search over latent actions. CoLA (Jia et al., 2025) is the canonical example.
- N3 — Synthesized executable simulator. The world model is an executable program (Python, DSL) synthesized by an LLM and run against ground-truth transitions during planning. GIF-MCTS (Dainese et al., 2024), WorldCoder, Executable WMs for ARC-AGI-3 (Rodionov, 2026).

Each N-subtype has a distinct architectural commitment. N1 commits to learned latent dynamics; N2 commits to a discrete latent-action space; N3 commits to program synthesis as the dynamics. Conflating them — as the loose “Dreamer-for-LLMs gesture” framing did — obscures which architectural bet is being made. Under any N-subtype, CWM, TRACED, SemCoder, and most of the corpus do *not* qualify.

The two definitions agree on the empirical fact that execution-grounded supervision helps. They disagree on whether that grounding is correctly named a *world model* in the model-based-RL sense.

Aspect	Definition D (descriptive)	Definition N (normative)
Granted to	Any LLM trained with execution-related signal	Only systems with explicit forward-prediction module
Includes CWM	Yes	No (architecturally a Llama-class decoder)
Includes TRACED, SemCoder, NExT	Yes	No (auxiliary heads on a Transformer)
Includes CoLA, GIF-MCTS	Yes	Yes
Includes PRMs (ExecVerify, ThinkPRM)	Often grouped	No (critics, not forward predictors)
Includes verifiers (Lean, Dafny, Z3)	Sometimes grouped	No (deterministic oracles, not learned WMs)
Includes Dreamer / V-JEPA (vision precedents)	Yes	Yes

This survey uses D throughout the catalog sections (§§6–16) because that is how the literature talks, and N in §17 because the critical synthesis depends on it. We mark the switch each time it matters. The two-definition framework, in our reading, is the cleanest way to resolve the tension.

0.5.2 3.2 What is modeled

Orthogonal to D-vs-N, *what* is modeled varies: variable values and stack frames (CWM, CodeExecutor); execution traces (NExT, SemCoder, TRACED); test outcomes (LEVER, RLEF); OS/web environment state (WebDreamer, Dyna-Think); repository state (RepoGraph); spec or developer intent (ATLAS, Re:Form); adversarial behavior (Double Life of CWMs). A system typically commits strongly to one or two of these.

CWM occupies an instructive position: behaviorally explicit (emits stack frames) but architecturally implicit (a 32B Llama-class decoder with no separate dynamics head). SemCoder, NExT, and TRACED follow the same pattern. The “explicit WM” label flattens this distinction; §17.1 develops the consequence.

0.6 4. Twelve Years of Code World Models

The lineage is best understood as a sequence of inheriting questions. Each era’s answer dissolved the previous era’s bottleneck and exposed the next.

(See Figure 1 at the start of the paper for the timeline diagram.)

Diamond markers in the figure tag the moments where “world model” enters the *name* of the contribution.

Pre-2020 — Can a network execute code? Learning to Execute (Zaremba and Sutskever, 2014) trained an LSTM to predict Python output on bounded-loop programs. Neural Programmer-Interpreters (Reed and de Freitas, 2015) and Differentiable Forth built differentiable program counters; Dynamic Neural Program Embedding (Wang et al., 2017) embedded real interpreter traces; Neural Code Fusion (Shi et al., 2019) and IPA-GNN (Bieber et al., 2020) used GNN attention as a program counter. By 2020 the answer was yes — but only with the interpreter encoded into architecture, which did not scale to real languages. Ha & Schmidhuber (Ha and Schmidhuber, 2018) had named the pattern for vision-RL; the coding community had not yet borrowed it.

2020–2022 — Can the training include execution? Codex (Chen et al., 2021) and MBPP (Austin et al., 2021) made code generation a target. Scratchpads (Nye et al., 2021) showed a Transformer could predict program output if allowed to emit intermediate computation first — the Dynamic-NPE move at LLM scale, without architectural surgery.

2023 — Trace pretraining as a named recipe. CodeExecutor (Liu et al., 2023) trained a Transformer to emit per-line state traces from source; TRACED (Ding et al., 2023) generalized this to a pretraining auxiliary; CRUXEval (Gu et al., 2024a) provided the eval that became standard. In parallel, Reflexion (Shinn et al., 2023) / Self-Debug (Chen et al., 2023) used execution feedback in-context; LEVER (Ni et al., 2023) used it during decoding; RAP (Hao et al., 2023) ran MCTS over an LLM-as-world-model.

2024 — From models that simulate to agents that act. SWE-bench (Jimenez et al., 2023) shifted the task to real GitHub repos. CodeAct (Wang et al., 2024b), SWE-agent (Yang et al., 2024), NExT (Ni et al., 2024), RLEF (Gehring et al., 2024), WebDreamer (Gu et al., 2024b), and Generating Code World Models via MCTS ((Dainese et al., 2024), source of the name) define the agentic turn. The WM moved from weights to loop.

2025 — The CWM moment. DeepSeek-R1 (DeepSeek-AI et al., 2025) and SWE-RL (Wei et al., 2025d) scaled reasoning-RL. CoLA (Jia et al., 2025) and LLM-JEPA (Huang et al., 2025) ported Dreamer- and JEPA-style ideas to LLMs. General Agents Contain World Models (Richens et al., 2025) provided the existence theorem. CWM (FAIR CodeGen Team et al., 2025) — 32B open-

weights, mid-trained on 5T tokens of Python traces plus 3M ForagerAgent trajectories — made the artifact downloadable.

2026 — Stress-testing and broadening. Debugging CWMs (Rahmani, 2026) and Demystifying Errors in LLM Reasoning Traces (Abdollahi et al., 2025) catalog failures. Industrial CWM (Yang et al., 2026), Parallel-Code WMs (Singh et al., 2026), and Executable WMs for ARC-AGI-3 (Rodionov, 2026) generalize the recipe. Reinforcement World Model Learning (Yu et al., 2026) trains the WM rather than the policy.

0.7 5. Taxonomy: Three Axes of Code World Models

Two cuts at the taxonomy are useful, and they are complementary. The first is a *lineage* cut — which research thread produced the system — and is the basis for §§6–13. The second, more durable cut adapts the three-axis framework of Li et al. ((Li et al., 2025b), *A Comprehensive Survey on World Models for Embodied AI*) to the code domain. The lineage map is below; the three axes are developed in §§5.1–5.3.

Representation	Vision analog	Code-WM exemplars
Token Sequence (TS)	Token-as-pixel (IRIS, Genie, Sora)	CWM, CodeExecutor, TRACED, NExT, SemCoder
Synthesized executable (N3)	—	GIF-MCTS, WorldCoder, ARC executable WMs
Global Latent Vector (GLV)	RSSM (Dreamer V1–V3)	white space — no code exemplar
Spatial / Structural Grid (SLG)	OccWorld, DriveWorld (BEV/voxel)	white space — no code exemplar
Decomposed Object / Slot (DOR)	SlotFormer, object-centric WMs	white space — no code exemplar

Verifiers (Lean / Dafny / Z3) and PRMs are intentionally absent: a verifier is a grounding oracle, not a representation; a PRM is a critic, not a forward predictor.

Figure 2: Representation taxonomy of code world models. Each row is one representation class (left), its vision precedent (centre), and the code-WM systems that instantiate it (right). Three classes — Global Latent Vector (Dreamer-style), Spatial/Structural Grid, and Decomposed Object/Slot — have no code exemplar in the corpus and are the survey’s white-space gaps.

Adjacent WM surveys converge on overlapping splits that the three-axis framework subsumes as projections. Ding et al. (Ding et al., 2024a) split top-level by *implicit representation* vs *future prediction*. JiahuaDong’s awesome-list organizes by *paradigm*: RL-based / observation-generative / latent-space / object-centric. knightnemo’s list surfaces *pixel* vs *mesh* vs *latent* as cross-cutting tags. The three axes — functionality, temporal modeling, and representation — capture the choices a system makes regardless of which lineage it belongs to.

0.7.1 5.1 Axis 1 — Functionality

- Decision-coupled WMs model only the slice of the world relevant to acting on it. CWM (FAIR CodeGen Team et al., 2025), RLEF (Gehring et al., 2024), and WebDreamer (Gu et al., 2024b) are decision-coupled — their WMs exist to enable code generation, RL planning, or web navigation respectively. CWM does not predict global filesystem state; it predicts the next Python frame *because* the next action depends on it.
- General-purpose WMs model the environment without reference to a particular task. The “general agents contain world models” theorem (Richens et al., 2025) is the abstract limit. In the corpus, only the largest CWM-class models with broad mid-training ap-

proach generality; most “code world models” are decision-coupled to a sub-task (repair, completion, agent control).

0.7.2 5.2 Axis 2 — Temporal Modeling

- Sequential simulation/inference. Step-by-step autoregressive rollout. CWM, NExT, SemCoder, all the trace-pretraining systems, and most LLM-as-WM planners (RAP, WebDreamer in its MPC loop) live here. The state is updated one timestep at a time. The vision analog is RSSM (Hafner et al., DreamerV1–V3).
- Global difference prediction. Predict the entire future state at once, in parallel. The vision analog is video-diffusion or masked-JEPA. In code, this fits diffusion code models (DiffuCoder, (Gong et al., 2025); Dream-Coder 7B, (Xie et al., 2025)) where the next state is sampled jointly rather than autoregressively, and the “specification is the program” framing (Monperrus, 2026) where the entire trace is the spec.
- Static, no-trace. Some systems (SemCoder’s static mode, the trace-free baselines in “Do Code Semantics Help?”) explicitly drop temporal modeling at inference, reducing to single-shot prediction.

0.7.3 5.3 Axis 3 — Representation

This is the axis where the WM literature has converged most strongly, and where the code-WM literature is most uneven. Adapting Li et al.’s four-category split (GLV / TFS / SLG / DRR) to code yields five classes, of which only two are well-populated.

Class	Encodes the world as	Vision analog	Code exemplars
Token Sequence (TS)	Discrete or continuous token streams with execution traces, variable bindings, or rationales interleaved with source	Token-as-pixel (IRIS, TWM, Genie, Sora)	CWM (FAIR CodeGen Team et al., 2025), CodeExecutor (Liu et al., 2023), TRACED (Ding et al., 2023), NExT (Ni et al., 2024), SemCoder (Ding et al., 2024b) — the dominant code-WM mode
Global Latent Vector (GLV)	A compact vector updated recurrently, encoding the entire program/agent state	RSSM (Hafner et al., DreamerV1–V3)	No clean exemplar. CoLA (Jia et al., 2025) introduces a learned action codebook but is otherwise a standard LLM, not RSSM-style
Spatial / Structural Grid (SLG)	A geometric or structural grid (BEV/voxel in vision; AST, call-graph, CFG in code)	OccWorld, DriveWorld	No exemplar. RepoGraph (Ouyang et al., 2024) uses a static dependency graph but does not predict over it as a WM
Decomposed Object / Slot (DOR)	Distinct persistent latent slots for objects in the world	SlotFormer and object-centric WMs	No exemplar. No code-WM models variables, scopes, or classes as discrete persistent slots

Class	Encodes the world as	Vision analog	Code exemplars
Synthesized executable (N3)	The world model is an executable program, synthesized rather than learned	(orthogonal to vision)	GIF-MCTS (Dainese et al., 2024), WorldCoder, Executable WMs for ARC-AGI-3 (Rodionov, 2026)

Verifiers (Lean, Dafny, Z3) and PRMs are intentionally absent from this table. A verifier is not a *representation* of the world; it is a *grounding oracle* over candidate artifacts. A PRM is a *backward-looking critic*, not a forward predictor. Both belong in §5.4 below.

Three white spaces. The Dreamer-style GLV, the object-centric DOR, and the spatial-grid SLG representations have not been instantiated for code, despite being mature for vision. §18 lists these as open problems, with the caveat that not all three are equally promising — the Dreamer-style gap is contestable (§17.3), while the object-centric and structural-grid gaps look more genuine.

0.7.4 5.4 Grounding mode (orthogonal to representation)

A second classification asks how each system’s predictions are validated. This is the analog of Li et al.’s “reality” column and addresses the decoupling between WM-head accuracy and downstream task success that DyMo (Guo et al., 2025) and §17.7 develop.

Grounding mode	Definition	Code exemplars
None / self-report	Predictions never validated against ground truth	RAP (Hao et al., 2023), basic LLM-as-WM planners
Execution-grounded	Predictions checked against real interpreter / runtime	CWM (FAIR CodeGen Team et al., 2025), TRACED (Ding et al., 2023), NExT (Ni et al., 2024), SemCoder (Ding et al., 2024b), RLEF (Gehring et al., 2024)
Verifier-grounded	Outputs checked by Lean / Dafny / Verus / Rocq / Z3	ATLAS (Baksys et al., 2025), Re:Form (Yan et al., 2025), CLEVER (Thakur et al., 2025), VeriStruct (Sun et al., 2025), AutoRocq (Tu et al., 2025)
Synthesized-simulator-checked	Synthesized world model checked against held-out transitions	GIF-MCTS (Dainese et al., 2024), Executable WMs for ARC-AGI-3 (Rodionov, 2026)
Critic-grounded (not a WM)	Backward-looking value/quality score, no rollout	ExecVerify (Tang et al., 2026), SWE-PRM (Gandhi et al., 2025), ThinkPRM (Khalifa et al., 2025) — listed for contrast; these are critics, not WMs (§17.4)

Grounding mode is orthogonal to the representation axis: a token-sequence representation can be execution-grounded (CWM) or self-report (RAP); an N3 synthesized-simulator can be checked against transitions (GIF-MCTS) or never validated. The “WM fidelity” question of §17.7 is, structurally, a question about whether a system’s grounding mode is strong enough to falsify a bad WM at training time.

0.8 6. Foundations: Neural Execution as Implicit World Modeling

Zaremba & Sutskever’s Learning to Execute (Zaremba and Sutskever, 2014) established both feasibility and brittleness. Show Your Work — Scratchpads (Nye et al., 2021) is the pivotal contribution: by training a Transformer to emit intermediate computation states, the authors recovered much of the LSTM-era execution-prediction performance at scale, presaging the trace-pretraining lineage of §6. CRUXEval (Gu et al., 2024a) and REval (Chen et al., 2024) provide the canonical execution-reasoning benchmarks. The lesson the field absorbed: *replacing* the interpreter with a neural network is harder than *augmenting* a transformer with interpreter-style supervision. Modern systems all take the latter path.

0.9 7. The Trace-Pretraining and CWM Lineage

0.9.1 7.1 Trace-pretraining as a recipe

CodeExecutor (Liu et al., 2023) trains a Transformer to simulate Python execution token-by-token. TRACED (Ding et al., 2023) adds dynamic-state supervision to a code-LLM pretraining mix. NExT (Ni et al., 2024) formats traces as natural-language rationales, letting a chat-style LLM reason about runtime behavior via chain-of-thought. SemCoder (Ding et al., 2024b) generalizes to “monologue reasoning” linking source-text to execution state.

The 2025 wave consolidated and stress-tested the approach. “What I cannot execute, I do not understand” (Armengol-Estapé et al., 2025) trains and evaluates LLMs explicitly on traces with dynamic scratchpads, pushing Llama-3.1-8B from 37.8% to ~80% on CRUXEval-O. Code Execution as Grounded Supervision (Jung et al., 2025) repurposes line-by-line traces as verifiable CoT. Self-Execution Simulation (Maimon et al., 2026) lets the model train on its own execution predictions. Demystifying Errors in LLM Reasoning Traces (Abdollahi et al., 2025) audits where trace-trained LLMs fail. “Do Code Semantics Help?” (Wang et al., 2025a) is the most damaging paper in the lineage: a comprehensive ablation across DeepSeek-Coder, LLaMA-3, and Gemma-2 with five trace representations finds that *no single representation consistently improves code generation*, and in 7 of 9 synthesis settings the no-trace baseline wins or ties.

0.9.2 7.2 TRACED (Ding et al., 2023)

TRACED augments a RoBERTa/UnixCoder pre-training mix with two execution-grounded heads on top of standard MLM: per-line program-state classification (variable type and quantized value over 30 bins) and per-line execution coverage. Trained on ~121k C traces from CodeNet collected via gdb, it shows that *quantized* variable-value prediction is a viable auxiliary signal — concrete values lose to discretized bins. On static execution estimation, full-path accuracy rises from UnixCoder’s 63.7% to 71.6%, and downstream clone retrieval and defect detection improve modestly. The essence: trace prediction as a pre-training side objective, not a separate model.

0.9.3 7.3 NExT (Ni et al., 2024)

NExT inlines execution traces into source as Python-style comments (`# (k) varA=...; varB=...`) and trains PaLM 2-L on (rationale, fix) candidates via STaR-style self-training — sample 32 candidates per problem, accept those that pass unit tests, SFT on the accepted set, repeat. After ten iterations Mbpp-R pass@1 climbs from 23.2% to 49.3% (+26.1 absolute). The result that matters most for the rest of the survey: NExT *retains* a +17 absolute gain (23 → 40.8) on Mbpp-R when traces are *removed at inference*, which makes it the cleanest example of trace pretraining transferring to non-trace inference on a repair task.

0.9.4 7.4 SemCoder (Ding et al., 2024b)

SemCoder formalizes “monologue reasoning” linking four code modalities — natural-language description, source, operational trace, and abstract input-invariant constraints — under a single NTP objective with rejection-sampled training data (the PYX corpus). Distinctive features: forward *and* backward monologues (NExT is forward-only), abstract-semantics constraints

rather than concrete state at every step, and entirely static inference. SemCoder-1.3B reaches CRUXEval-I/O 63.6/65.1 vs GPT-3.5-turbo’s 50.3/59.0, and a monologue-format ablation beats Scratchpad and NExT.

0.9.5 7.5 CWM (FAIR CodeGen Team et al., 2025)

CWM is a 32B dense decoder-only Transformer with grouped-query attention, sliding-window blocks, RoPE — architecturally Llama-class. What earns it the “world model” name is the mid-training datamix: 5T tokens of Python observation-action traces (120M traced functions, 262k CodeContests traces, 70k repo-level traced commits, 75M natural-language rewrites) plus 3M ForagerAgent SWE trajectories from 10.2k Dockerized repositories. Tokens are formatted so next-token prediction is next-state prediction at line granularity. CWM reaches 65.8% on SWE-bench Verified with test-time scaling (best@16 over 40 verifier-reranked samples), 94.3% on CRUXEval-Output. Its second technical contribution is *Activ*, which uses GitHub Actions CI to scale executable repository images. CWM is behaviorally explicit (emits stack frames) but architecturally implicit (no separate dynamics head).

Important caveat (developed further in §17): the 65.8% headline is *not* pure pass@1 but best-of-16 with verifier reranking. Pure pass@1 is approximately 53–55%. The trace-mid-training contribution is not causally isolated from the ForagerAgent-trajectory contribution and from the joint-RL contribution. Without an ablation removing one while holding the others fixed, the “world model” component’s causal role is unfalsifiable.

0.9.6 7.7 GIF-MCTS / Generating Code World Models via MCTS (Dainese et al., 2024)

GIF-MCTS treats the world model itself as a Python program — an `Environment.step(s,a) → (s', r, done)` class synthesized by an LLM to match a small batch of pre-collected (s, a, r, s', d) transitions. MCTS over partial programs uses three action types (*generate* lines, *improve* full program given a failing transition, *fix* runtime/syntax errors), with reward equal to the fraction of transitions reproduced correctly. The synthesized world model, once compiled, runs 4–6 orders of magnitude faster than calling an LLM as world model. On APPS-Competition it reaches 28.3% strict pass@20 (Llama-3-70B), beating WorldCoder’s 25.1%. The conceptual contribution is to *search over candidate Python world-model programs* rather than train a neural one.

0.10 8. World Models for Code Agents

Once an LLM is an *agent* taking actions in a non-trivial environment, the world-model question becomes whether the agent simulates the environment’s response. Three sub-environments dominate.

0.10.1 8.1 Web agents

Web Agents with World Models (Chae et al., 2024) systematizes the thread. DyMo / World Modeling Improves LM Agents (Guo et al., 2025) adds a next-state prediction head to function-calling agents and reports gains on BFCL-V2 — though with a caveat (§17): the WM head reaches 90–94% state-prediction accuracy while the underlying policy reaches only 72.8% task success, illustrating that WM-head accuracy and agent accuracy can decouple.

WebDreamer (Gu et al., 2024b). WebDreamer treats the web as a POMDP in which the LLM imagines natural-language state-change descriptions for each candidate click, type, or select. A specialist *Dreamer-7B* (Qwen2-VL-7B fine-tuned on 3.1M synthesized (initial visual state, action, state-change) tuples from random walks over Common Crawl URLs) provides cheap roll-outs. At inference, model-predictive control samples actions, scores simulated trajectories with GPT-4o on a 3-scale rubric, and executes the argmax. On VisualWebArena, Online-Mind2Web, and Mind2Web-Live, this beats the reactive baseline by +34/+42/+24% relative (⌈+6–11 absolute) while running 4–5× faster than tree search. The bet: when actions are irreversible (forms, purchases), one-step MPC over an LLM-as-WM beats backtracking search.

0.10.2 8.2 OS / computer-use agents

Reinforcement World Model Learning for LLM-based Agents (Yu et al., 2026) and World Models as an Intermediary between Agents and the Real World (Yang, 2026) generalize the lens: a learned WM mediates between LLM and expensive environment.

Dyna-Think (Yu et al., 2025). Dyna-Think trains a single Qwen2.5-32B to internalize world-model simulation inside its `<think>` block for OSWorld and WindowsAgentArena. Two stages: DIT (imitation learning on R1 traces cleaned to keep only WM-simulation text) and DDT (Dyna-Q-style joint training over three WM heads — next state, state-diff, critic-prediction — with rejection-sampled policy updates). On OSWorld BoN the 32B model reaches 43.1, essentially matching DeepSeek-R1 at 685B with half the tokens. World-model accuracy correlates with task success at $r=0.32$ across models. This is the cleanest instance in the corpus of policy and learned world model hosted in the *same* LLM.

0.10.3 8.3 SWE agents

SWE-bench (Jimenez et al., 2023) and SWE-Gym (Pan et al., 2024) defined the eval and training environment respectively. CodeAct (Wang et al., 2024b) made the Python interpreter the unified action space. Reflexion (Shinn et al., 2023) was the earliest entry with episodic verbal RL. Nanbeige SWE-World (Sun et al., 2026) trains a learned Docker-free execution surrogate. Understanding by Reconstruction (Zeng et al., 2026) reverses the development process to harvest agentic pretraining traces. SWE-TRACE (Han et al., 2026) provides process-level reward modeling over trajectories. Self-Play SWE-RL (Wei et al., 2025e) introduces adversarial bug-injection/repair self-play. Bootstrapping Coding Agents — The Specification Is the Program (Monperrus, 2026) reframes the SWE task itself as a programmatic spec.

The §16 empirical synthesis separates three regimes: execution-grounded open-weight model training (CWM, SWE-RL), execution-grounded agents with learned simulators (Nanbeige SWE-World), and scaffold evolution around closed-model executors (Darwin GM, Huxley GM). Under best-of-k or verifier-reranked protocols, several systems report 60–68% on SWE-bench Verified, but those numbers are not directly comparable to pass@1 model-training results, and the scaffold-evolved systems are not “32B open-weight world-model-trained” — they run frontier closed models inside an evolved harness.

0.11 9. RL with Execution as the World Signal

The model-based-RL framing — world model is what the policy plans over — has produced a clean lineage.

0.11.1 9.1 RLEF (Gehring et al., 2024)

RLEF formulates iterative code synthesis as a POMDP — actions are full code responses, observations are formatted public-test execution feedback, rewards come from held-out *private* tests. Standard PPO with KL regularization and a 3-turn limit. Llama-3.1-70B+RLEF reaches 37.5/40.1 pass@1 valid/test on CodeContests at budget 1@3 (vs 25.9/27.5 baseline), matching AlphaCodium-GPT-4 with 5 samples; at 10@100 it reaches 54.5/54.5, surpassing the AlphaCode 41B+clustering baseline. Critically, a random-feedback ablation removes the entire gain, isolating that the model is learning to *use* execution feedback rather than just sample more.

0.11.2 9.2 SWE-RL (Wei et al., 2025d)

SWE-RL applies GRPO to 273k high-quality PR seeds with a rule-based, continuous reward (`difflib.SequenceMatcher` similarity between predicted and oracle patch) — *no code execution at training time*. Llama-3.3-70B fine-tuned this way (Llama3-SWE-RL-70B) hits 41.0% pass@1 on SWE-bench Verified with the Agentless Mini scaffold. The surprising result is OOD transfer: HumanEval+ 76.2% 79.9, CRUXEval-O 61.9% 75.5, MATH 70.9% 73.7, while SFT on the same data *degrades* on these. Continuous reward beats discrete in ablation (34.8 vs 29.0 oracle-

repair). The thesis: partial-credit similarity rewards on real PR patches induce reasoning patterns that transfer beyond the training distribution.

0.11.3 9.3 Process Reward Models

ExecVerify (Tang et al., 2026), SWE-PRM (Gandhi et al., 2025), DataPRM (Qiu et al., 2026), ThinkPRM (Khalifa et al., 2025) form a cluster where the WM is a learned *evaluator* of partial trajectories. As §17 develops critically, this is not the same object as a forward world model — PRMs are critics with execution grounding. They cannot roll out, cannot simulate counterfactuals. Survey hygiene argues for keeping the distinction.

0.12 10. Planning and Search with Code World Models

0.12.1 10.1 RAP (Hao et al., 2023)

RAP frames reasoning as MCTS in a self-consistent MDP where the same frozen LLM serves as both policy and transition model — a state is a textual configuration (blocks layout, intermediate variables, current fact), an action is a step proposed by the LLM, and the transition is obtained by re-prompting. Rewards combine action likelihood, state confidence (majority voting), self-evaluation (“Is this correct?”), and task heuristics. On Blocksworld 4-step, RAP@10 reaches 0.86 with LLaMA-33B, surpassing GPT-4+CoT’s 0.63 by 33% relative. The conceptual template — *repurpose the LLM as both policy and transition model under MCTS* — is what every later LLM-as-WM paper extends.

Tree of Thoughts (Yao et al., 2023), AlphaZero-like Tree Search for LLM Decoding (Feng et al., 2023), Tree Search for LM Agents (Koh et al., 2024), and Mastering Board Games by External/Internal Planning with LMs (Schultz et al., 2024) develop the search frame; the last gives the clearest contemporary recipe for learned tree-search with LLM-as-WM, straightforwardly transferable to code.

0.12.2 10.2 Execution-conditioned generation

Execution Guided Line-by-Line Code Generation (Lavon et al., 2025) uses classifier-free guidance to condition next-token prediction on candidate-runtime outcomes. Jupiter (Li et al., 2025a) formulates notebook state as MCTS nodes. REPL-Plan (Liu et al., 2024) reuses a REPL state pool across tasks. Substrate is well-developed for short-horizon code-gen; less so for long-horizon multi-file SWE.

0.13 11. JEPA, Dreamer, and the Latent-Action Gap

LeCun’s Joint Embedding Predictive Architecture (I-JEPA, (Assran et al., 2023)) predicts in embedding space rather than pixel space. The Dreamer family — Hafner et al.’s DreamerV1 (Hafner et al., 2019), DreamerV2 (Hafner et al., 2020), and DreamerV3 (Hafner et al., 2023), built around the Recurrent State-Space Model — has near-zero direct application to code. Two papers occupy the gap.

0.13.1 11.1 LLM-JEPA (Huang et al., 2025)

LLM-JEPA adds a joint-embedding predictive objective to standard NTP training, using (text, code) as the two JEPA views with the LLM’s last-layer last-token hidden state as encoder and a tied-weights [PRED] token as predictor. The loss is $\mathcal{L}_{\text{NTP}}(\text{text}) + \lambda \cdot d(\text{Pred}(\text{Enc}(\text{Text})), \text{Enc}(\text{Code}))$ with cosine distance. On Llama-3.2-1B fine-tuned on NL-RX-SYNTH the gain is 57.3 $\pm$ 71.5 (+14.2 absolute); on Spider, GSM8K, and HellaSwag the wins are smaller. The top-100 singular values of $\text{Enc}(\text{Text}) - \text{Enc}(\text{Code})$ collapse by orders of magnitude, indicating a low-rank text $\rightarrow$ code mapping. Whether this is “JEPA in the LeCun sense” or a regularizer on top of NTP is the live question raised in §17.

0.13.2 11.2 CoLA (Jia et al., 2025)

CoLA replaces the 128k-token action space of an LLM with a small learned latent-action code-book. Three modules: a VQ-VAE-style inverse-dynamics model that infers latent action a_t from $(x_1:t, x_{t+1})$; a language world model that inserts the chosen latent action into the LLM embedding stream and decodes the next token; and a policy $\pi(a_t | x_1:t)$ behavior-cloned from inverse-dynamics labels then RL-tuned. Action-level MCTS over the learned code-book (with a Double-DQN Q-function) reaches Math-500 68.2 vs 63.0 baseline MCTS-Q. CoLA is the corpus’s clearest “Dreamer-for-LLMs” instance — the action space is genuinely compressed, and rollout/search operate in that compressed space.

0.13.3 11.3 The gap

Despite CWM and dozens of LLM-as-world-model papers, *no public Dreamer/RSSM-style latent-imagination world model has been trained for SWE agents*. CWM rolls out in token space. CoLA is the closest concrete instance. UniZero (Pu et al., 2024) generalizes MuZero with transformers but is rarely instantiated on code. Genie (Bruce et al., 2024) gives the vision-side template. JEPA for RL (Kenneweg et al., 2025) extends the energy-based objective to RL.

Whether the gap matters is a live question developed critically in §17. The vision-domain pressure that motivated Dreamer’s RSSM design (pixel-space rollout cost) does not exist for code, where state is small and the simulator is available. The argument *for* latent imagination rests on inference-time speed and the action-space compression CoLA demonstrates, not on rollout cost per se.

0.14 12. Specialized Domains

Diffusion code models. DiffuCoder (Gong et al., 2025), Dream-Coder 7B (Xie et al., 2025). Iterative denoising naturally accommodates plan-then-refine generation.

Decompilation and cross-language. SK2Decompile (Tan et al., 2025), SALT4Decompile (Wang et al., 2025c). Translation as semantic-simulation task. EquiBench (Wei et al., 2025a) supplies the equivalence eval.

Hardware / RTL. VeriRL (Teng et al., 2025), ChipSeek (Chen et al., 2025b), VeriCoder (Wei et al., 2025c) form a cluster where the simulator is the world model. Hardware is an attractive domain because simulators are precise, fast, and deterministic — closer to Atari than to Python.

ARC and abstract synthesis. Executable World Models for ARC-AGI-3 (Rodionov, 2026) instantiates literal-WM-per-task: synthesize a Python world model verified against observations. SOAR (Pourcel et al., 2025) evolves programs over ARC. Darwin / Huxley Godel Machines ((Zhang et al., 2025a), (Wang et al., 2025b)) close the self-improvement loop.

0.15 13. Reasoning, Process Rewards, Memory

Long-CoT reasoning for code. o1-Coder (Zhang et al., 2024c) replicates o1 with MCTS+RL. R1-Code-Interpreter (Chen et al., 2025a) supplies the open SFT+RL recipe across 144 tasks. Scaling Test-Time Compute to Achieve IOI Gold Medal (Samadi et al., 2025) shows open-weight gpt-oss-120b matching closed reasoning models via inference-time scaling.

Long-CoT reasoning is *mental execution* — the chain-of-thought simulates the world model the network never explicitly trained. CWM-style explicit world modeling and R1-style reasoning are partial substitutes; whether they compose multiplicatively is open.

Memory. Episodic Memory is the Missing Piece for Long-Term LLM Agents (Pink et al., 2025) frames the gap. Memory as Action (Zhang et al., 2025b) treats memory operations as RL-learnable actions. RepoGraph (Ouyang et al., 2024) provides a durable repo-level dependency graph.

0.16 14. Verification, Probing, Safety

0.16.1 14.1 Formal verification: the leading edge

The verifier-grounded lineage is the only research direction in the corpus that does not rely on LLM self-report for correctness — the verifier provides ground truth.

ATLAS (Baksys et al., 2025) is the cleanest end-to-end pipeline in the corpus for scaling verified-code data. From TACO-verified (12.8k LeetCode-style problems with Python references and tests) ATLAS produces 2,751 verified Dafny programs decomposed into 19,385 training examples across six tasks (NL-to-Code, NL-to-Spec, Spec-to-Code, Spec-Repair, Impl-Repair, Proof-Infilling). Spec quality is filtered by three SMT-discharged lemma types — soundness, completeness-contradiction, and completeness-perturbation. Qwen-2.5-Coder-7B fine-tuned this way reaches DafnyBench Pass@1 of 55.8% (from 32.4%) and DafnySynthesis Pass@5 of 65.8% (from 15.8%, surpassing GPT-4’s 53.4). ATLAS does for verified Dafny what auto-formalization did for Lean theorems.

The cluster also includes Re:Form ((Yan et al., 2025), Dafny+RL), CLEVER ((Thakur et al., 2025), Lean), VeriStruct ((Sun et al., 2025), Verus), AutoRocq ((Tu et al., 2025), Rocq), and Semantic Equivalence Self-Play with Formal Verification ((Barone and Nok, 2026), Liquid Haskell).

CLEVER’s 71/161 end-to-end Lean result is the most sobering number in the corpus: frontier models, with Lean type-checker access for self-verification, still fail on >99% of HumanEval-derived problems requiring joint spec + implementation verification. Understanding Formal Reasoning Failures in LLMs as Abstract Interpreters (Mitchell et al., 2025) is the diagnostic: when asked to reason in the style of formal abstract interpretation over 22 SV-COMP programs, all frontier reasoning models make systematic errors in widening, fixpoint termination, and join operations.

0.16.2 14.2 Symbolic execution and LLMs

AutoBug (Li et al., 2025c), SESpec (Yang et al., 2025), LLM-Sym (Wang et al., 2024a), Loop Invariant Generation via Reasoning LLMs + SMT (Bharti et al., 2025) combine LLMs with concrete or symbolic engines. The unifying pattern: the LLM hypothesizes the world model; symbolic execution verifies or extends it.

0.16.3 14.3 Probing and mechanistic interpretability

Mechanistic Interpretability of Code Correctness via SAEs (Tahimic and Cheng, 2025) and On LLMs’ Internal Representation of Code Correctness (Ribeiro et al., 2025) ask what code LLMs actually represent. Findings: *partial, brittle* internal execution representations — a vindication of explicit trace pretraining.

0.16.4 14.4 Repair and debugging as world-model probing

Self-Debug (Chen et al., 2023), InspectCoder (Wang et al., 2025d), Agent That Debugs — Dynamic State-Guided Vulnerability Repair (Liu et al., 2025), Agentic Code Reasoning ((Ugare and Chandra, 2026), semi-formal execution-path reasoning without running code). Shared pattern: maintain a belief over program state, query the runtime to update the belief, act on the posterior — Bayesian world-modeling in everything but name.

0.16.5 14.5 Safety and malicious code

The Double Life of Code World Models (Sahoo, 2025) repurposes CWM-style trace predictions for malicious-behavior detection. CodeBreaker (Yan et al., 2024) is the offensive analogue. Concolic Execution + LLM for Zero-Day Malware Detection (Edwards and Eslamimehr, 2026) pairs path-prioritization with concrete execution.

0.17 15. Benchmarks and the Evaluation Gap

Benchmarks split cleanly by what they measure.

- Static code quality — HumanEval, MBPP, LiveCodeBench (Jain et al., 2024) measure code-LLM output without exercising the world-model claim.
- Execution reasoning — CRUXEval (Gu et al., 2024a), REval (Chen et al., 2024), CRUXEval-X (Xu et al., 2024), TraceEval (Li et al., 2026), PLSemanticsBench (Thimmaiah et al., 2025). The canonical WM evals.
- Semantic equivalence — EquiBench (Wei et al., 2025a), CodeARC (Wei et al., 2025b).
- Agentic — SWE-bench, SWE-Gym, WebArena (Zhou et al., 2023), Mind2Web (Deng et al., 2023), PyBench (Zhang et al., 2024a), OSWorld, WindowsAgentArena.

The eval gap. No widely adopted benchmark directly measures world-model fidelity by holding the policy fixed and varying the WM. Every measurement of WM capability is mediated through downstream task performance, so we cannot distinguish *the model has internalized program semantics from the model is exploiting trace-token shortcuts in the test distribution*. Demystifying Errors in LLM Reasoning Traces (Abdollahi et al., 2025) supplies the diagnostic: even DeepSeek-R1, o4-mini, Gemini 2.5 Flash, and Claude 4, when asked to simulate execution and explain their reasoning, produce traces with errors clustering into nine categories (Computation, Indexing, Control Flow, Skip Statements, Misvaluation of Native API, Hallucination, Input Misread, etc.). Models with 85–98% final-answer accuracy on output prediction produce traces with systematic errors throughout. The decoupling — high outcome accuracy, low process fidelity — is the signature of a system that has learned to predict outcomes without faithfully simulating dynamics.

0.18 16. Empirical Landscape

0.18.1 16.1 SWE-bench scoreboard

Protocols mix in SWE-bench reporting and are non-comparable across the obvious axes. We split the table into three protocol classes and ask readers to compare within, not across, classes. Training / Evaluation regime: *Frozen-model scaffold* = no model training (agent loop around a frozen closed model); *EG-LLM SFT* = supervised fine-tune on agent trajectories; *EG-LLM + RL* = RL with execution rewards; *EG-agent + learned simulator* = SFT+RL plus a learned execution surrogate; *EG-LLM + trace mid-train + RL* = the CWM lineage; *Scaffold-evolved* = recursive scaffold/agent self-improvement around a frozen closed model. WM-Arch is reserved for systems with explicit transition/rollout machinery (N1/N2/N3 from §3); no row in this table qualifies.

Table 16.1a — Single-sample resolution (pass@1 on Verified unless noted)

System	Base model	Bench	Pass@1	Regime	Source
SWE-agent	GPT-4 Turbo	full	12.5%	Frozen-model scaffold	(Yang et al., 2024)
AutoCodeRover	GPT-4	Lite	19.0%	Frozen-model scaffold	(Zhang et al., 2024b)
Agentless	GPT-4o	Lite	32.0%	Frozen-model scaffold	(Xia et al., 2024)
SWE-Gym	Qwen-2.5-Coder-32B	Verified	20.6%	EG-LLM SFT	(Pan et al., 2024)
Agent-RLVR (no RM)	Qwen-2.5-72B	Verified	22.4%	EG-LLM + RL	(Da et al., 2025)

System	Base model	Bench	Pass@1	Regime	Source
Long-Context Multi-Turn RL	Qwen-2.5-72B	Verified	39.0%	EG-LLM + RL	(Golubev et al., 2025)
SWE-RL (Llama3-SWE-RL-70B)	Llama-3.3-70B	Verified	41.0%	EG-LLM + RL	(Wei et al., 2025d)
Nanbeige SWE-World (RL)	Qwen-2.5-Coder-32B	Verified	55.0%	EG-agent + learned simulator	(Sun et al., 2026)
CWM	CWM-32B	Verified	<i>not directly reported</i> †	EG-LLM + trace mid-train + RL	(FAIR CodeGen Team et al., 2025)

† The CWM paper does not report a pure pass@1 number under the standard protocol; the headline 65.8% in Table 16.1b is best@16 with verifier reranking. Inferred pass@1 from the paper’s reported figures is approximately 53–55%, but this is *our estimate, not a directly reported number*, and we do not place it in the main row.

Table 16.1b — Best-of-k with verifier reranking / TTS

System	Base	Bench	Score	Sample budget	Source
Agent-RLVR + RM rerank	Qwen-2.5-72B	Verified	27.8%	rerank over multi-sample	(Da et al., 2025)
Nanbeige SWE-World (TTS@8)	Qwen-2.5-Coder-32B	Verified	68.2%	best@8	(Sun et al., 2026)
CWM (TTS)	CWM-32B	Verified	65.8%	best@16 over 40 reranked	(FAIR CodeGen Team et al., 2025)

Table 16.1c — Scaffold-evolution and self-play around closed-model executors

System	Backbone executor	Bench	Score (start ? end)	Note	Source
Darwin Godel Machine	Claude-3.5-Sonnet	Verified	20% ? 50%	Scaffold evolves over 80 iterations	(Zhang et al., 2025a)
Huxley Godel Machine	GPT-5-mini	Verified (500)	— ? 61.4%	Scaffold optimized for Verified-60	(Wang et al., 2025b)
SICA	Claude-3.5-Sonnet + o3-mini	Verified (subset)	17% ? 53%	Scaffold evolves; not a model-training method	(Robeyns et al., 2025)

Citations that quote “CWM 65.8% SWE-bench Verified” without disclosing the best@16/TTS protocol should be read as misreporting, not as direct pass@1 numbers.

What can and cannot be compared. Within Table 16.1a the comparable axis is pass@1. SWE-RL’s 41.0%, Long-Context-MT-RL’s 39.0%, and Nanbeige SWE-World’s 55.0% are roughly comparable as “open-weight WM-trained pass@1 on Verified.” Table 16.1c sits in a different category — these are scaffold-evolution methods over closed models, and treating them as evidence for “world-model training works” conflates scaffold search with model improvement. Reading them inside Table 16.1a alongside model-trained pass@1 numbers is apples-to-oranges.

Two empirical claims, separately defensible. (i) On *pass@1*, open-weight WM-trained 32B systems (Nanbeige 55.0%, Long-Context-MT-RL 39.0%, SWE-RL 41.0%) outperform frozen open-weight baselines (Qwen-2.5-Coder-32B raw 6.2%) by 30–50 absolute points; this gain is the strongest case for training on agent trajectories with execution feedback. (ii) On *best@k with verifier reranking* (Table 16.1b), the same models reach 65–68%, but the additional 13–15 points are attributable to TTS reranking infrastructure, not to the WM training. Conflating (i) and (ii) by quoting CWM’s 65.8% alongside SWE-RL’s 41.0% as both “world-model-trained pass@1 SOTA” is exactly the misreporting the protocol split is designed to prevent.

0.18.2 16.2 Trace pretraining gains on execution-reasoning

Paper	Backbone	Baseline	After trace pretrain/FT	Delta	Benchmark
TRACED	UnixCoder	—	+12.4% rel branch-coverage; +25.2% rel variable-value	rel	CodeNet exec
NExT	PaLM 2-L	23.2	49.3	+26.1 abs	MBPP-R
NExT	PaLM 2-L	32.2	42.5	+10.3 abs	HumanEvalFix-Plus
SemCoder (1.3B)	DS-Coder 1.3B	base	63.6 / 63.9	+23 abs	CRUXEval-I / O
“What I cannot execute”	Llama-3.1-8B	37.8%	~80%	+42 abs	CRUXEval-O
Do Code Semantics Help?	DSCoder, Llama-3, Gemma-2	various	⌈ a few abs; some regressions	mixed	comprehensive

For under-trained ⌈8B open-weights, trace pretraining delivers +15 to +42 absolute on CRUXEval-O. The “Do Code Semantics Help?” ablation (Wang et al., 2025a) is the disconfirming evidence: across multiple backbones and five trace representations, no single representation consistently outperforms others, and several downstream tasks regress under trace augmentation. The gain shrinks rapidly with base-model quality. Trace pretraining is a remedial intervention for weak code models; whether frontier models still benefit is unsettled.

0.18.3 16.3 Web/OS agents from WMs

Paper	Benchmark	Base	With WM	Delta	Mechanism
WebDreamer (GPT-4o WM)	VisualWebArena	7.6%	23.6%	+34.1% rel (⌈+6 abs)	LLM-as-WM + MPC
WebDreamer	Online-Mind2Web	26.0%	37.0%	+42.3% rel (⌈+11 abs)	same
Dreamer-7B (trained WM)	VisualWebArena	base	+4.7 abs	—	trained WM

Paper	Benchmark	Base	With WM	Delta	Mechanism
WMA (Chae et al., 2024)	WebArena	base	action-selection 52.7% 70%	—	trained transition WM
Dyna-Think DDT (32B)	OSWorld	RFT~28%	43.1%	7+15 abs	Dyna-Q + WM head same
Dyna-Think DDT	WindowsAgentArena	28.4%	34.9%	+6.5 abs	

WM gains on web/OS are real but quantitatively small (7+5–10 absolute task success rate on most benchmarks), and partly confounded with the extra synthetic data the WM generates. DyMo’s 90%+ state-prediction accuracy versus 72.8% task success rate exemplifies the decoupling: WM heads can be accurate without the agent being accurate.

0.18.4 16.4 Formal verification vs LLM-only

System	Language	Benchmark	Baseline	With system	Source
ATLAS	Dafny	DafnyBench Pass@1	32.4%	55.8%	(Baksys et al., 2025)
ATLAS	Dafny	DafnySynthesis Pass@5	15.8%	65.8% (>GPT-4 53.4)	(Baksys et al., 2025)
CLEVER	Lean 4	161 problems end-to-end	best frontier	71/161	(Thakur et al., 2025)
VeriStruct	Verus / Rust	11 modules	—	99.2% (128/129 fns)	(Sun et al., 2025)
AutoRocq	Rocq	math + verif lemmas	5 baselines	48.0% math / 30.9% verif	(Tu et al., 2025)
Semantic Equiv Self-Play	Liquid Haskell	EquiBench	base	+13.3 pp	(Barone and Nok, 2026)

Verified codegen has the steepest training-data sensitivity in the survey: small synthetic datasets (2.7K verified Dafny programs in ATLAS) produce +25–50 absolute gains because LLM-only baselines start near zero. CLEVER’s 71/161 shows that without explicit data/scaffolding, frontier models cannot reliably produce verified code. VeriStruct shows that on curated targets near-perfect is reachable.

0.18.5 16.5 Reasoning-model competitive programming

Model	Codeforces	IOI / ICPC
gpt-4o	808 (11th pct)	—
o1-preview	1258 (62nd pct)	—
o1	1673 (89th pct)	—
o1-ioi	1807 (~93rd pct)	IOI 2024 49th pct live
o3	elite-human-class	IOI 2024 gold
gpt-oss-120b + GenCluster	—	IOI 2025 gold (open-weight)
Gemini 2.5 Pro Exp	—	ICPC-Eval Pass@1 22.0%
DeepSeek-R1	—	ICPC-Eval Pass@1 14.4%
Claude 3.7 Sonnet	—	ICPC-Eval Pass@1 11.8%
GPT-4o	—	ICPC-Eval Pass@1 5.9%

Codeforces +999 rating points in 14 months on the o-series. The same line shows that frontier reasoning models gain almost everything from RL scaling, not from domain-specialized scaff-

folds — which complicates the case that “trace-style world models” are doing causal work for frontier systems.

0.19 17. Critical Perspectives

This section names where the field overclaims, where the consensus is fragile, and where vocabulary is doing more work than evidence. We develop seven theses.

0.19.1 17.1 The “world model” label has become marketing for any code LLM trained on something other than raw source

This is a terminology argument, not a contradiction with §3. §3 deliberately admits “implicit WMs in token policies” as a fourth flavor — including CWM, TRACED, SemCoder under that flavor — because that *is* how the field currently uses the term. The claim of §17.1 is that this permissive definition is what allows the rhetorical sleight of hand, and that a stricter definition would be more useful going forward.

Concretely: read CWM (FAIR CodeGen Team et al., 2025) carefully and the architecture it ships is a 32B decoder-only Transformer with GQA, sliding-window blocks, RoPE, AdamW — a Llama-class model. What earns it the “world model” badge is the mid-training datamix: 5T tokens of Python observation-action traces plus ForagerAgent SWE trajectories. There is no separate dynamics head, no inverse model, no recurrent latent. The same observation lands on LLM-JEPA, DyMo, and most “world model” papers from 2024–2026: the artifact is a standard LLM with an enriched objective.

A useful purity test: *can we ablate the supposed world-modeling component without changing the architecture?* If yes, the system is a trace-trained LLM. If no, there is a genuine architectural commitment. By that test, CoLA (Jia et al., 2025) and the Dreamer-for-LLMs gestures pass. CWM, SemCoder, NExT, and most of the “explicit WM” cluster fail. We propose that “world model” be reserved for systems with the architectural commitment, and that *execution-grounded code LLM* serve for the rest. The §3 permissive definition can stay as descriptive — what the field currently calls a code WM — but a normative tightening is warranted, and the survey from §17 onward uses the stricter sense.

0.19.2 17.2 Trace pretraining has a causal-isolation problem the surface numbers obscure

“Do Code Semantics Help?” (Wang et al., 2025a) is the most damaging paper for the prevailing optimism. It runs a comprehensive ablation across DeepSeek-Coder, LLaMA-3, and Gemma-2 with five representations (Scratchpad, NExT, CodeExecutor, Concise, SemCoder) on program repair, code synthesis, BigCodeBench, LiveCodeBench, and CRUXEval. Its headline: integrating trace-based semantic information into SFT *cannot significantly enhance* code-generation ability. In 7 of 9 synthesis settings the no-trace baseline wins or ties. At inference, in 36 of 56 test-scaling configurations, trace prompts hurt.

“What I cannot execute, I do not understand” (Armengol-Estapé et al., 2025) is gentler — Execution Tuning reaches ~80% CRUXEval-O — but the same paper’s downstream evaluations on HumanEval, MBPP, and GSM8K show *negligible* gains from trace data in the SFT mix.

The strongest counterexample to “barely transfers” is NExT (Ni et al., 2024), which pushes PaLM 2-L on Mbpp-R from 23.2% to 40.8% *with traces removed at test time* — a +17.6 absolute gain on the no-trace-at-inference setting, which is exactly the transfer pattern the skeptical reading says shouldn’t happen. The honest version of the thesis is therefore narrower: trace pretraining transfers to *program-repair* tasks where the model needs to reason about a buggy program’s behavior (where NExT and TraceFixer-style work shines), and barely transfers to *fresh code synthesis* on benchmarks like HumanEval and MBPP that are closer to the model’s prior. The “Do Code Semantics Help?” disconfirmation is on synthesis; NExT’s transfer is on repair. Both can be true.

The honest reading: trace pretraining helps execution prediction (the thing trained on), transfers to runtime-reasoning-heavy downstream tasks like repair, and barely transfers to fresh code

synthesis, exactly the pattern you would expect if dense execution supervision is teaching a runtime-tracking skill rather than a generative model of program intent.

CWM’s 65.8% on SWE-bench is the apparent counterexample, but the Meta team reports it after trace mid-training *and* 3M ForagerAgent SWE trajectories *and* multi-task RL with verifiable rewards — three interventions stacked. The “world model” component is not causally isolated from the SWE-trajectory and RL components. Without an ablation that removes trace mid-training while holding ForagerAgent and RL fixed, the headline is unfalsifiable. The field has agreed to call CWM a world-model success because the name is on the model card, not because the experimental design demonstrates it.

0.19.3 17.3 The Dreamer-for-code gap may be a non-problem

The conventional framing treats the absence of latent-imagination world models for SWE as the field’s largest architectural gap. The empirical record argues the opposite. CWM in token space reaches 65.8% SWE-bench. CoLA produces respectable but not field-shifting results, and even there the WM is fine-tuned on top of a standard LLM rather than replacing it.

Vision world models needed latent rollouts because pixel-space rollouts were too expensive — a frame is $\sim 10^6$ dimensions, dynamics are partially observed. Program execution is the opposite: a Python frame is small, dynamics are observable, and the simulator (CPython) is available for free at training time. The pressure that drove Dreamer’s RSSM design does not exist for code.

Debugging Code World Models (Rahmani, 2026) shows CWM’s long-horizon failures are dominated by *action hallucination*, not state-propagation error — under teacher forcing CWM tracks state correctly for 128 steps. A latent rollout would compress states but not fix the action policy, which is the actual bottleneck.

The counterargument is fair: latent rollouts permit faster planning at inference, and for multi-agent or population-scale search the speedup is asymptotically meaningful. But the survey should retire “single largest architectural gap” framing and replace it with “an interesting open question whose payoff is not yet demonstrated.”

0.19.4 17.4 PRMs are critics, not world models

Process reward models — ExecVerify (Tang et al., 2026), SWE-PRM (Gandhi et al., 2025), ThinkPRM (Khalifa et al., 2025), FunPRM (Zhang et al., 2026), DataPRM (Qiu et al., 2026) — are often grouped under the world-modeling umbrella. This is wrong in a way that matters. A world model is, by every definition the survey uses, a *forward* predictor of `(state, action) → next_state`. A PRM is a *backward-looking evaluator*: given a partial trajectory, score it. In classical model-based RL these are different objects — Dreamer has both a world model (RSSM) and a critic (value function).

PRMs cannot roll out. Cannot simulate counterfactuals. Cannot be used by a planner that wants to score a hypothesized future. Conflating them dilutes the world-model concept until it means “any neural network trained on execution-related signals” — at which point the term is useless. Vocabulary discipline is cheap and the field would benefit from it.

The same critique applies, less severely, to verifier-grounded systems: a Lean proof checker is not a world model, it is a deterministic verifier of a candidate output. Calling it “the world model” when ATLAS, Re:Form, or AutoRocq use it makes for a tidy survey arc but blurs the actual computational structure.

0.19.5 17.5 “General Agents Contain World Models” is much weaker than its title suggests

Richens et al. (Richens et al., 2025) prove that any goal-conditioned policy satisfying a regret bound δ for sufficiently deep composite goals (depth $n \geq 1$) must encode an extractable approximation of the transition function with bounded error. Genuine and elegant.

But read the assumptions: fully observed environment, finite communicating stationary controlled MDP, goal-conditioned policy satisfying a regret bound for a specific class of LTL composite goals of depth n . Theorem 2 of the same paper explicitly shows that for myopic agents (depth-1 goals), *no world model is needed*. Real SWE agents are myopic-ish over short turns and

approximately competent over longer ones; their environments are partially observed (rarely full filesystem state); they violate stationarity (the repository changes under their actions); and their regret bound for arbitrary composite goals is unknown and almost certainly not satisfied. The authors caveat this in §6 (“Limitations”) of their paper.

The theorem is a beautiful existence proof for an idealized agent class. It is *not* an empirical statement that SWE coding agents have learned world models, and it provides no guidance about the fidelity of any world model they may have learned.

0.19.6 17.6 The verifier-grounded lineage is the actual leading edge

ATLAS, Re:Form, CLEVER, VeriStruct, AutoRocq, and the Liquid Haskell self-play paper (Barone and Nok, 2026) share a property no LLM-only system possesses: code whose correctness is *machine-checked* against a formal specification. Compare to the SWE-bench paradigm, where “correctness” means “hidden unit tests pass” — a weaker guarantee, since unit tests cover specific inputs and the system can pass them while being wrong on adjacent inputs.

The abstract-interpreter paper (Mitchell et al., 2025) is the diagnostic: when frontier reasoning LLMs are asked to reason in the style of formal abstract interpretation over 22 SV-COMP programs, they make systematic errors in widening, fixpoint termination, control-flow propagation, and meet/join operations. They generated unsound invariants on programs as small as `count_by_2.c`. If LLMs cannot reliably perform interval-domain abstract interpretation on toy C programs, claims that they have learned faithful internal world models of program semantics are doing a lot of inferential work.

The verifier-grounded line is the only research direction that does not rely on LLM self-report for correctness. Scoped carefully, it leads on *synthesis-from-spec* (ATLAS reaches 65.8% Pass@5 on DafnySynthesis at 7B, surpassing GPT-4) and on *near-perfect verified completion of curated targets* (VeriStruct 99.2% on 11 Verus modules). It does *not* lead on end-to-end verified codegen from natural language: CLEVER reports 7/161 on Lean problems requiring joint spec + implementation verification. The future of *correct* code is plausibly hybrid — neural proposal, symbolic verification, with the verifier providing ground truth that learned world models cannot — but the hybrid is not yet a finished pillar. We catalog the verifier line in §14.1 alongside symbolic execution and abstract-interpretation probing; readers who accept this thesis should weight that subsection more heavily than its sequence position suggests.

0.19.7 17.7 The evaluation gap is the structural reason the field looks confused

Across CRUXEval, REval, CRUXEval-X, PLSemanticsBench, TraceEval, and EquiBench, no benchmark holds policy fixed and varies world-model quality. Every measurement of “world-modeling capability” is mediated through downstream task performance, so we cannot distinguish *the model has internalized program semantics* from *the model is exploiting trace-token shortcuts in the test distribution*.

“Demystifying Errors in LLM Reasoning Traces” (Abdollahi et al., 2025) is the diagnostic: even DeepSeek-R1, o4-mini, Gemini 2.5 Flash, and Claude 4, when asked to simulate execution, produce traces with errors in nine systematic categories. Models with 85–98% final-answer accuracy on output prediction produce traces with systematic errors throughout — high outcome accuracy, low process fidelity, the signature of a system that predicts outcomes without faithfully simulating dynamics.

Self-repair literature exhibits the same pathology: Olausson et al. (Olausson et al., 2023) showed that GPT-4 self-repair on APPS and HumanEval, normalized by compute, often performs *worse* than i.i.d. resampling. The bottleneck is the model’s feedback quality, not its repair capability — human-written feedback boosts repair success by 1.58×. The model can generate code, can sometimes recognize bugs, but cannot reliably simulate why its code is wrong — which is exactly what a faithful world model would let it do. The empirical bound on LLM self-repair is, in effect, an empirical bound on the fidelity of the implicit world model the LLM is running. Calling that world model internal is fine; calling it good is not.

Until benchmarks measure process fidelity independently of outcome, “is this system actually building a world model?” remains scientifically undecidable.

0.20 18. Open Problems

The critical perspectives of §17 reshape the conventional open-problems list. We propose six problems where the literature is thinnest *and* the upside is largest.

1. Causal isolation of trace-pretraining contributions. Every claim of the form “this WM-trained model achieves X” should be paired with an ablation removing the WM component while holding training data and RL fixed. CWM in particular needs this. Without it, the headline numbers underdetermine whether the WM did the work.
2. World-model fidelity as a first-class metric. §15’s eval gap is concrete: build a benchmark where holding policy fixed and varying WM quality causes measurable variation in planning quality, independent of downstream task. This benchmark would clarify the field more than any single new model.
3. Hybrid neural-symbolic systems. §17.6 argues the verifier-grounded line is the leading edge. The natural integration is *neural proposal, symbolic verification*, with the verifier providing gradient-free correctness signal and the neural component providing proposals at scale. Differentiable surrogates of symbolic verifiers (Lean / Dafny / Rocq) that pass verifier-style gradients during training are open.
4. Multi-modal WMs for coding. GUI agents need pixel-level WMs (Neural Computers, (Zhuge et al., 2026), is a first attempt). Tying pixel WMs to code-state WMs through a shared latent is essentially unsolved.
5. Long-horizon credit assignment with execution-grounded rewards. PRMs (§9.3) are early, and §17.4 argues they should be conceptually separated from world models. The right structure for rewarding an agent across hundreds of execution-grounded steps is a live question.
6. World models of the developer, not just the program. All current WMs model the *machine*. Few model the *developer intent* with comparable fidelity. ATLAS and Re:Form gesture in this direction by treating the spec as the WM. A full developer-intent WM would close the agentic loop.

We do not list “Dreamer-for-SWE-agents” as the field’s largest gap. §17.3 argues the pressure motivating that direction in vision does not transfer to code. It remains an interesting research question, not the highest-leverage one.

Three under-explored representations. The §5.3 taxonomy table flags three classes of representation, mature in vision-WM, that have no code-WM exemplar:

- *Global Latent Vector (Dreamer-style RSSM)* — discussed and contested above. The Hafner et al. DreamerV1–V3 line proves the design space exists; whether it pays off for code is the question §17.3 leaves open.
- *Spatial / Structural Grid* — the analog of OccWorld / BEV for code would be a learned predictive grid over AST nodes, call-graph edges, or CFG states. RepoGraph (Ouyang et al., 2024) shows the static version is useful as agent state; the predictive version is unexplored.
- *Decomposed Object / Slot (object-centric WMs)* — the analog for code would model variables, scopes, or classes as discrete persistent slots whose state propagates independently. No paper in the corpus instantiates this, despite obvious mappings (each variable is an object, each frame is a scene).

The object-centric and structural-grid gaps look more genuine than the Dreamer one, in our reading, because they exploit structure that code *already has* (objects = variables, grids = AST/CFG) rather than borrowing pressure from a domain (vision) where the structural assumptions differ.

0.21 19. Conclusion

Across the literature surveyed here, a single trajectory is visible: from neural execution (modeling the machine), through trace pretraining (modeling execution implicitly), to CWM and its descendants (modeling execution explicitly with a named artifact), to agentic SWE and RL (modeling the environment), to JEPA and latent-action models (modeling in compressed space), and on toward formal verification, probing, and safety (modeling reliably). What was a scattered set of insights in 2014 has by 2026 cohered into a recognizable research program with a recognizable artifact — the code world model.

The remaining work splits into two halves. The first is empirical: close the eval gap, isolate the causal contribution of WM-training, build hybrid neural-symbolic systems whose correctness is verifier-checkable rather than test-checkable. The second is rhetorical: hold the term “world model” to a strict definition so the literature can distinguish architectural commitments from training-data choices, and resist the temptation to oversell extractability theorems and latent-imagination analogies whose premises do not transfer to code.

The opportunity is large precisely because the framework is now clear enough to identify what is missing. The work to do is the work this survey has tried to make visible.

0.22 Appendix A · References

All citations in this survey use arxiv identifiers inline (e.g. [[@arxiv2510_02387](#)]). Each cited identifier resolves to one entry below.

- (Zaremba and Sutskever, 2014) — Learning to Execute. [https://arxiv.org/abs/\(Zaremba and Sutskever, 2014\)](https://arxiv.org/abs/(Zaremba and Sutskever, 2014))
- (Reed and de Freitas, 2015) — Neural Programmer-Interpreters. [https://arxiv.org/abs/\(Reed and de Freitas, 2015\)](https://arxiv.org/abs/(Reed and de Freitas, 2015))
- (Wang et al., 2017) — Dynamic Neural Program Embedding for Program Repair. [https://arxiv.org/abs/\(Wang et al., 2017\)](https://arxiv.org/abs/(Wang et al., 2017))
- (Ha and Schmidhuber, 2018) — World Models. [https://arxiv.org/abs/\(Ha and Schmidhuber, 2018\)](https://arxiv.org/abs/(Ha and Schmidhuber, 2018))
- (Shi et al., 2019) — Learning Execution through Neural Code Fusion. [https://arxiv.org/abs/\(Shi et al., 2019\)](https://arxiv.org/abs/(Shi et al., 2019))
- (Hafner et al., 2019) — Dream to Control - Learning Behaviors by Latent Imagination (DreamerV1). [https://arxiv.org/abs/\(Hafner et al., 2019\)](https://arxiv.org/abs/(Hafner et al., 2019))
- (Hafner et al., 2020) — Mastering Atari with Discrete World Models (DreamerV2). [https://arxiv.org/abs/\(Hafner et al., 2020\)](https://arxiv.org/abs/(Hafner et al., 2020))
- (Bieber et al., 2020) — Learning to Execute Programs with Instruction Pointer Attention Graph Neural Networks. [https://arxiv.org/abs/\(Bieber et al., 2020\)](https://arxiv.org/abs/(Bieber et al., 2020))
- (Chen et al., 2021) — Evaluating Large Language Models Trained on Code. [https://arxiv.org/abs/\(Chen et al., 2021\)](https://arxiv.org/abs/(Chen et al., 2021))
- (Austin et al., 2021) — Program Synthesis with Large Language Models. [https://arxiv.org/abs/\(Austin et al., 2021\)](https://arxiv.org/abs/(Austin et al., 2021))
- (Nye et al., 2021) — Show Your Work - Scratchpads for Intermediate Computation with Language Models. [https://arxiv.org/abs/\(Nye et al., 2021\)](https://arxiv.org/abs/(Nye et al., 2021))
- (Le et al., 2022) — CodeRL - Mastering Code Generation through Pretrained Models and Deep Reinforcement Learning. [https://arxiv.org/abs/\(Le et al., 2022\)](https://arxiv.org/abs/(Le et al., 2022))
- (Hafner et al., 2023) — Mastering Diverse Domains through World Models (DreamerV3). [https://arxiv.org/abs/\(Hafner et al., 2023\)](https://arxiv.org/abs/(Hafner et al., 2023))
- (Assran et al., 2023) — I-JEPA - Self-Supervised Learning from Images with a Joint-Embedding Predictive Architecture. [https://arxiv.org/abs/\(Assran et al., 2023\)](https://arxiv.org/abs/(Assran et al., 2023))
- (Ni et al., 2023) — LEVER - Learning to Verify Language-to-Code Generation with Execution. [https://arxiv.org/abs/\(Ni et al., 2023\)](https://arxiv.org/abs/(Ni et al., 2023))
- (Shinn et al., 2023) — Reflexion - Language Agents with Verbal Reinforcement Learning. [https://arxiv.org/abs/\(Shinn et al., 2023\)](https://arxiv.org/abs/(Shinn et al., 2023))

- (Chen et al., 2023) — Teaching Large Language Models to Self-Debug. [https://arxiv.org/abs/\(Chen et al., 2023\)](https://arxiv.org/abs/(Chen et al., 2023))
- (Liu et al., 2023) — Code Execution with Pre-trained Language Models. [https://arxiv.org/abs/\(Liu et al., 2023\)](https://arxiv.org/abs/(Liu et al., 2023))
- (Yao et al., 2023) — Tree of Thoughts - Deliberate Problem Solving with Large Language Models. [https://arxiv.org/abs/\(Yao et al., 2023\)](https://arxiv.org/abs/(Yao et al., 2023))
- (Hao et al., 2023) — Reasoning with Language Model is Planning with World Model - RAP. [https://arxiv.org/abs/\(Hao et al., 2023\)](https://arxiv.org/abs/(Hao et al., 2023))
- (Deng et al., 2023) — Mind2Web - Towards a Generalist Agent for the Web. [https://arxiv.org/abs/\(Deng et al., 2023\)](https://arxiv.org/abs/(Deng et al., 2023))
- (Ding et al., 2023) — TRACED - Execution-aware Pre-training for Source Code. [https://arxiv.org/abs/\(Ding et al., 2023\)](https://arxiv.org/abs/(Ding et al., 2023))
- (Olausson et al., 2023) — Is Self-Repair a Silver Bullet for Code Generation. [https://arxiv.org/abs/\(Olausson et al., 2023\)](https://arxiv.org/abs/(Olausson et al., 2023))
- (Zhou et al., 2023) — WebArena - A Realistic Web Environment for Building Autonomous Agents. [https://arxiv.org/abs/\(Zhou et al., 2023\)](https://arxiv.org/abs/(Zhou et al., 2023))
- (Feng et al., 2023) — AlphaZero-like Tree-Search can Guide Large Language Model Decoding and Training. [https://arxiv.org/abs/\(Feng et al., 2023\)](https://arxiv.org/abs/(Feng et al., 2023))
- (Jimenez et al., 2023) — SWE-bench - Can Language Models Resolve Real-World GitHub Issues. [https://arxiv.org/abs/\(Jimenez et al., 2023\)](https://arxiv.org/abs/(Jimenez et al., 2023))
- (Gu et al., 2024a) — CRUXEval - A Benchmark for Code Reasoning, Understanding and Execution. [https://arxiv.org/abs/\(Gu et al., 2024a\)](https://arxiv.org/abs/(Gu et al., 2024a))
- (Wang et al., 2024b) — Executable Code Actions Elicit Better LLM Agents. [https://arxiv.org/abs/\(Wang et al., 2024b\)](https://arxiv.org/abs/(Wang et al., 2024b))
- (Bruce et al., 2024) — Genie - Generative Interactive Environments. [https://arxiv.org/abs/\(Bruce et al., 2024\)](https://arxiv.org/abs/(Bruce et al., 2024))
- (Jain et al., 2024) — LiveCodeBench - Holistic and Contamination Free Evaluation of Large Language Models for Code. [https://arxiv.org/abs/\(Jain et al., 2024\)](https://arxiv.org/abs/(Jain et al., 2024))
- (Chen et al., 2024) — Evaluating Large Language Models with Runtime Behavior of Program Execution. [https://arxiv.org/abs/\(Chen et al., 2024\)](https://arxiv.org/abs/(Chen et al., 2024))
- (Zhang et al., 2024b) — AutoCodeRover - Autonomous Program Improvement. [https://arxiv.org/abs/\(Zhang et al., 2024b\)](https://arxiv.org/abs/(Zhang et al., 2024b))
- (Ni et al., 2024) — NExT - Teaching Large Language Models to Reason about Code Execution. [https://arxiv.org/abs/\(Ni et al., 2024\)](https://arxiv.org/abs/(Ni et al., 2024))
- (Dainese et al., 2024) — Generating Code World Models with Large Language Models Guided by Monte Carlo Tree Search. [https://arxiv.org/abs/\(Dainese et al., 2024\)](https://arxiv.org/abs/(Dainese et al., 2024))
- (Yang et al., 2024) — SWE-agent - Agent-Computer Interfaces Enable Automated Software Engineering. [https://arxiv.org/abs/\(Yang et al., 2024\)](https://arxiv.org/abs/(Yang et al., 2024))
- (Jiang et al., 2024) — A Survey on Large Language Models for Code Generation. [https://arxiv.org/abs/\(Jiang et al., 2024\)](https://arxiv.org/abs/(Jiang et al., 2024))
- (Ding et al., 2024b) — SemCoder - Training Code Language Models with Comprehensive Semantics Reasoning. [https://arxiv.org/abs/\(Ding et al., 2024b\)](https://arxiv.org/abs/(Ding et al., 2024b))
- (Yan et al., 2024) — An LLM-Assisted Easy-to-Trigger Backdoor Attack on Code Completion Models. [https://arxiv.org/abs/\(Yan et al., 2024\)](https://arxiv.org/abs/(Yan et al., 2024))
- (Pu et al., 2024) — UniZero - Generalized and Efficient Planning with Scalable Latent World Models. [https://arxiv.org/abs/\(Pu et al., 2024\)](https://arxiv.org/abs/(Pu et al., 2024))
- (Koh et al., 2024) — Tree Search for Language Model Agents. [https://arxiv.org/abs/\(Koh et al., 2024\)](https://arxiv.org/abs/(Koh et al., 2024))
- (Xia et al., 2024) — Agentless - Demystifying LLM-based Software Engineering Agents. [https://arxiv.org/abs/\(Xia et al., 2024\)](https://arxiv.org/abs/(Xia et al., 2024))
- (Zhang et al., 2024a) — PyBench - Evaluating LLM Agent on Various Real-World Coding Tasks. [https://arxiv.org/abs/\(Zhang et al., 2024a\)](https://arxiv.org/abs/(Zhang et al., 2024a))
- (Xu et al., 2024) — CRUXEval-X - A Benchmark for Multilingual Code Reasoning Understanding and Execution. [https://arxiv.org/abs/\(Xu et al., 2024\)](https://arxiv.org/abs/(Xu et al., 2024))
- (Wang et al., 2024a) — Python Symbolic Execution with LLM-powered Code Generation. [https://arxiv.org/abs/\(Wang et al., 2024a\)](https://arxiv.org/abs/(Wang et al., 2024a))
- (Gehring et al., 2024) — RLEF - Grounding Code LLMs in Execution Feedback with Reinforcement Learning. [https://arxiv.org/abs/\(Gehring et al., 2024\)](https://arxiv.org/abs/(Gehring et al., 2024))

- (Chae et al., 2024) — Web Agents with World Models - Learning and Leveraging Environment Dynamics in Web Navigation. [https://arxiv.org/abs/\(Chae et al., 2024\)](https://arxiv.org/abs/(Chae et al., 2024))
- (Ouyang et al., 2024) — RepoGraph - Enhancing AI Software Engineering with Repository-level Code Graph. [https://arxiv.org/abs/\(Ouyang et al., 2024\)](https://arxiv.org/abs/(Ouyang et al., 2024))
- (Gu et al., 2024b) — Is Your LLM Secretly a World Model of the Internet - WebDreamer. [https://arxiv.org/abs/\(Gu et al., 2024b\)](https://arxiv.org/abs/(Gu et al., 2024b))
- (Liu et al., 2024) — REPL-Plan - Interactive and Expressive Code-Augmented Planning with Large Language Models. [https://arxiv.org/abs/\(Liu et al., 2024\)](https://arxiv.org/abs/(Liu et al., 2024))
- (Ding et al., 2024a) — Understanding World or Predicting Future - A Comprehensive Survey of World Models. [https://arxiv.org/abs/\(Ding et al., 2024a\)](https://arxiv.org/abs/(Ding et al., 2024a))
- (Zhang et al., 2024c) — o1-Coder - an o1 Replication for Coding. [https://arxiv.org/abs/\(Zhang et al., 2024c\)](https://arxiv.org/abs/(Zhang et al., 2024c))
- (Schultz et al., 2024) — Mastering Board Games by External and Internal Planning with Language Models. [https://arxiv.org/abs/\(Schultz et al., 2024\)](https://arxiv.org/abs/(Schultz et al., 2024))
- (Pan et al., 2024) — Training Software Engineering Agents and Verifiers with SWE-Gym. [https://arxiv.org/abs/\(Pan et al., 2024\)](https://arxiv.org/abs/(Pan et al., 2024))
- (DeepSeek-AI et al., 2025) — DeepSeek-R1 incentivizes reasoning in LLMs through reinforcement learning. [https://arxiv.org/abs/\(DeepSeek-AI et al., 2025\)](https://arxiv.org/abs/(DeepSeek-AI et al., 2025))
- (Pink et al., 2025) — Position - Episodic Memory is the Missing Piece for Long-Term LLM Agents. [https://arxiv.org/abs/\(Pink et al., 2025\)](https://arxiv.org/abs/(Pink et al., 2025))
- (Wei et al., 2025a) — EquiBench - Benchmarking LLMs' Reasoning about Program Semantics via Equivalence Checking. [https://arxiv.org/abs/\(Wei et al., 2025a\)](https://arxiv.org/abs/(Wei et al., 2025a))
- (Wei et al., 2025d) — SWE-RL - Advancing LLM Reasoning via Reinforcement Learning on Open Software Evolution. [https://arxiv.org/abs/\(Wei et al., 2025d\)](https://arxiv.org/abs/(Wei et al., 2025d))
- (Armengol-Estapé et al., 2025) — What I cannot execute, I do not understand - Training and Evaluating LLMs on Program Execution Traces. [https://arxiv.org/abs/\(Armengol-Estapé et al., 2025\)](https://arxiv.org/abs/(Armengol-Estapé et al., 2025))
- (Mitchell et al., 2025) — Understanding Formal Reasoning Failures in LLMs as Abstract Interpreters. [https://arxiv.org/abs/\(Mitchell et al., 2025\)](https://arxiv.org/abs/(Mitchell et al., 2025))
- (Jia et al., 2025) — CoLA - Controlling Large Language Models with Latent Actions. [https://arxiv.org/abs/\(Jia et al., 2025\)](https://arxiv.org/abs/(Jia et al., 2025))
- (Wei et al., 2025b) — CodeARC - Benchmarking Reasoning Capabilities of LLM Agents for Inductive Program Synthesis. [https://arxiv.org/abs/\(Wei et al., 2025b\)](https://arxiv.org/abs/(Wei et al., 2025b))
- (Liu et al., 2025) — Agent That Debugs - Dynamic State-Guided Vulnerability Repair. [https://arxiv.org/abs/\(Liu et al., 2025\)](https://arxiv.org/abs/(Liu et al., 2025))
- (Robeyns et al., 2025) — A Self-Improving Coding Agent. [https://arxiv.org/abs/\(Robeyns et al., 2025\)](https://arxiv.org/abs/(Robeyns et al., 2025))
- (Wei et al., 2025c) — VeriCoder - Enhancing LLM-Based RTL Code Generation through Functional Correctness Validation. [https://arxiv.org/abs/\(Wei et al., 2025c\)](https://arxiv.org/abs/(Wei et al., 2025c))
- (Kenneweg et al., 2025) — JEPA for RL - Investigating Joint-Embedding Predictive Architectures for Reinforcement Learning. [https://arxiv.org/abs/\(Kenneweg et al., 2025\)](https://arxiv.org/abs/(Kenneweg et al., 2025))
- (Khalifa et al., 2025) — Process Reward Models That Think. [https://arxiv.org/abs/\(Khalifa et al., 2025\)](https://arxiv.org/abs/(Khalifa et al., 2025))
- (Li et al., 2025c) — Large Language Model Powered Symbolic Execution. [https://arxiv.org/abs/\(Li et al., 2025c\)](https://arxiv.org/abs/(Li et al., 2025c))
- (Thakur et al., 2025) — CLEVER - A Curated Benchmark for Formally Verified Code Generation. [https://arxiv.org/abs/\(Thakur et al., 2025\)](https://arxiv.org/abs/(Thakur et al., 2025))
- (Chen et al., 2025a) — R1-Code-Interpreter - LLMs Reason with Code via Supervised and Multi-stage Reinforcement Learning. [https://arxiv.org/abs/\(Chen et al., 2025a\)](https://arxiv.org/abs/(Chen et al., 2025a))
- (Zhang et al., 2025a) — Darwin Godel Machine - Open-Ended Evolution of Self-Improving Agents. [https://arxiv.org/abs/\(Zhang et al., 2025a\)](https://arxiv.org/abs/(Zhang et al., 2025a))
- (Yu et al., 2025) — Dyna-Think - Synergizing Reasoning, Acting, and World Model Simulation in AI Agents. [https://arxiv.org/abs/\(Yu et al., 2025\)](https://arxiv.org/abs/(Yu et al., 2025))
- (Richens et al., 2025) — General Agents Contain World Models. [https://arxiv.org/abs/\(Richens et al., 2025\)](https://arxiv.org/abs/(Richens et al., 2025))
- (Guo et al., 2025) — World Modeling Improves Language Model Agents. [https://arxiv.org/abs/\(Guo et al., 2025\)](https://arxiv.org/abs/(Guo et al., 2025))
- (Yang et al., 2025) — Integrating Symbolic Execution with LLMs for Automated Generation of Program Specifications. [https://arxiv.org/abs/\(Yang et al., 2025\)](https://arxiv.org/abs/(Yang et al., 2025))

- (Jung et al., 2025) — Code Execution as Grounded Supervision for LLM Reasoning. [https://arxiv.org/abs/\(Jung et al., 2025\)](https://arxiv.org/abs/(Jung et al., 2025))
- (Lavon et al., 2025) — Execution Guided Line-by-Line Code Generation. [https://arxiv.org/abs/\(Lavon et al., 2025\)](https://arxiv.org/abs/(Lavon et al., 2025))
- (Da et al., 2025) — Agent-RLVR - Training Software Engineering Agents via Guidance and Environment Rewards. [https://arxiv.org/abs/\(Da et al., 2025\)](https://arxiv.org/abs/(Da et al., 2025))
- (Gong et al., 2025) — DiffuCoder - Understanding and Improving Masked Diffusion Models for Code Generation. [https://arxiv.org/abs/\(Gong et al., 2025\)](https://arxiv.org/abs/(Gong et al., 2025))
- (Chen et al., 2025b) — ChipSeek - Optimizing Verilog Generation via EDA-Integrated Reinforcement Learning. [https://arxiv.org/abs/\(Chen et al., 2025b\)](https://arxiv.org/abs/(Chen et al., 2025b))
- (Pourcel et al., 2025) — SOAR - Self-Improving Language Models for Evolutionary Program Synthesis on ARC-AGI. [https://arxiv.org/abs/\(Pourcel et al., 2025\)](https://arxiv.org/abs/(Pourcel et al., 2025))
- (Yan et al., 2025) — Re-Form - Reducing Human Priors in Scalable Formal Software Verification with RL in LLMs on Dafny. [https://arxiv.org/abs/\(Yan et al., 2025\)](https://arxiv.org/abs/(Yan et al., 2025))
- (Bharti et al., 2025) — Loop Invariant Generation - A Hybrid Framework of Reasoning Optimised LLMs and SMT Solvers. [https://arxiv.org/abs/\(Bharti et al., 2025\)](https://arxiv.org/abs/(Bharti et al., 2025))
- (Golubev et al., 2025) — Training Long-Context, Multi-Turn Software Engineering Agents with Reinforcement Learning. [https://arxiv.org/abs/\(Golubev et al., 2025\)](https://arxiv.org/abs/(Golubev et al., 2025))
- (Teng et al., 2025) — VeriRL - Boosting LLM-based Verilog Code Generation via Reinforcement Learning. [https://arxiv.org/abs/\(Teng et al., 2025\)](https://arxiv.org/abs/(Teng et al., 2025))
- (Xie et al., 2025) — Dream-Coder 7B - An Open Diffusion Language Model for Code. [https://arxiv.org/abs/\(Xie et al., 2025\)](https://arxiv.org/abs/(Xie et al., 2025))
- (Gandhi et al., 2025) — Act Like You're Paying for This - Course-Correcting Code Agents with PRMs. [https://arxiv.org/abs/\(Gandhi et al., 2025\)](https://arxiv.org/abs/(Gandhi et al., 2025))
- (Li et al., 2025a) — Jupiter - Enhancing LLM Data Analysis Capabilities via Notebook and Inference-Time Value-Guided Search. [https://arxiv.org/abs/\(Li et al., 2025a\)](https://arxiv.org/abs/(Li et al., 2025a))
- (Wang et al., 2025a) — Do Code Semantics Help - A Comprehensive Study on Execution Trace-Based Information for Code LLMs. [https://arxiv.org/abs/\(Wang et al., 2025a\)](https://arxiv.org/abs/(Wang et al., 2025a))
- (Huang et al., 2025) — LLM-JEPA - Large Language Models Meet Joint Embedding Predictive Architectures. [https://arxiv.org/abs/\(Huang et al., 2025\)](https://arxiv.org/abs/(Huang et al., 2025))
- (Wang et al., 2025c) — SALT4Decompile - Inferring Source-level Abstract Logic Tree for LLM-Based Binary Decompilation. [https://arxiv.org/abs/\(Wang et al., 2025c\)](https://arxiv.org/abs/(Wang et al., 2025c))
- (Tan et al., 2025) — SK2Decompile - LLM-Based Two-Phase Binary Decompilation. [https://arxiv.org/abs/\(Tan et al., 2025\)](https://arxiv.org/abs/(Tan et al., 2025))
- (FAIR CodeGen Team et al., 2025) — CWM - An Open-Weights LLM for Research on Code Generation with World Models. [https://arxiv.org/abs/\(FAIR CodeGen Team et al., 2025\)](https://arxiv.org/abs/(FAIR CodeGen Team et al., 2025))
- (Tahimic and Cheng, 2025) — Mechanistic Interpretability of Code Correctness in LLMs via Sparse Autoencoders. [https://arxiv.org/abs/\(Tahimic and Cheng, 2025\)](https://arxiv.org/abs/(Tahimic and Cheng, 2025))
- (Thimmaiah et al., 2025) — PLSemanticsBench - LLMs as Programming Language Interpreters. [https://arxiv.org/abs/\(Thimmaiah et al., 2025\)](https://arxiv.org/abs/(Thimmaiah et al., 2025))
- (Zhang et al., 2025b) — Memory as Action - Autonomous Context Curation for Long-Horizon Agentic Tasks. [https://arxiv.org/abs/\(Zhang et al., 2025b\)](https://arxiv.org/abs/(Zhang et al., 2025b))
- (Samadi et al., 2025) — Scaling Test-Time Compute to Achieve IOI Gold Medal with Open-Weight Models. [https://arxiv.org/abs/\(Samadi et al., 2025\)](https://arxiv.org/abs/(Samadi et al., 2025))
- (Li et al., 2025b) — A Comprehensive Survey on World Models for Embodied AI. [https://arxiv.org/abs/\(Li et al., 2025b\)](https://arxiv.org/abs/(Li et al., 2025b))
- (Wang et al., 2025d) — InspectCoder - Dynamic Analysis-Enabled Self Repair through Interactive LLM-Debugger Collaboration. [https://arxiv.org/abs/\(Wang et al., 2025d\)](https://arxiv.org/abs/(Wang et al., 2025d))
- (Wang et al., 2025b) — Huxley-Godel Machine - Human-Level Coding Agent Development. [https://arxiv.org/abs/\(Wang et al., 2025b\)](https://arxiv.org/abs/(Wang et al., 2025b))
- (Sun et al., 2025) — VeriStruct - AI-assisted Automated Verification of Data-Structure Modules in Verus. [https://arxiv.org/abs/\(Sun et al., 2025\)](https://arxiv.org/abs/(Sun et al., 2025))
- (Tu et al., 2025) — Agentic Program Verification. [https://arxiv.org/abs/\(Tu et al., 2025\)](https://arxiv.org/abs/(Tu et al., 2025))
- (Abdollahi et al., 2025) — Demystifying Errors in LLM Reasoning Traces - An Empirical Study of Code Execution Simulation. [https://arxiv.org/abs/\(Abdollahi et al., 2025\)](https://arxiv.org/abs/(Abdollahi et al., 2025))
- (Ribeiro et al., 2025) — On LLMs' Internal Representation of Code Correctness. [https://arxiv.org/abs/\(Ribeiro et al., 2025\)](https://arxiv.org/abs/(Ribeiro et al., 2025))

- (Baksys et al., 2025) — ATLAS - Automated Toolkit for Large-Scale Verified Code Synthesis. [https://arxiv.org/abs/\(Baksys et al., 2025\)](https://arxiv.org/abs/(Baksys et al., 2025))
- (Sahoo, 2025) — The Double Life of Code World Models - Provably Unmasking Malicious Behavior Through Execution Traces. [https://arxiv.org/abs/\(Sahoo, 2025\)](https://arxiv.org/abs/(Sahoo, 2025))
- (Wei et al., 2025e) — Toward Training Superintelligent Software Agents through Self-Play SWE-RL. [https://arxiv.org/abs/\(Wei et al., 2025e\)](https://arxiv.org/abs/(Wei et al., 2025e))
- (Zhang et al., 2026) — FunPRM - Function-as-Step Process Reward Model with Meta Reward Correction for Code Generation. [https://arxiv.org/abs/\(Zhang et al., 2026\)](https://arxiv.org/abs/(Zhang et al., 2026))
- (Yang, 2026) — World Models as an Intermediary between Agents and the Real World. [https://arxiv.org/abs/\(Yang, 2026\)](https://arxiv.org/abs/(Yang, 2026))
- (Sun et al., 2026) — Nanbeige SWE-World - Building Software Engineering Agents in Docker-Free Environments. [https://arxiv.org/abs/\(Sun et al., 2026\)](https://arxiv.org/abs/(Sun et al., 2026))
- (Yu et al., 2026) — Reinforcement World Model Learning for LLM-based Agents. [https://arxiv.org/abs/\(Yu et al., 2026\)](https://arxiv.org/abs/(Yu et al., 2026))
- (Rahmani, 2026) — Debugging Code World Models. [https://arxiv.org/abs/\(Rahmani, 2026\)](https://arxiv.org/abs/(Rahmani, 2026))
- (Ugare and Chandra, 2026) — Agentic Code Reasoning. [https://arxiv.org/abs/\(Ugare and Chandra, 2026\)](https://arxiv.org/abs/(Ugare and Chandra, 2026))
- (Edwards and Eslamimehr, 2026) — Synergistic Directed Execution and LLM-Driven Analysis for Zero-Day AI-Generated Malware Detection. [https://arxiv.org/abs/\(Edwards and Eslamimehr, 2026\)](https://arxiv.org/abs/(Edwards and Eslamimehr, 2026))
- (Beck et al., 2026) — Towards a Neural Debugger for Python. [https://arxiv.org/abs/\(Beck et al., 2026\)](https://arxiv.org/abs/(Beck et al., 2026))
- (Zeng et al., 2026) — Understanding by Reconstruction - Reversing the Software Development Process for LLM Pretraining. [https://arxiv.org/abs/\(Zeng et al., 2026\)](https://arxiv.org/abs/(Zeng et al., 2026))
- (Tang et al., 2026) — ExecVerify - White-Box RL with Verifiable Stepwise Rewards for Code Execution Reasoning. [https://arxiv.org/abs/\(Tang et al., 2026\)](https://arxiv.org/abs/(Tang et al., 2026))
- (Monperrus, 2026) — Bootstrapping Coding Agents - The Specification Is the Program. [https://arxiv.org/abs/\(Monperrus, 2026\)](https://arxiv.org/abs/(Monperrus, 2026))
- (Yang et al., 2026) — InCoder-32B-Thinking - Industrial Code World Model for Thinking. [https://arxiv.org/abs/\(Yang et al., 2026\)](https://arxiv.org/abs/(Yang et al., 2026))
- (Maimon et al., 2026) — Self-Execution Simulation. [https://arxiv.org/abs/\(Maimon et al., 2026\)](https://arxiv.org/abs/(Maimon et al., 2026))
- (Zhuge et al., 2026) — Neural Computers. [https://arxiv.org/abs/\(Zhuge et al., 2026\)](https://arxiv.org/abs/(Zhuge et al., 2026))
- (Han et al., 2026) — SWE-TRACE - Optimizing Long-Horizon SWE Agents Through Rubric Process Reward Models and Heuristic Test-Time Scaling. [https://arxiv.org/abs/\(Han et al., 2026\)](https://arxiv.org/abs/(Han et al., 2026))
- (Barone and Nok, 2026) — Improving LLM Code Reasoning via Semantic Equivalence Self-Play with Formal Verification. [https://arxiv.org/abs/\(Barone and Nok, 2026\)](https://arxiv.org/abs/(Barone and Nok, 2026))
- (Singh et al., 2026) — Learning Reasoning World Models for Parallel Code. [https://arxiv.org/abs/\(Singh et al., 2026\)](https://arxiv.org/abs/(Singh et al., 2026))
- (Qiu et al., 2026) — DataPRM - Process-Level Reward Modeling for Agentic Data Analysis. [https://arxiv.org/abs/\(Qiu et al., 2026\)](https://arxiv.org/abs/(Qiu et al., 2026))
- (Rodionov, 2026) — Executable World Models for ARC-AGI-3 in the Era of Coding Agents. [https://arxiv.org/abs/\(Rodionov, 2026\)](https://arxiv.org/abs/(Rodionov, 2026))
- (Li et al., 2026) — TraceEval - An Execution-Verified Multi-Language Benchmark for Code Semantic Reasoning. [https://arxiv.org/abs/\(Li et al., 2026\)](https://arxiv.org/abs/(Li et al., 2026))

References

- Abdollahi, M., Tasnia, K. R., Saha, S. K., Yang, J., Wang, S., and Hemmati, H. (2025). Demystifying errors in llm reasoning traces: An empirical study of code execution simulation.
- Armengol-Estapé, J., Carbonneaux, Q., Zhang, T., Markosyan, A. H., Seeker, V., Cummins, C., Kambadur, M., O’Boyle, M. F. P., Wang, S., Synnaeve, G., and Leather, H. J. (2025). What i cannot execute, i do not understand: Training and evaluating llms on program execution traces.

- Assran, M., Duval, Q., Misra, I., Bojanowski, P., Vincent, P., Rabbat, M., LeCun, Y., and Ballas, N. (2023). Self-supervised learning from images with a joint-embedding predictive architecture.
- Austin, J., Odena, A., Nye, M., Bosma, M., Michalewski, H., Dohan, D., Jiang, E., Cai, C., Terry, M., Le, Q., and Sutton, C. (2021). Program synthesis with large language models.
- Baksys, M., Zetsche, S., Bouissou, O., Delmas, R., Kong, S., and Holden, S. B. (2025). Atlas: Automated toolkit for large-scale verified code synthesis.
- Barone, A. V. M. and Nok, P. T. (2026). Improving llm code reasoning via semantic equivalence self-play with formal verification.
- Beck, M., Gehring, J., Kossen, J., and Synnaeve, G. (2026). Towards a neural debugger for python.
- Bharti, V., Jha, S., Kumar, D., and Jalote, P. (2025). Loop invariant generation: A hybrid framework of reasoning optimised llms and smt solvers.
- Bieber, D., Sutton, C., Larochelle, H., and Tarlow, D. (2020). Learning to execute programs with instruction pointer attention graph neural networks.
- Bruce, J., Dennis, M., Edwards, A., Parker-Holder, J., Shi, Y., Hughes, E., Lai, M., Mavalankar, A., Steigerwald, R., Apps, C., Aytar, Y., Bechtle, S., Behbahani, F., Chan, S., Heess, N., Gonzalez, L., Osindero, S., Ozair, S., Reed, S., Zhang, J., Zolna, K., Clune, J., de Freitas, N., Singh, S., and Rocktäschel, T. (2024). Genie: Generative interactive environments.
- Chae, H., Kim, N., iunn Ong, K. T., Gwak, M., Song, G., Kim, J., Kim, S., Lee, D., and Yeo, J. (2024). Web agents with world models: Learning and leveraging environment dynamics in web navigation.
- Chen, J., Pan, Z., Hu, X., Li, Z., Li, G., and Xia, X. (2024). Reasoning runtime behavior of a program with llm: How far are we?
- Chen, M., Tworek, J., Jun, H., Yuan, Q., de Oliveira Pinto, H. P., Kaplan, J., Edwards, H., Burda, Y., Joseph, N., Brockman, G., Ray, A., Puri, R., Krueger, G., Petrov, M., Khlaaf, H., Sastry, G., Mishkin, P., Chan, B., Gray, S., Ryder, N., Pavlov, M., Power, A., Kaiser, L., Bavarian, M., Winter, C., Tillet, P., Such, F. P., Cummings, D., Plappert, M., Chantzis, F., Barnes, E., Herbert-Voss, A., Guss, W. H., Nichol, A., Paino, A., Tezak, N., Tang, J., Babuschkin, I., Balaji, S., Jain, S., Saunders, W., Hesse, C., Carr, A. N., Leike, J., Achiam, J., Misra, V., Morikawa, E., Radford, A., Knight, M., Brundage, M., Murati, M., Mayer, K., Welinder, P., McGrew, B., Amodei, D., McCandlish, S., Sutskever, I., and Zaremba, W. (2021). Evaluating large language models trained on code.
- Chen, X., Lin, M., Schärli, N., and Zhou, D. (2023). Teaching large language models to self-debug.
- Chen, Y., Liu, Y., Zhou, J., Hao, Y., Wang, J., Zhang, Y., Li, N., and Fan, C. (2025a). R1-code-interpreter: Llms reason with code via supervised and multi-stage reinforcement learning.
- Chen, Z., Chang, K., Li, Z., Li, C., He, X., Chen, C., Wang, M., Xu, H., Han, Y., Li, H., and Wang, Y. (2025b). Chipseek: Optimizing verilog generation via eda-integrated reinforcement learning.
- Da, J., Wang, C., Deng, X., Ma, Y., Barhate, N., and Hendryx, S. (2025). Agent-rlvr: Training software engineering agents via guidance and environment rewards.
- Dainese, N., Merler, M., Alakuijala, M., and Marttinen, P. (2024). Generating code world models with large language models guided by monte carlo tree search.
- DeepSeek-AI, Guo, D., Yang, D., Zhang, H., Song, J., Wang, P., Zhu, Q., Xu, R., Zhang, R., Ma, S., Bi, X., Zhang, X., Yu, X., Wu, Y., Wu, Z. F., Gou, Z., Shao, Z., Li, Z., Gao, Z., Liu, A., Xue, B., Wang, B., Wu, B., Feng, B., Lu, C., Zhao, C., Deng, C., Zhang, C., Ruan, C., Dai, D., Chen, D., Ji, D., Li, E., Lin, F., Dai, F., Luo, F., Hao, G., Chen, G., Li, G., Zhang, H., Bao, H., Xu, H., Wang, H., Ding, H., Xin, H., Gao, H., Qu, H., Li, H., Guo, J., Li, J., Wang, J., Chen, J., Yuan, J., Qiu, J., Li, J., Cai, J. L., Ni, J., Liang, J., Chen, J., Dong, K., Hu, K., Gao, K., Guan, K., Huang, K., Yu,

- K., Wang, L., Zhang, L., Zhao, L., Wang, L., Zhang, L., Xu, L., Xia, L., Zhang, M., Zhang, M., Tang, M., Li, M., Wang, M., Li, M., Tian, N., Huang, P., Zhang, P., Wang, Q., Chen, Q., Du, Q., Ge, R., Zhang, R., Pan, R., Wang, R., Chen, R. J., Jin, R. L., Chen, R., Lu, S., Zhou, S., Chen, S., Ye, S., Wang, S., Yu, S., Zhou, S., Pan, S., Li, S. S., Zhou, S., Wu, S., Ye, S., Yun, T., Pei, T., Sun, T., Wang, T., Zeng, W., Zhao, W., Liu, W., Liang, W., Gao, W., Yu, W., Zhang, W., Xiao, W. L., An, W., Liu, X., Wang, X., Chen, X., Nie, X., Cheng, X., Liu, X., Xie, X., Liu, X., Yang, X., Li, X., Su, X., Lin, X., Li, X. Q., Jin, X., Shen, X., Chen, X., Sun, X., Wang, X., Song, X., Zhou, X., Wang, X., Shan, X., Li, Y. K., Wang, Y. Q., Wei, Y. X., Zhang, Y., Xu, Y., Li, Y., Zhao, Y., Sun, Y., Wang, Y., Yu, Y., Zhang, Y., Shi, Y., Xiong, Y., He, Y., Piao, Y., Wang, Y., Tan, Y., Ma, Y., Liu, Y., Guo, Y., Ou, Y., Wang, Y., Gong, Y., Zou, Y., He, Y., Xiong, Y., Luo, Y., You, Y., Liu, Y., Zhou, Y., Zhu, Y. X., Xu, Y., Huang, Y., Li, Y., Zheng, Y., Zhu, Y., Ma, Y., Tang, Y., Zha, Y., Yan, Y., Ren, Z. Z., Ren, Z., Sha, Z., Fu, Z., Xu, Z., Xie, Z., Zhang, Z., Hao, Z., Ma, Z., Yan, Z., Wu, Z., Gu, Z., Zhu, Z., Liu, Z., Li, Z., Xie, Z., Song, Z., Pan, Z., Huang, Z., Xu, Z., Zhang, Z., and Zhang, Z. (2025). Deepseek-r1: Incentivizing reasoning capability in llms via reinforcement learning.
- Deng, X., Gu, Y., Zheng, B., Chen, S., Stevens, S., Wang, B., Sun, H., and Su, Y. (2023). Mind2web: Towards a generalist agent for the web.
- Ding, J., Zhang, Y., Shang, Y., Feng, J., Zhang, Y., Zong, Z., Yuan, Y., Su, H., Li, N., Piao, J., Deng, Y., Sukiennik, N., Gao, C., Xu, F., and Li, Y. (2024a). Understanding world or predicting future? a comprehensive survey of world models.
- Ding, Y., Peng, J., Min, M. J., Kaiser, G., Yang, J., and Ray, B. (2024b). Semcoder: Training code language models with comprehensive semantics reasoning.
- Ding, Y., Steenhoek, B., Pei, K., Kaiser, G., Le, W., and Ray, B. (2023). Traced: Execution-aware pre-training for source code.
- Edwards, G. and Eslamimehr, M. (2026). Synergistic directed execution and llm-driven analysis for zero-day ai-generated malware detection.
- FAIR CodeGen Team, Copet, J., Carbonneaux, Q., Cohen, G., Gehring, J., Kahn, J., Kossen, J., Kreuk, F., McMilin, E., Meyer, M., Wei, Y., Zhang, D., Zheng, K., Armengol-Estapé, J., Bashiri, P., Beck, M., Chambon, P., Charnalia, A., Cummins, C., Decugis, J., Fisches, Z. V., Fleuret, F., Gloeckle, F., Gu, A., Hassid, M., Haziza, D., Idrissi, B. Y., Keller, C., Kindi, R., Leather, H., Maimon, G., Markosyan, A., Massa, F., Mazaré, P.-E., Mella, V., Murray, N., Muzumdar, K., O’Hearn, P., Pagliardini, M., Pedchenko, D., Remez, T., Seeker, V., Selvi, M., Sultan, O., Wang, S., Wehrstedt, L., Yoran, O., Zhang, L., Cohen, T., Adi, Y., and Synnaeve, G. (2025). Cwm: An open-weights llm for research on code generation with world models.
- Feng, X., Wan, Z., Wen, M., McAleer, S. M., Wen, Y., Zhang, W., and Wang, J. (2023). Alphazero-like tree-search can guide large language model decoding and training.
- Gandhi, S., Tsay, J., Ganhotra, J., Kate, K., and Rizk, Y. (2025). When agents go astray: Course-correcting swe agents with prms.
- Gehring, J., Zheng, K., Copet, J., Mella, V., Carbonneaux, Q., Cohen, T., and Synnaeve, G. (2024). Rlef: Grounding code llms in execution feedback with reinforcement learning.
- Golubev, A., Trofimova, M., Polezhaev, S., Badertdinov, I., Nekrashevich, M., Shevtsov, A., Karasik, S., Abramov, S., Andriushchenko, A., Fisin, F., Skvortsov, S., and Yangel, B. (2025). Training long-context, multi-turn software engineering agents with reinforcement learning.
- Gong, S., Zhang, R., Zheng, H., Gu, J., Jaitly, N., Kong, L., and Zhang, Y. (2025). Diffucoder: Understanding and improving masked diffusion models for code generation.
- Gu, A., Rozière, B., Leather, H., Solar-Lezama, A., Synnaeve, G., and Wang, S. I. (2024a). Crux-eval: A benchmark for code reasoning, understanding and execution.
- Gu, Y., Zhang, K., Ning, Y., Zheng, B., Gou, B., Xue, T., Chang, C., Srivastava, S., Xie, Y., Qi, P., Sun, H., and Su, Y. (2024b). Is your llm secretly a world model of the internet? model-based planning for web agents.

- Guo, S., Domingues, O. D., Avalos, R., Courville, A., and Strub, F. (2025). World modelling improves language model agents.
- Ha, D. and Schmidhuber, J. (2018). World models.
- Hafner, D., Lillicrap, T., Ba, J., and Norouzi, M. (2019). Dream to control: Learning behaviors by latent imagination.
- Hafner, D., Lillicrap, T., Norouzi, M., and Ba, J. (2020). Mastering atari with discrete world models.
- Hafner, D., Pasukonis, J., Ba, J., and Lillicrap, T. (2023). Mastering diverse domains through world models.
- Han, H., Xie, J., Ma, X., Zhu, W., Zhang, Z., Long, Z., Chen, H., and Ye, Q. (2026). Swe-trace: Optimizing long-horizon swe agents through rubric process reward models and heuristic test-time scaling.
- Hao, S., Gu, Y., Ma, H., Hong, J. J., Wang, Z., Wang, D. Z., and Hu, Z. (2023). Reasoning with language model is planning with world model.
- Huang, H., LeCun, Y., and Balestriero, R. (2025). Llm-jepa: Large language models meet joint embedding predictive architectures.
- Jain, N., Han, K., Gu, A., Li, W.-D., Yan, F., Zhang, T., Wang, S., Solar-Lezama, A., Sen, K., and Stoica, I. (2024). Livecodebench: Holistic and contamination free evaluation of large language models for code.
- Jia, C., Li, Z., Wang, P., Li, Y.-C., Hou, Z., Dong, Y., and Yu, Y. (2025). Controlling large language model with latent actions.
- Jiang, J., Wang, F., Shen, J., Kim, S., and Kim, S. (2024). A survey on large language models for code generation.
- Jimenez, C. E., Yang, J., Wettig, A., Yao, S., Pei, K., Press, O., and Narasimhan, K. (2023). Swe-bench: Can language models resolve real-world github issues?
- Jung, D., Zhou, W., and Chen, M. (2025). Code execution as grounded supervision for llm reasoning.
- Kenneweg, T., Kenneweg, P., and Hammer, B. (2025). Jepa for rl: Investigating joint-embedding predictive architectures for reinforcement learning.
- Khalifa, M., Agarwal, R., Logeswaran, L., Kim, J., Peng, H., Lee, M., Lee, H., and Wang, L. (2025). Process reward models that think.
- Koh, J. Y., McAleer, S., Fried, D., and Salakhutdinov, R. (2024). Tree search for language model agents.
- Lavon, B., Katz, S., and Wolf, L. (2025). Execution guided line-by-line code generation.
- Le, H., Wang, Y., Gotmare, A. D., Savarese, S., and Hoi, S. C. H. (2022). Coderl: Mastering code generation through pretrained models and deep reinforcement learning.
- Li, S., Liu, Y., Du, S., Zeng, W., Xu, Z., Zhou, M., He, Y., Dong, H., Han, S., and Zhang, D. (2025a). Jupiter: Enhancing llm data analysis capabilities via notebook and inference-time value-guided search.
- Li, X., He, X., Zhang, L., Wu, M., Li, X., and Liu, Y. (2025b). A comprehensive survey on world models for embodied ai.
- Li, Y., Jiang, J., Zhang, T., Yang, C., Zhong, C., Yide, Y., Bin, L. W., Ouh, E. L., Shar, L. K., and Lo, D. (2026). An execution-verified multi-language benchmark for code semantic reasoning.
- Li, Y., Meng, R., and Duck, G. J. (2025c). Large language model powered symbolic execution.

- Liu, A. Z., Wang, X., Sansom, J., Fu, Y., Choi, J., Sohn, S., Kim, J., and Lee, H. (2024). Interactive and expressive code-augmented planning with large language models.
- Liu, C., Lu, S., Chen, W., Jiang, D., Svyatkovskiy, A., Fu, S., Sundaresan, N., and Duan, N. (2023). Code execution with pre-trained language models.
- Liu, Z., Ma, Y., Xu, J., Ai, J., Gao, X., Sun, H., and Roychoudhury, A. (2025). Agent that debugs: Dynamic state-guided vulnerability repair.
- Maimon, G., Yoran, O., Kreuk, F., Hassid, M., Cohen, G., Chambon, P., and Adi, Y. (2026). Self-execution simulation improves coding models.
- Mitchell, J. L., Kim, B. H., Zhou, C., and Wang, C. (2025). Understanding formal reasoning failures in llms as abstract interpreters.
- Monperrus, M. (2026). Bootstrapping coding agents: The specification is the program.
- Ni, A., Allamanis, M., Cohan, A., Deng, Y., Shi, K., Sutton, C., and Yin, P. (2024). Next: Teaching large language models to reason about code execution.
- Ni, A., Iyer, S., Radev, D., Stoyanov, V., tau Yih, W., Wang, S. I., and Lin, X. V. (2023). Lever: Learning to verify language-to-code generation with execution.
- Nye, M., Andreassen, A. J., Gur-Ari, G., Michalewski, H., Austin, J., Bieber, D., Dohan, D., Lewkowycz, A., Bosma, M., Luan, D., Sutton, C., and Odena, A. (2021). Show your work: Scratchpads for intermediate computation with language models.
- Olausson, T. X., Inala, J. P., Wang, C., Gao, J., and Solar-Lezama, A. (2023). Is self-repair a silver bullet for code generation?
- Ouyang, S., Yu, W., Ma, K., Xiao, Z., Zhang, Z., Jia, M., Han, J., Zhang, H., and Yu, D. (2024). Repograph: Enhancing ai software engineering with repository-level code graph.
- Pan, J., Wang, X., Neubig, G., Jaitly, N., Ji, H., Suhr, A., and Zhang, Y. (2024). Training software engineering agents and verifiers with swe-gym.
- Pink, M., Wu, Q., Vo, V. A., Turek, J., Mu, J., Huth, A., and Toneva, M. (2025). Position: Episodic memory is the missing piece for long-term llm agents.
- Pourcel, J., Colas, C., and Oudeyer, P.-Y. (2025). Self-improving language models for evolutionary program synthesis: A case study on arc-agi.
- Pu, Y., Niu, Y., Yang, Z., Ren, J., Li, H., and Liu, Y. (2024). Unizero: Generalized and efficient planning with scalable latent world models.
- Qiu, Z., Qiao, S., Xu, K., Zhu, Y., Du, L., Zhang, N., and Chen, H. (2026). Rewarding the scientific process: Process-level reward modeling for agentic data analysis.
- Rahmani, B. (2026). Debugging code world models.
- Reed, S. and de Freitas, N. (2015). Neural programmer-interpreters.
- Ribeiro, F., Spiess, C., Devanbu, P., and Nadi, S. (2025). On llms’ internal representation of code correctness.
- Richens, J., Abel, D., Bellot, A., and Everitt, T. (2025). General agents contain world models.
- Robeyns, M., Szummer, M., and Aitchison, L. (2025). A self-improving coding agent.
- Rodionov, S. (2026). Executable world models for arc-agi-3 in the era of coding agents.
- Sahoo, S. (2025). The double life of code world models: Provably unmasking malicious behavior through execution traces.
- Samadi, M., Ficek, A., Narenthiran, S., Jain, S., Ahmad, W. U., Majumdar, S., Noroozi, V., and Ginsburg, B. (2025). Scaling test-time compute to achieve ioi gold medal with open-weight models.

- Schultz, J., Adamek, J., Jusup, M., Lanctot, M., Kaisers, M., Perrin, S., Hennes, D., Shar, J., Lewis, C., Ruoss, A., Zahavy, T., Veličković, P., Prince, L., Singh, S., Malmi, E., and Tomašev, N. (2024). Mastering board games by external and internal planning with language models.
- Shi, Z., Swersky, K., Tarlow, D., Ranganathan, P., and Hashemi, M. (2019). Learning execution through neural code fusion.
- Shinn, N., Cassano, F., Berman, E., Gopinath, A., Narasimhan, K., and Yao, S. (2023). Reflexion: Language agents with verbal reinforcement learning.
- Singh, G., Guha, A., Kailkhura, B., and Menon, H. (2026). Learning reasoning world models for parallel code.
- Sun, C., Sun, Y., Amrollahi, D., Zhang, E., Lahiri, S., Lu, S., Dill, D., and Barrett, C. (2025). Veristruct: Ai-assisted automated verification of data-structure modules in verus.
- Sun, S., Song, H., Huang, L., Jiang, J., Le, R., Lv, Z., Chen, Z., Hu, Y., Luo, W., Zhao, W. X., Song, Y., Xu, H., Zhang, T., and Wen, J.-R. (2026). Swe-world: Building software engineering agents in docker-free environments.
- Tahimic, K. and Cheng, C. (2025). Mechanistic interpretability of code correctness in llms via sparse autoencoders.
- Tan, H., Li, W., Tian, X., Wang, S., Liu, J., Li, J., and Zhang, Y. (2025). Sk2decompile: Llm-based two-phase binary decompilation from skeleton to skin.
- Tang, L., Ye, H., Chu, Z., Ye, M., Liu, Z., Ren, X., and Bao, L. (2026). Execverify: White-box rl with verifiable stepwise rewards for code execution reasoning.
- Teng, F., Pan, M., Zhang, X., He, Z., Yang, Y., Chai, X., Qi, M., Lu, L., and Yin, J. (2025). Verirl: Boosting the llm-based verilog code generation via reinforcement learning.
- Thakur, A., Lee, J., Tsoukalas, G., Sistla, M., Zhao, M., Zetsche, S., Durrett, G., Yue, Y., and Chaudhuri, S. (2025). Clever: A curated benchmark for formally verified code generation.
- Thimmaiah, A., Zhang, J., Srinivasa, J., Li, J. J., and Gligoric, M. (2025). Plsemanticsbench: Large language models as programming language interpreters.
- Tu, H., Zhao, H., Song, Y., Zafar, M., Meng, R., and Roychoudhury, A. (2025). Agentic verification of software systems.
- Ugare, S. and Chandra, S. (2026). Agentic code reasoning.
- Wang, J., Xie, X., Hu, Q., Liu, S., and Li, Y. (2025a). Do code semantics help? a comprehensive study on execution trace-based information for code large language models.
- Wang, K., Singh, R., and Su, Z. (2017). Dynamic neural program embedding for program repair.
- Wang, W., Liu, K., Chen, A. R., Li, G., Jin, Z., Huang, G., and Ma, L. (2024a). Python symbolic execution with llm-powered code generation.
- Wang, W., Piękos, P., Nanbo, L., Laakom, F., Chen, Y., Ostaszewski, M., Zhuge, M., and Schmidhuber, J. (2025b). Huxley-gödel machine: Human-level coding agent development by an approximation of the optimal self-improving machine.
- Wang, X., Chen, Y., Yuan, L., Zhang, Y., Li, Y., Peng, H., and Ji, H. (2024b). Executable code actions elicit better llm agents.
- Wang, Y., Xu, X., Zhu, X., Gu, X., and Shen, B. (2025c). Salt4decompile: Inferring source-level abstract logic tree for llm-based binary decompilation.
- Wang, Y., Zhang, Y., Li, G., Zhi, C., Li, B., Huang, F., Li, Y., and Deng, S. (2025d). Inspectcoder: Dynamic analysis-enabled self repair through interactive llm-debugger collaboration.

- Wei, A., Cao, J., Li, R., Chen, H., Zhang, Y., Wang, Z., Liu, Y., Teixeira, T. S. F. X., Yang, D., Wang, K., and Aiken, A. (2025a). Equibench: Benchmarking large language models’ reasoning about program semantics via equivalence checking.
- Wei, A., Suresh, T., Cao, J., Kannan, N., Wu, Y., Yan, K., Teixeira, T. S. F. X., Wang, K., and Aiken, A. (2025b). Codearc: Benchmarking reasoning capabilities of llm agents for inductive program synthesis.
- Wei, A., Tan, H., Suresh, T., Mendoza, D., Teixeira, T. S. F. X., Wang, K., Trippel, C., and Aiken, A. (2025c). Vericoder: Enhancing llm-based rtl code generation through functional correctness validation.
- Wei, Y., Duchenne, O., Copet, J., Carbonneaux, Q., Zhang, L., Fried, D., Synnaeve, G., Singh, R., and Wang, S. I. (2025d). Swe-rl: Advancing llm reasoning via reinforcement learning on open software evolution.
- Wei, Y., Sun, Z., McMilin, E., Gehring, J., Zhang, D., Synnaeve, G., Fried, D., Zhang, L., and Wang, S. (2025e). Toward training superintelligent software agents through self-play swe-rl.
- Xia, C. S., Deng, Y., Dunn, S., and Zhang, L. (2024). Agentless: Demystifying llm-based software engineering agents.
- Xie, Z., Ye, J., Zheng, L., Gao, J., Dong, J., Wu, Z., Zhao, X., Gong, S., Jiang, X., Li, Z., and Kong, L. (2025). Dream-coder 7b: An open diffusion language model for code.
- Xu, R., Cao, J., Lu, Y., Wen, M., Lin, H., Han, X., He, B., Cheung, S.-C., and Sun, L. (2024). Cruxeval-x: A benchmark for multilingual code reasoning, understanding and execution.
- Yan, C., Che, F., Huang, X., Xu, X., Li, X., Li, Y., Qu, X., Shi, J., Lin, C., Yang, Y., Yuan, B., Zhao, H., Qiao, Y., Zhou, B., and Fu, J. (2025). Re:form – reducing human priors in scalable formal software verification with rl in llms: A preliminary study on dafny.
- Yan, S., Wang, S., Duan, Y., Hong, H., Lee, K., Kim, D., and Hong, Y. (2024). An llm-assisted easy-to-trigger backdoor attack on code completion models: Injecting disguised vulnerabilities against strong detection.
- Yang, F., Ma, X., Wang, S., Xu, X., Cao, Q., Zhan, N., Li, X., and Gu, B. (2025). Integrating symbolic execution with llms for automated generation of program specifications.
- Yang, J., Jimenez, C. E., Wettig, A., Lieret, K., Yao, S., Narasimhan, K., and Press, O. (2024). Swe-agent: Agent-computer interfaces enable automated software engineering.
- Yang, J., Zhang, W., Wu, J., Cheng, J., Zheng, T., Xu, F., Gu, W., Jing, L., Du, Y., Li, J., Li, Y., Xing, Y., Hao, C., Tao, R., Gong, R., Liu, A., Li, Z., Tang, M., Lin, C., Chen, S., Zhao, W. X., Liu, X., Zhou, M., Dai, B., and Lv, W. (2026). Incoder-32b-thinking: Industrial code world model for thinking.
- Yang, S. (2026). World models as an intermediary between agents and the real world.
- Yao, S., Yu, D., Zhao, J., Shafran, I., Griffiths, T. L., Cao, Y., and Narasimhan, K. (2023). Tree of thoughts: Deliberate problem solving with large language models.
- Yu, X., Peng, B., Xu, R., Galley, M., Cheng, H., Nath, S., Gao, J., and Yu, Z. (2025). Dyna-think: Synergizing reasoning, acting, and world model simulation in ai agents.
- Yu, X., Peng, B., Xu, R., Shen, Y., He, P., Nath, S., Singh, N., Gao, J., and Yu, Z. (2026). Reinforcement world model learning for llm-based agents.
- Zaremba, W. and Sutskever, I. (2014). Learning to execute.
- Zeng, Z., Zhang, Y., Shan, Y., Hua, K., Fang, S., Liu, Z., Liu, J., Wang, H., Zheng, Y., Ding, M., Shen, K., Zhang, G., Huang, W., and Qiu, X. (2026). Understanding by reconstruction: Reversing the software development process for llm pretraining.

- Zhang, J., Hu, S., Lu, C., Lange, R., and Clune, J. (2025a). Darwin godel machine: Open-ended evolution of self-improving agents.
- Zhang, R., Qin, P., Cao, Q., Xue, E., and Xie, P. (2026). Funprm: Function-as-step process reward model with meta reward correction for code generation.
- Zhang, Y., Pan, Y., Wang, Y., and Cai, J. (2024a). Pybench: Evaluating llm agent on various real-world coding tasks.
- Zhang, Y., Ruan, H., Fan, Z., and Roychoudhury, A. (2024b). Autocoderover: Autonomous program improvement.
- Zhang, Y., Shu, J., Ma, Y., Lin, X., Wu, S., and Sang, J. (2025b). Memory as action: Autonomous context curation for long-horizon agentic tasks.
- Zhang, Y., Wu, S., Yang, Y., Shu, J., Xiao, J., Kong, C., and Sang, J. (2024c). o1-coder: an o1 replication for coding.
- Zhou, S., Xu, F. F., Zhu, H., Zhou, X., Lo, R., Sridhar, A., Cheng, X., Ou, T., Bisk, Y., Fried, D., Alon, U., and Neubig, G. (2023). Webarena: A realistic web environment for building autonomous agents.
- Zhuge, M., Zhao, C., Liu, H., Zhou, Z., Liu, S., Wang, W., Chang, E., Lan, G. L., Fei, J., Zhang, W., Sun, Y., Cai, Z., Liu, Z., Xiong, Y., Yang, Y., Tian, Y., Shi, Y., Chandra, V., and Schmidhuber, J. (2026). Neural computers.